

# CICN

# Community

# Information-Centric Networking



Tutorial at ACM SIGCOMM ICN, Berlin, Germany  
26<sup>th</sup> of September 2017



# Tutorial agenda

- Project overview
- Vector Packet Processing
- vICN: automation of virtual ICN network deployment
- The consumer/producer socket API with applications to HTTP





# CICN project overview

- CCNx Internet documents are specified at the ICNRG define the architecture.
- The rest is just software development, testing and experimentation.
- Focus on VPP and application development:
  - Vector Packet Processing as the Universal Data Plane for vRouting and vSwitching
  - vICN automation of virtual networks deployment
  - The Consumer/Producer Socket API and HTTP



# What is FD.io (pronounced “fido”)?





# FD.io: The Universal Dataplane

- Project at Linux Foundation

- Multi-party
- Multi-project

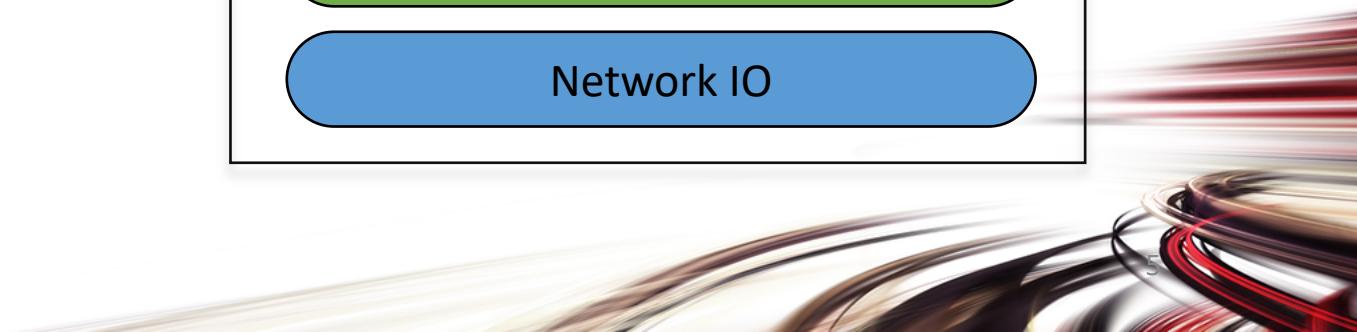
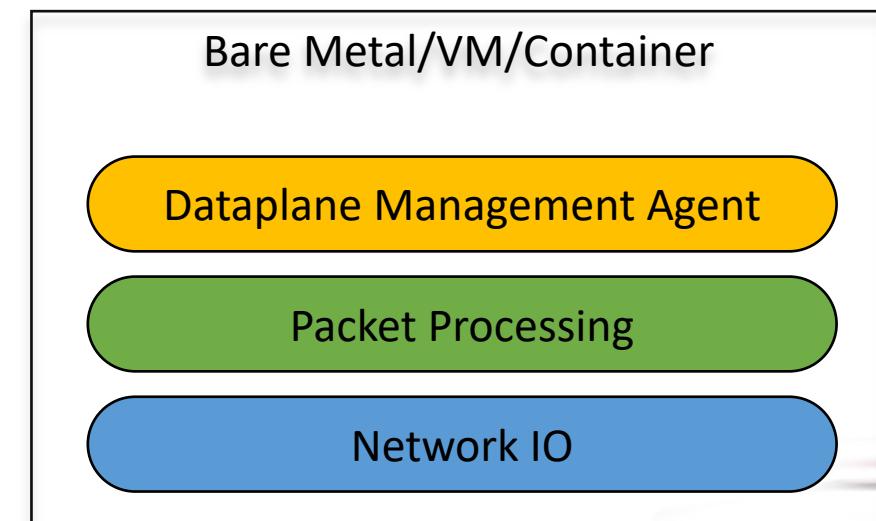
- Software Dataplane

- High throughput
- Low Latency
- Feature Rich
- Resource Efficient
- Bare Metal/VM/Container
- Multiplatform

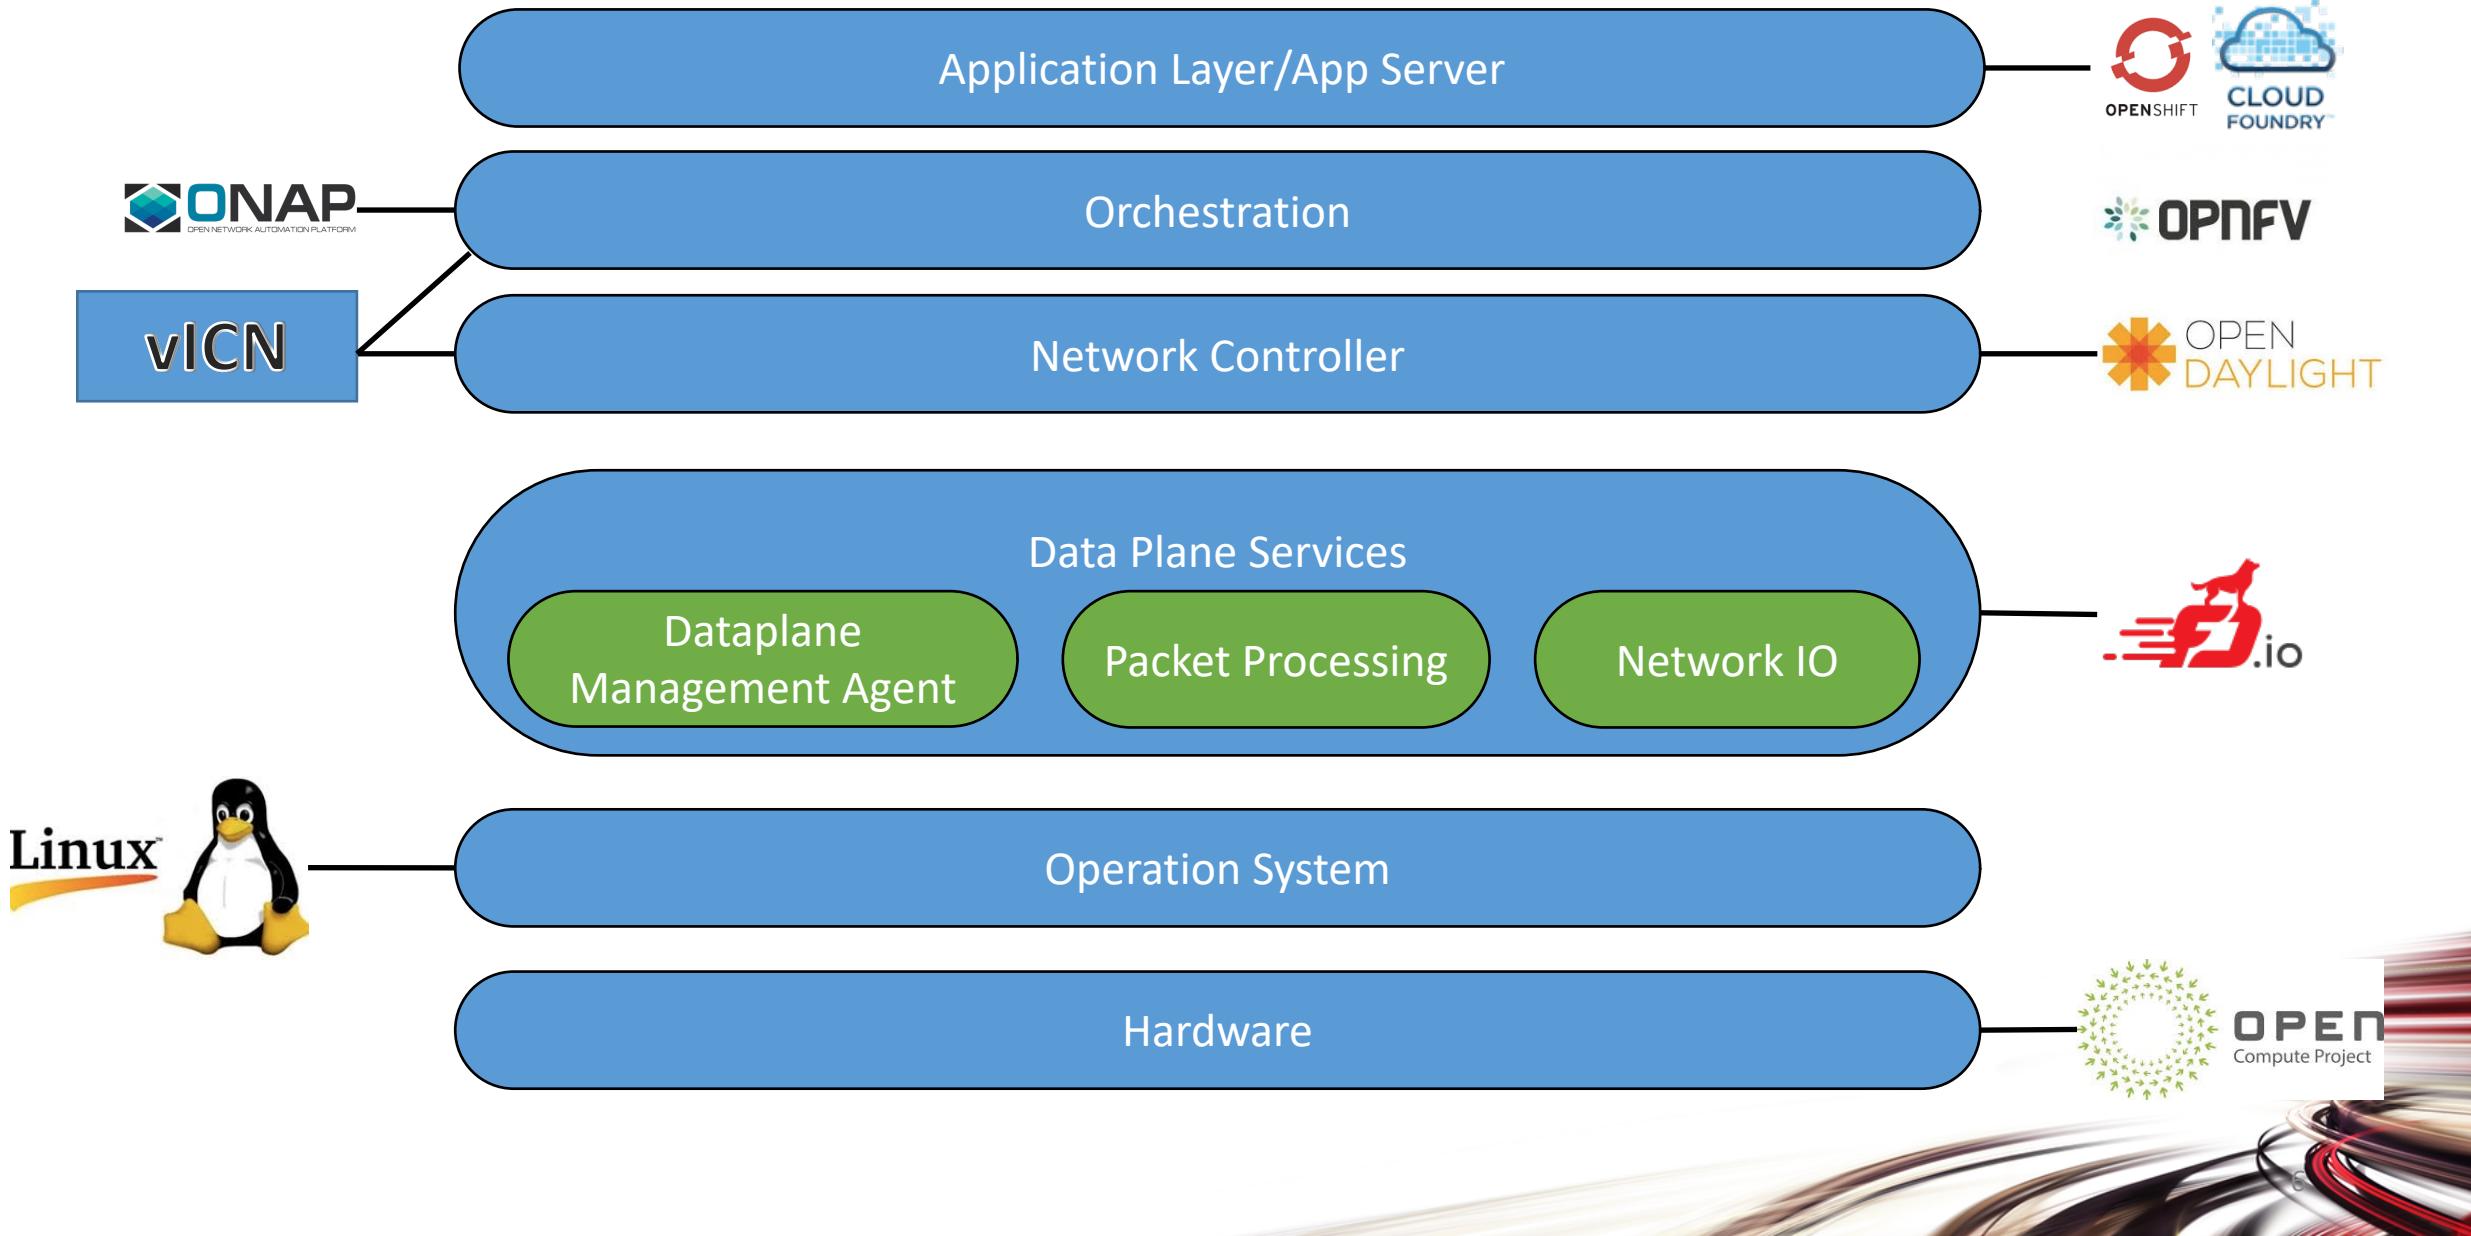


- Fd.io Scope:

- **Network IO** - NIC/vNIC <-> cores/threads
- **Packet Processing** – Classify/Transform/Prioritize/Forward/Terminate
- **Dataplane Management Agents** - ControlPlane



# Fd.io in the overall stack





# Multiparty: Broad Membership

## Service Providers



## Network Vendors



## Chip Vendors



## Integrators





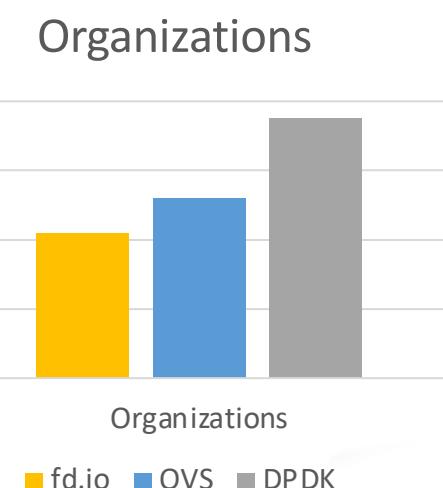
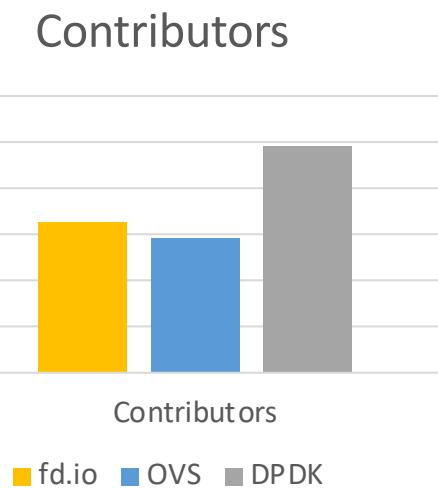
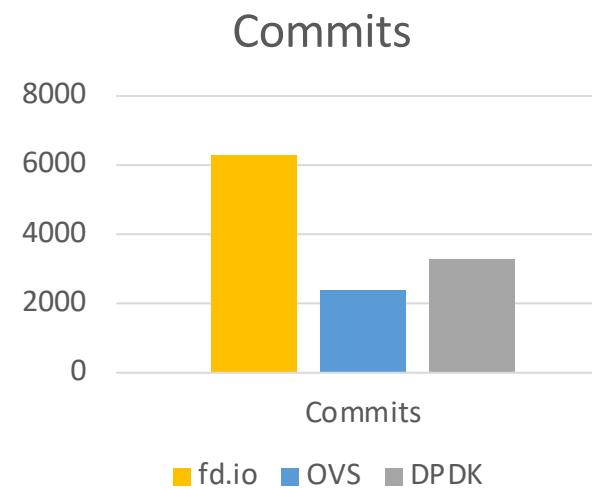
# Multiparty: Broad Contribution



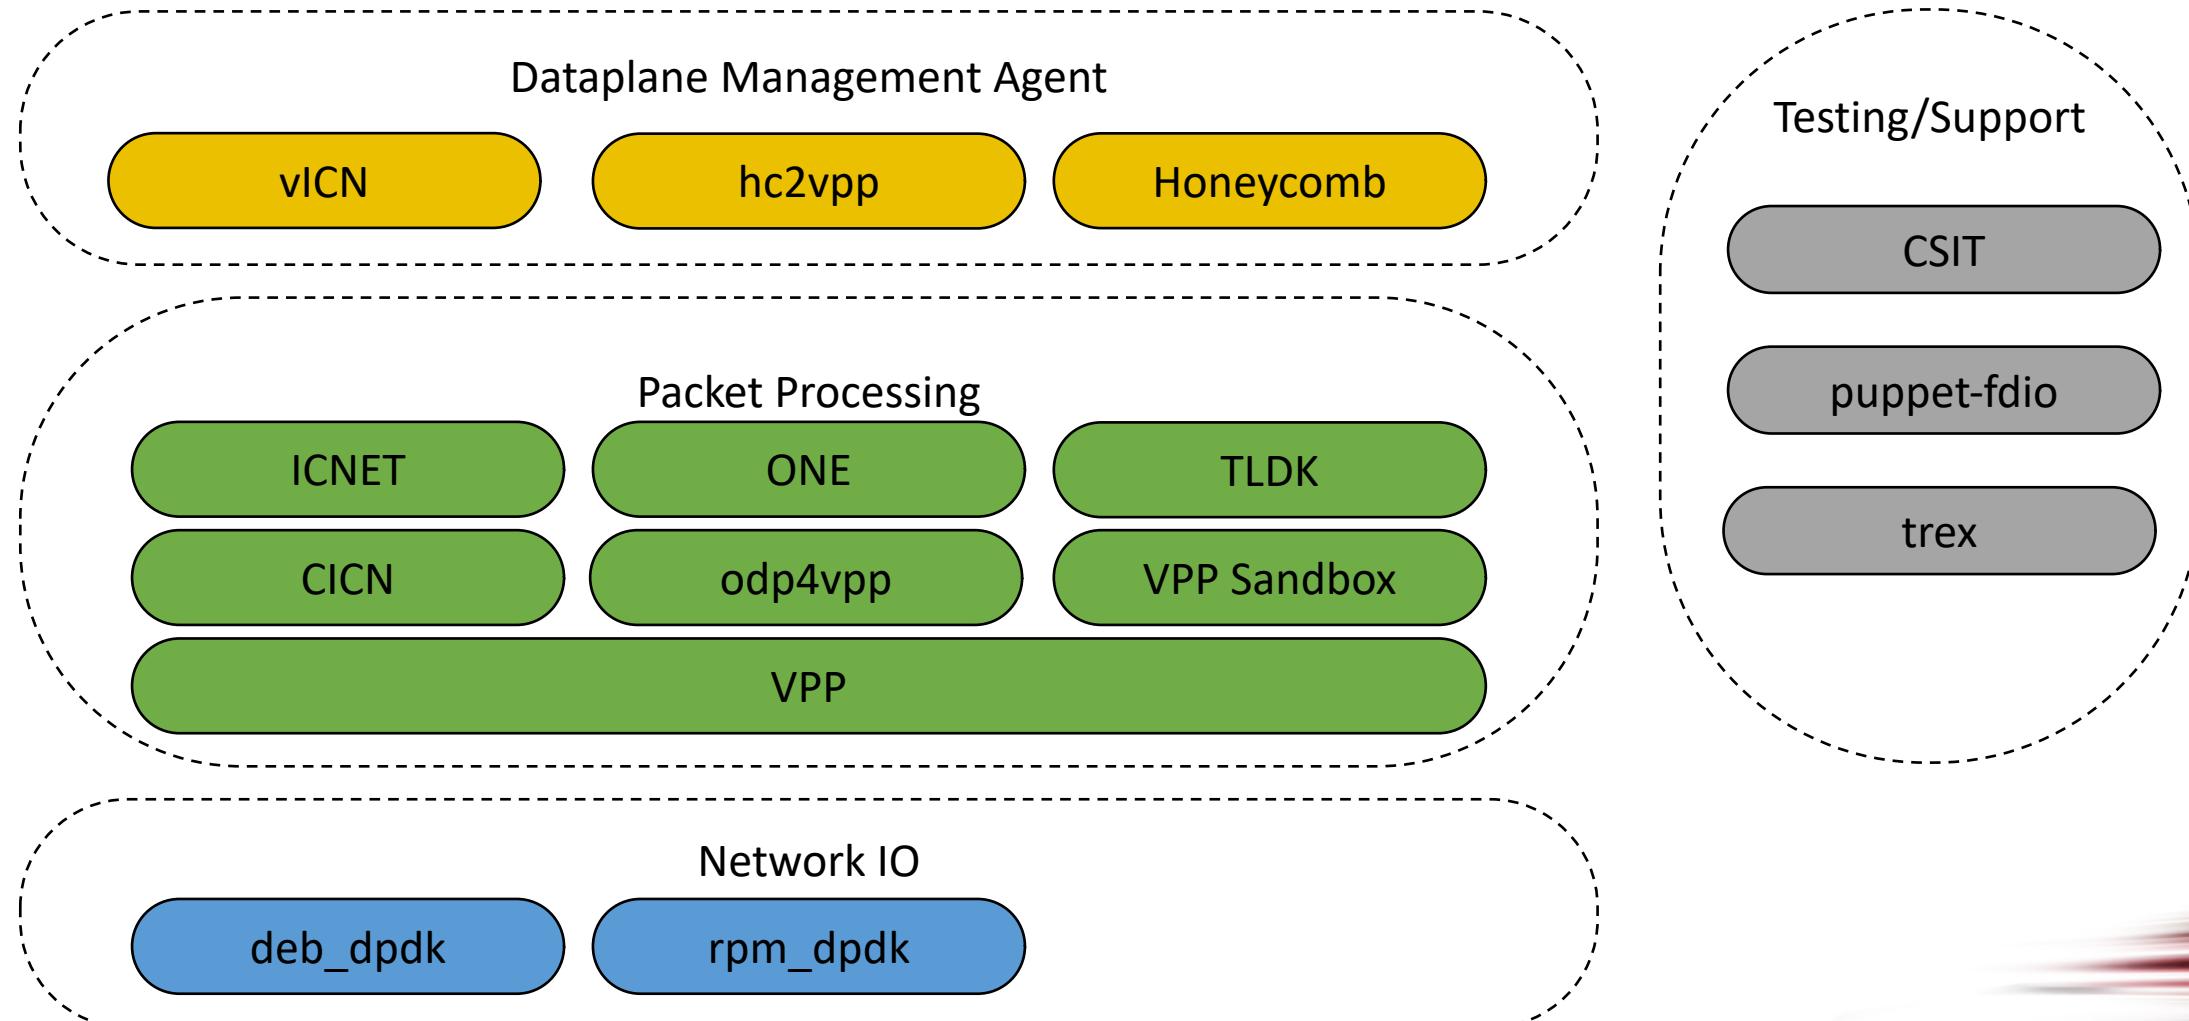
# Code Activity

- In the period since its inception, fd.io has more commits than OVS and DPDK combined, and more contributors than OVS

2016-02-11 to 2017-04-03	Fd.io	OVS	DPDK
Commits	6283	2395	3289
Contributors	163	146	245
Organizations	42	52	78



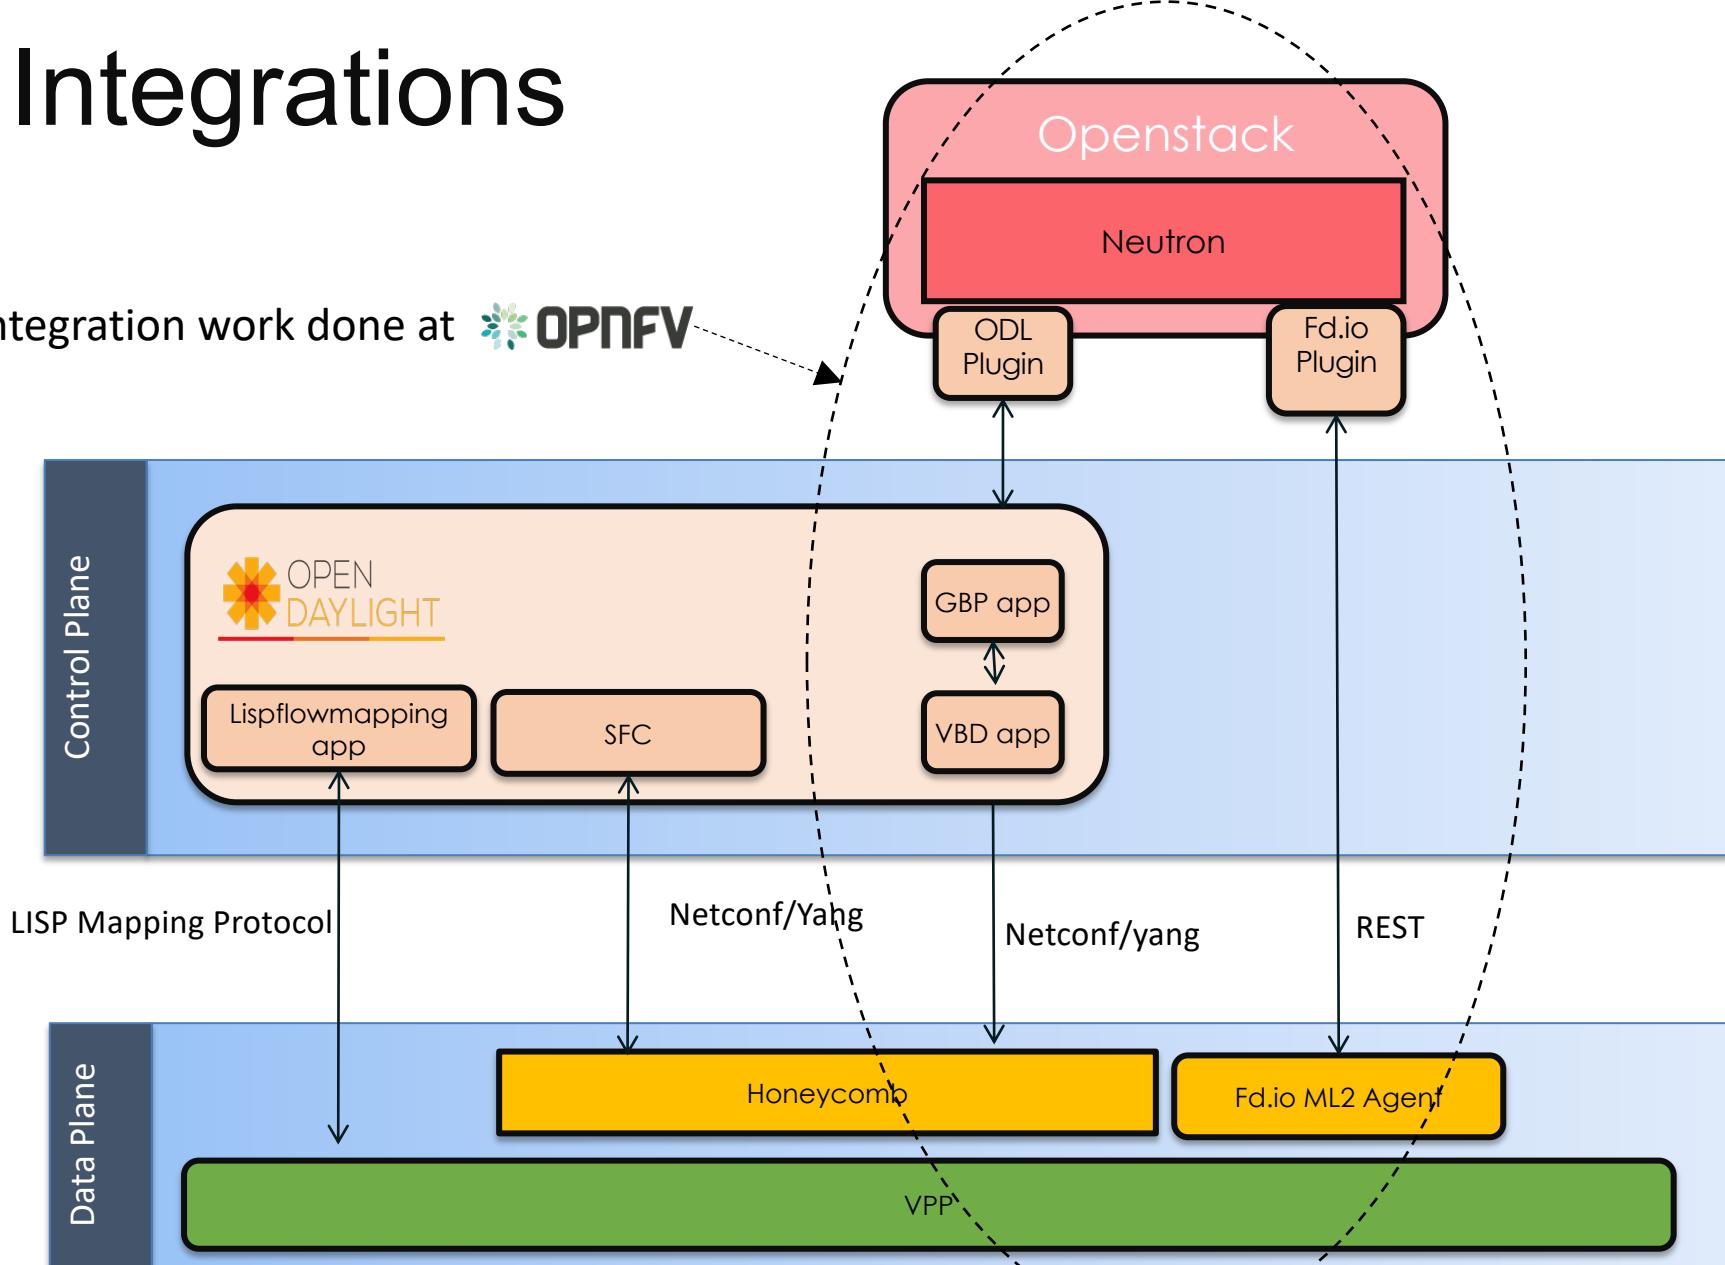
# Multiproject: Fd.io Projects



# Fd.io Integrations

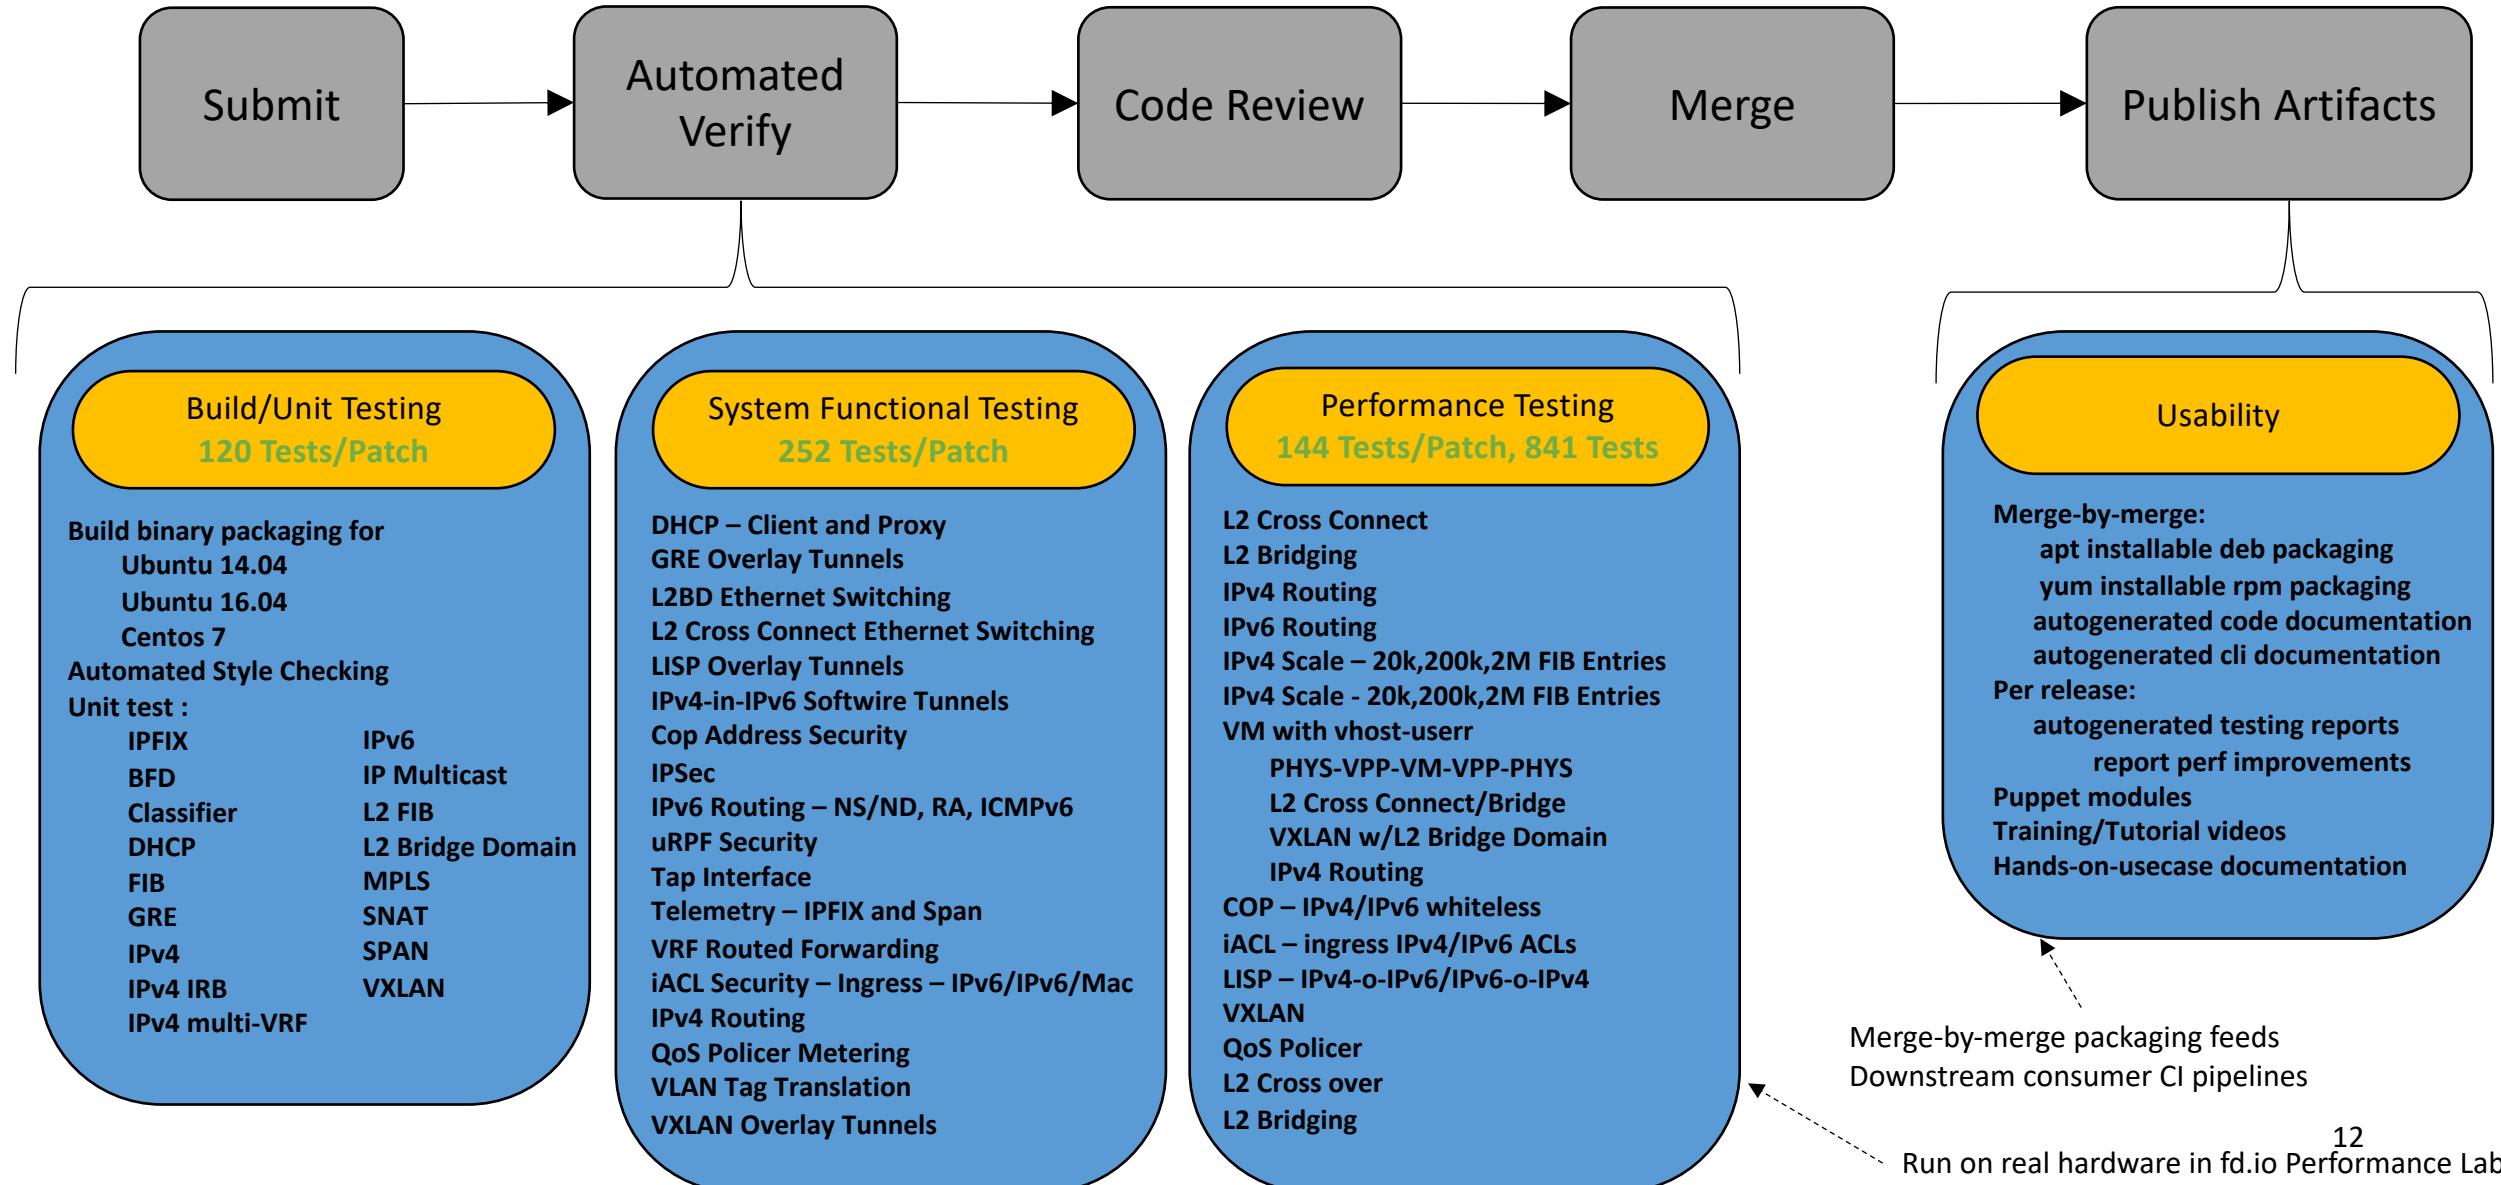


Integration work done at  OPNFV

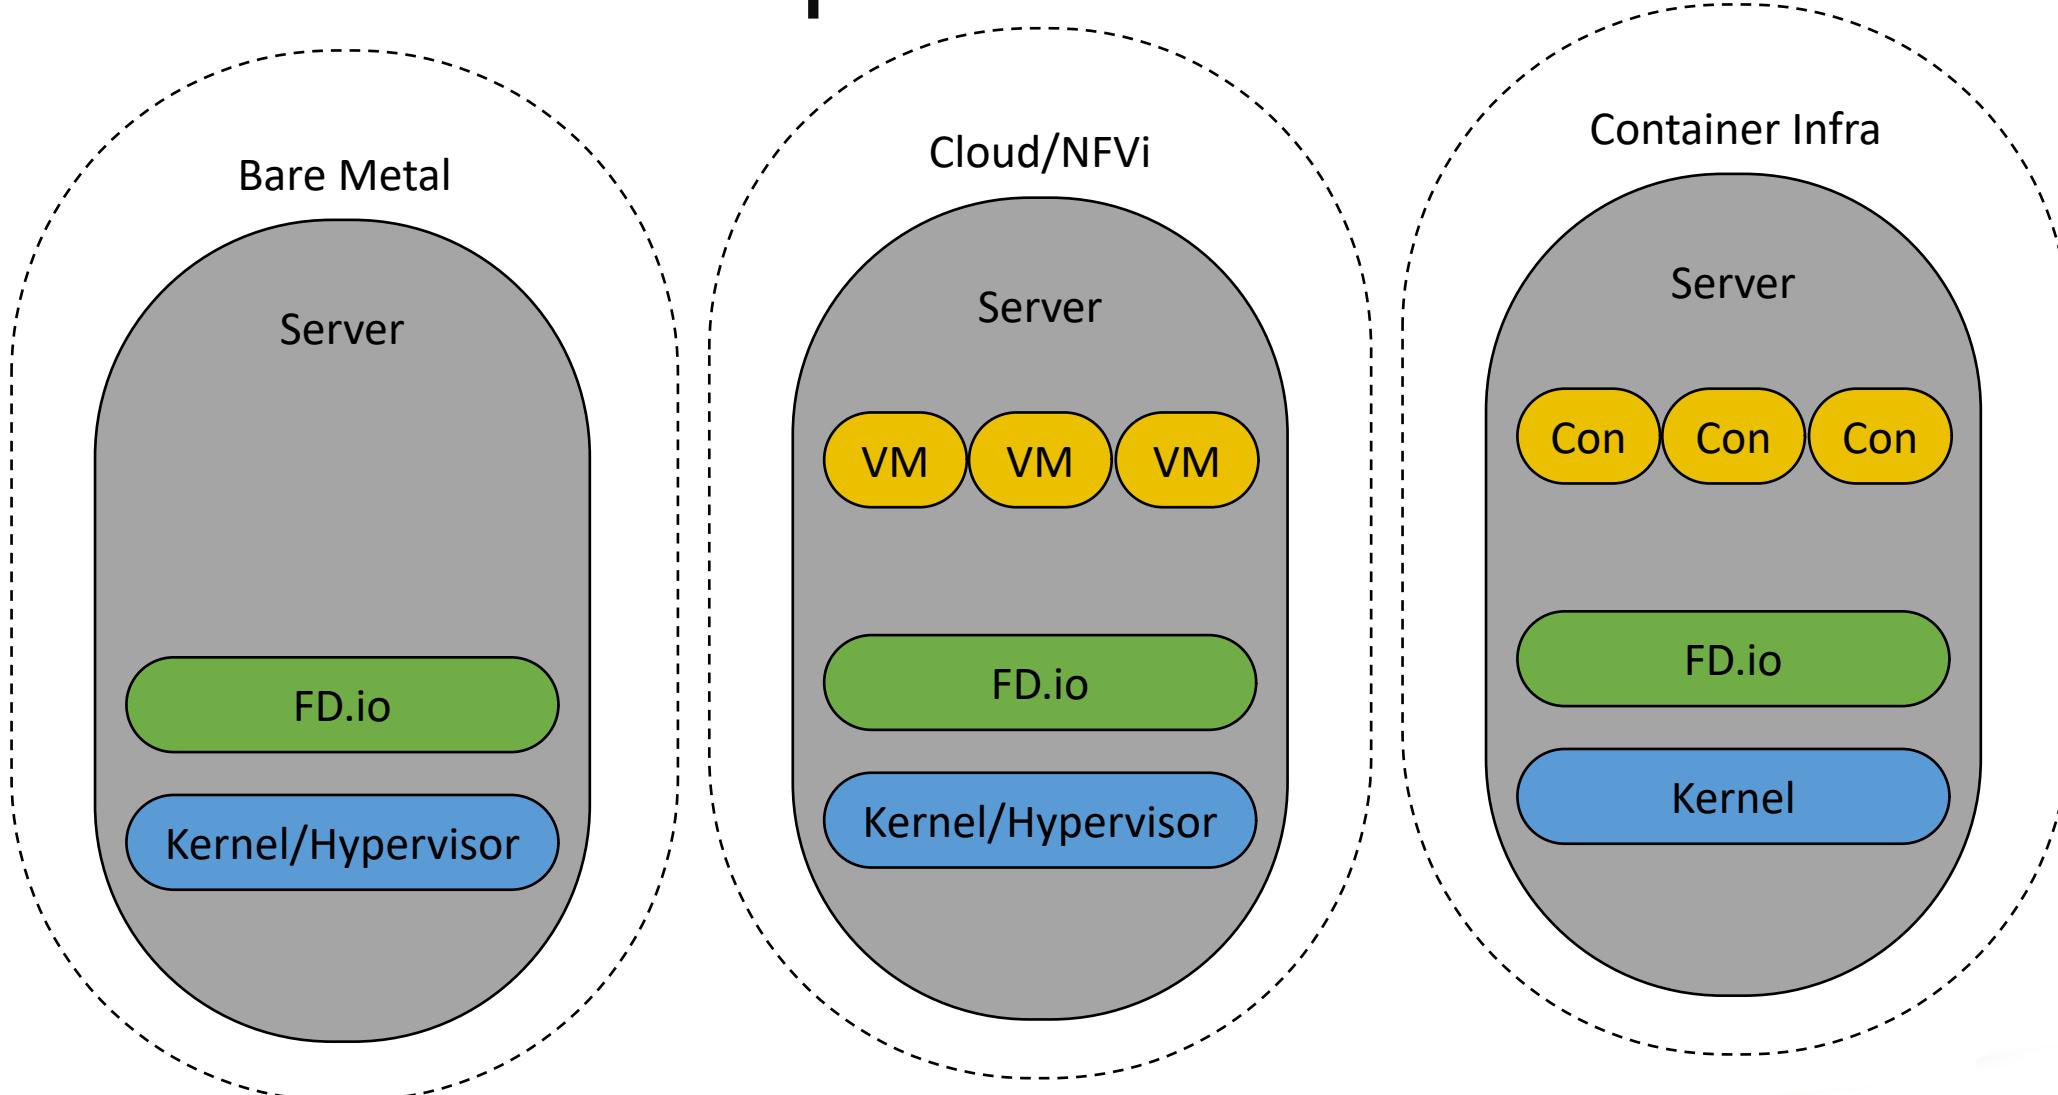


# Continuous Quality, Performance, Usability

Built into the development process – patch by patch

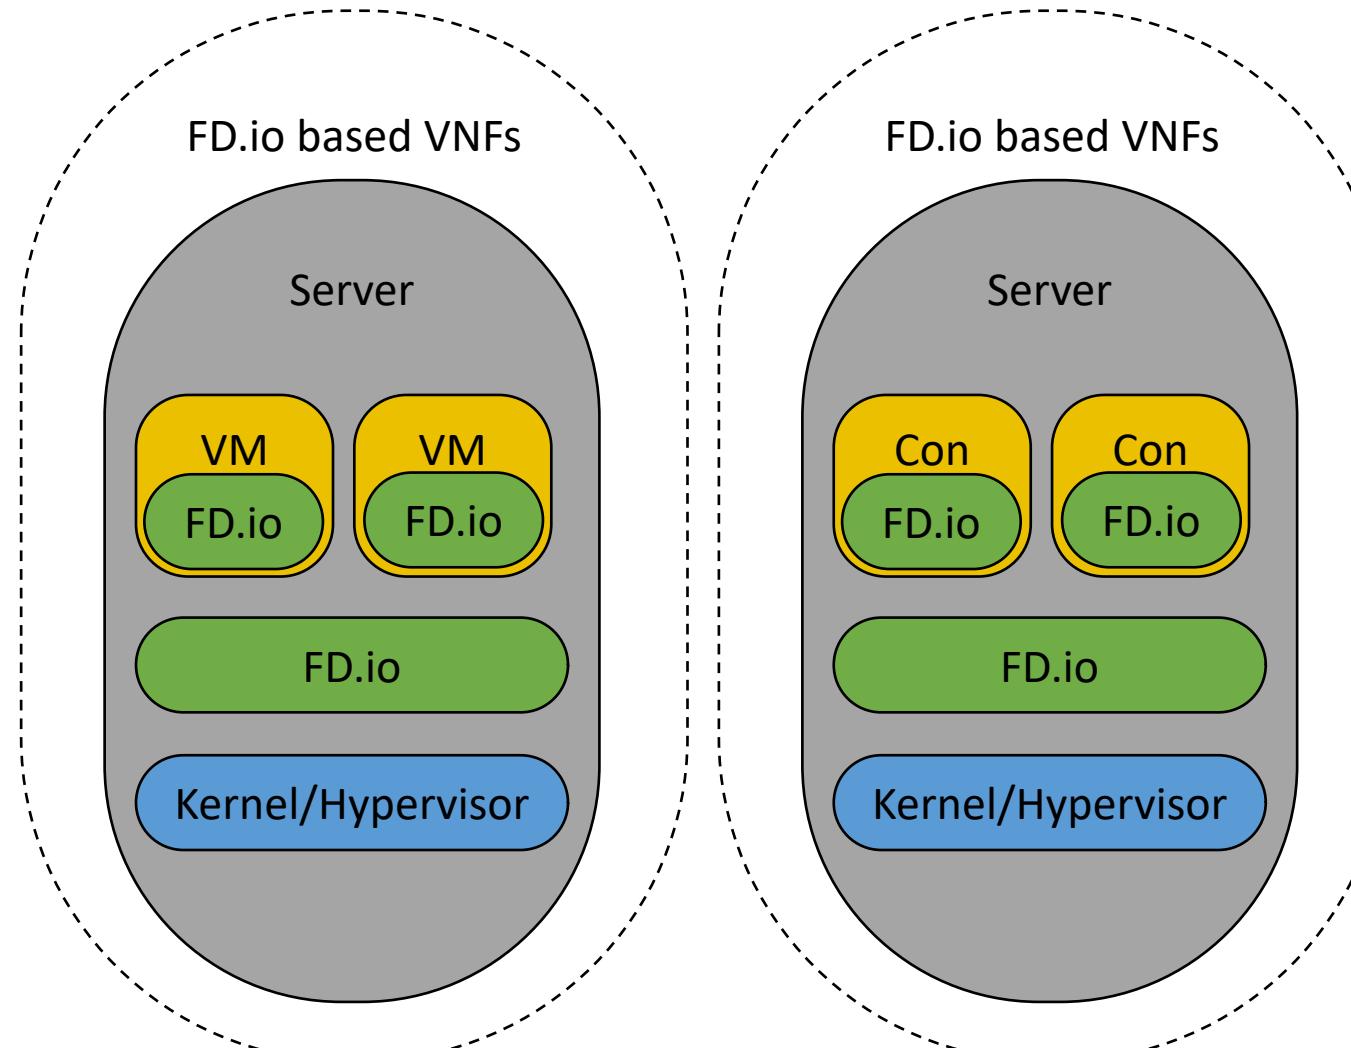


# Universal Dataplane: Infrastructure



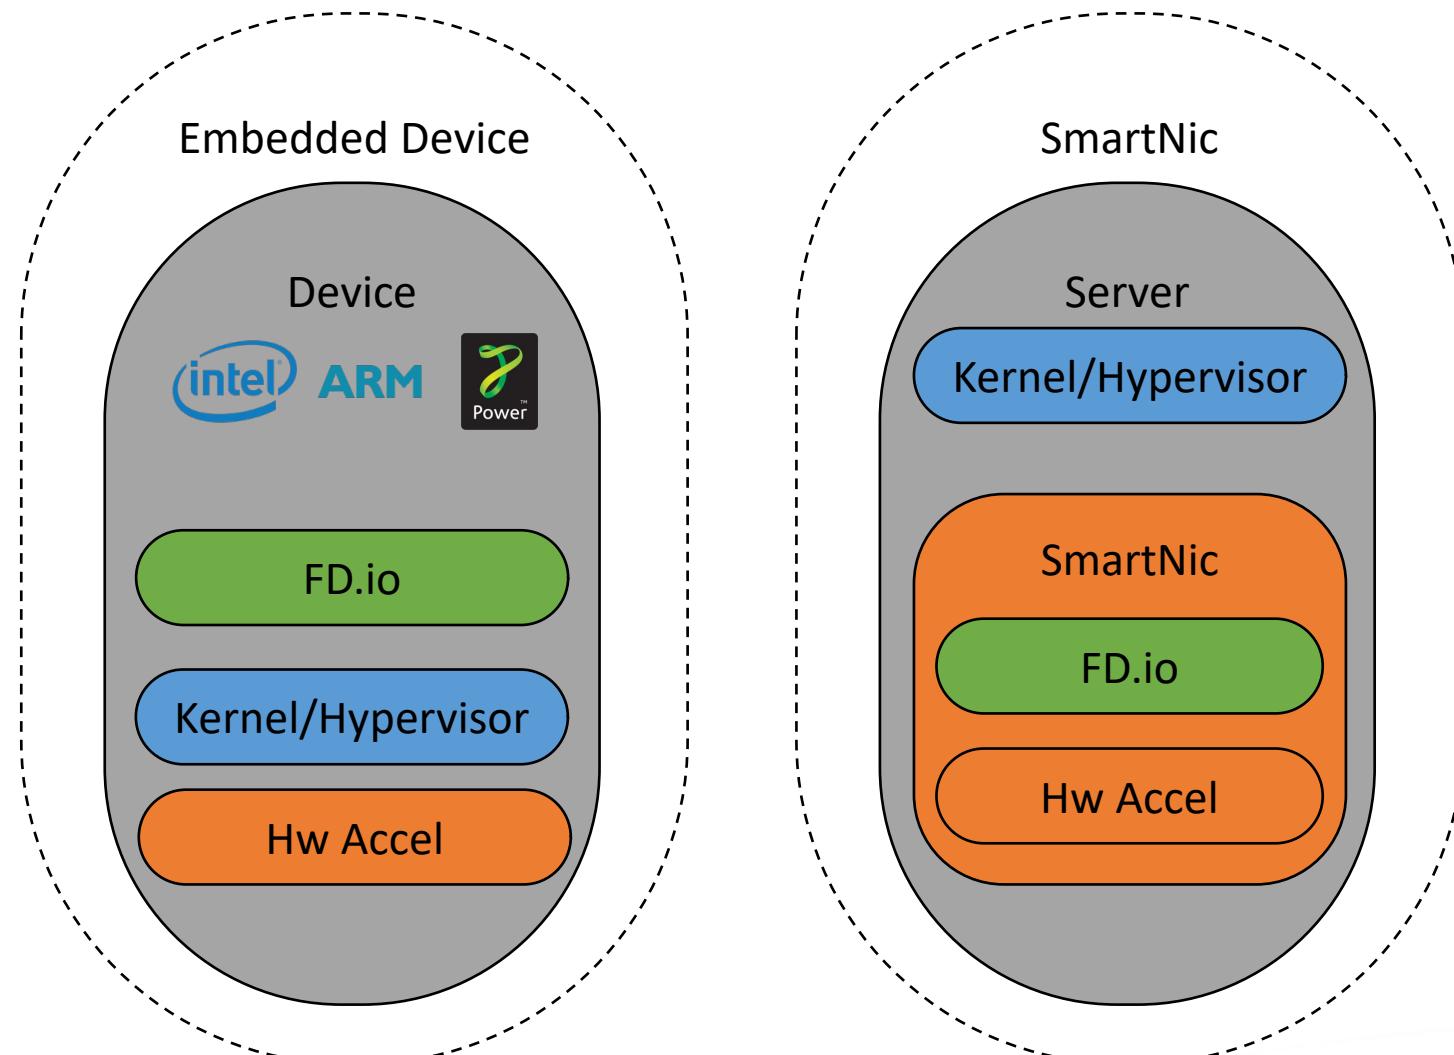


# Universal Dataplane: VNFs



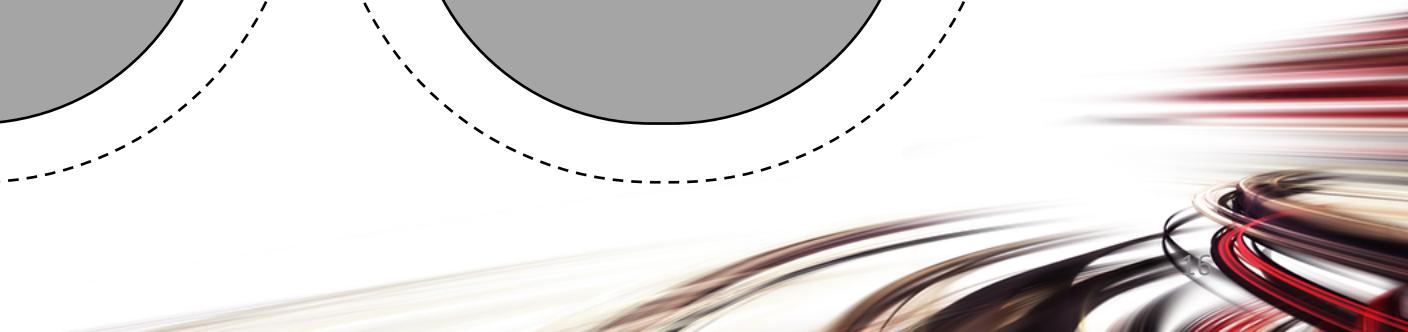
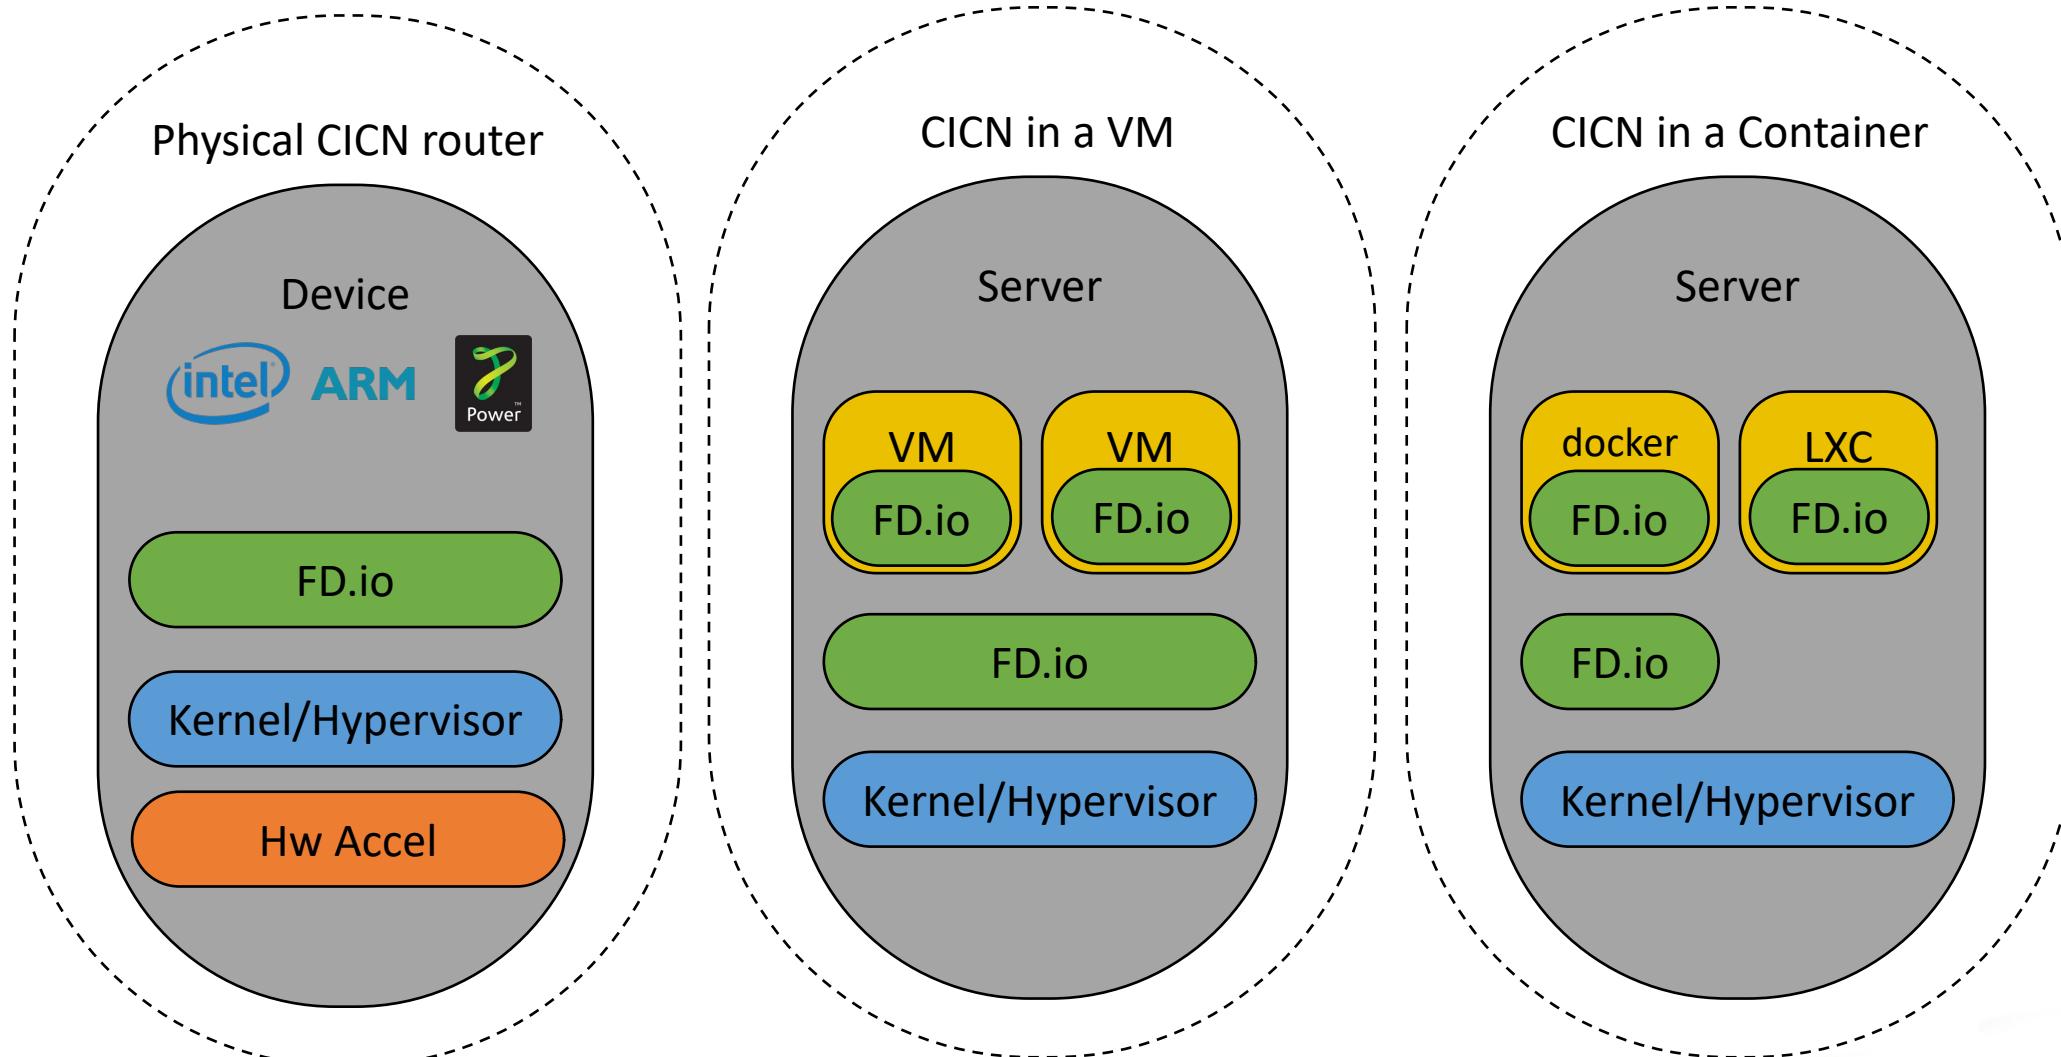


# Universal Dataplane: Embedded

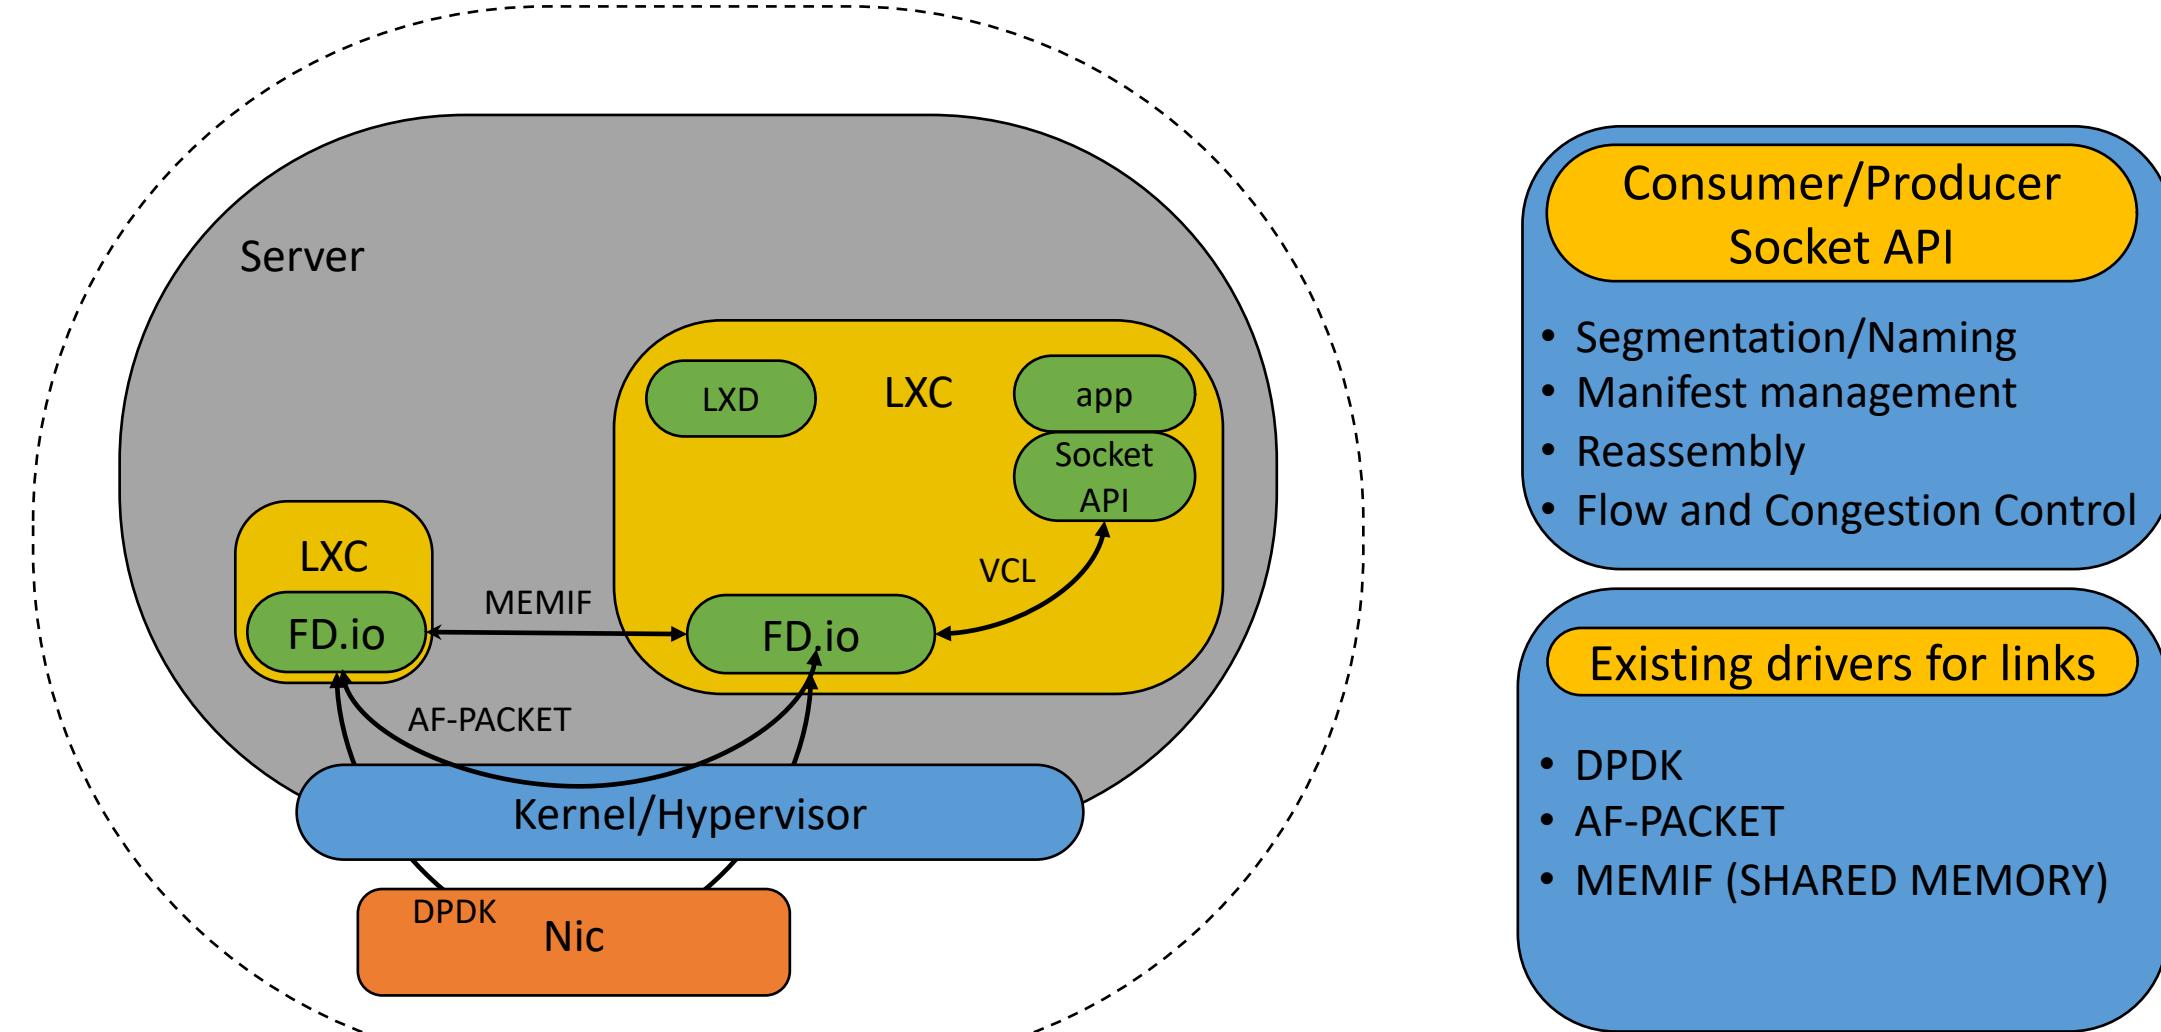




# Universal Dataplane: CICN Example



# Universal Dataplane: communication/API



## Consumer/Producer Socket API

- Segmentation/Naming
- Manifest management
- Reassembly
- Flow and Congestion Control

## Existing drivers for links

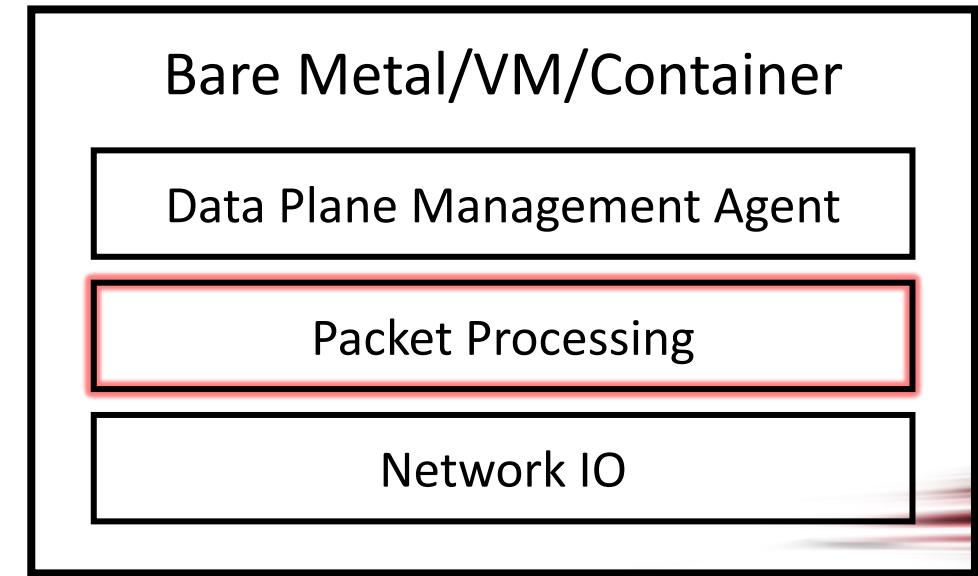
- DPDK
- AF-PACKET
- MEMIF (SHARED MEMORY)



# What is Vector Packet Processing?

An open-source software that provides out-of-the-box production quality switch/router functionality running under commodity CPUs

- High Throughput
  - 14+ Mpps per core
- Multiplatform
- Feature rich
  - L2, L3, L4, local and remote programmability
- Modular and Extensible
  - Through plugins





# Why VPP?

- NFV goals
  - Software flexibility without giving up to hardware level performance
- What about existing solutions?
  - Linux Kernel
    - Too slow for high throughput
    - Evolve slowly
  - Click
    - In principle similar to VPP, no V(ector)



# CICN distribution

- Core libraries
  - Consumer/Producer Socket API, CCNx libs, PARC C libraries
- Server and Router
  - VPP cicn plugin for Ubuntu 16, CentOS 7
  - HTTP video server, Apache Traffic Server Plugin coming soon
- Client
  - Metis Forwarder
  - VIPER MPEG-DASH video player
  - Android 7/8, MacOS X 10.12, iOS 10/11, Ubuntu 16, CentOS 7
  - Soon Apple Store and Google Play
- vICN
  - intent-based networking
  - model driven programmable framework
  - monitoring and streaming for BigData support (PNDA.io)



# Opportunities to Contribute



- Forwarding strategies
- Mobility management
- Hardware Accelerators
- vICN, configuration/management/control
- Consumer/Producer Socket API
- Reliable Transport
- Instrumentation tools
- HTTP integration

We invite you to Participate in [fd.io](#)

- [Get the Code, Build the Code, Run the Code, install from binaries](#)
- [from binary packages](#)
- [Read/Watch the Tutorials](#)
- [Join the Mailing Lists](#)
- [Join the IRC Channels](#)
- [Explore the wiki](#)
- [Join fd.io as a member](#)
- <https://wiki.fd.io/view/cicn>
- <https://wiki.fd.io/view/vicn>
- <https://fd.io/>



# Vector Packet Processing for ICN

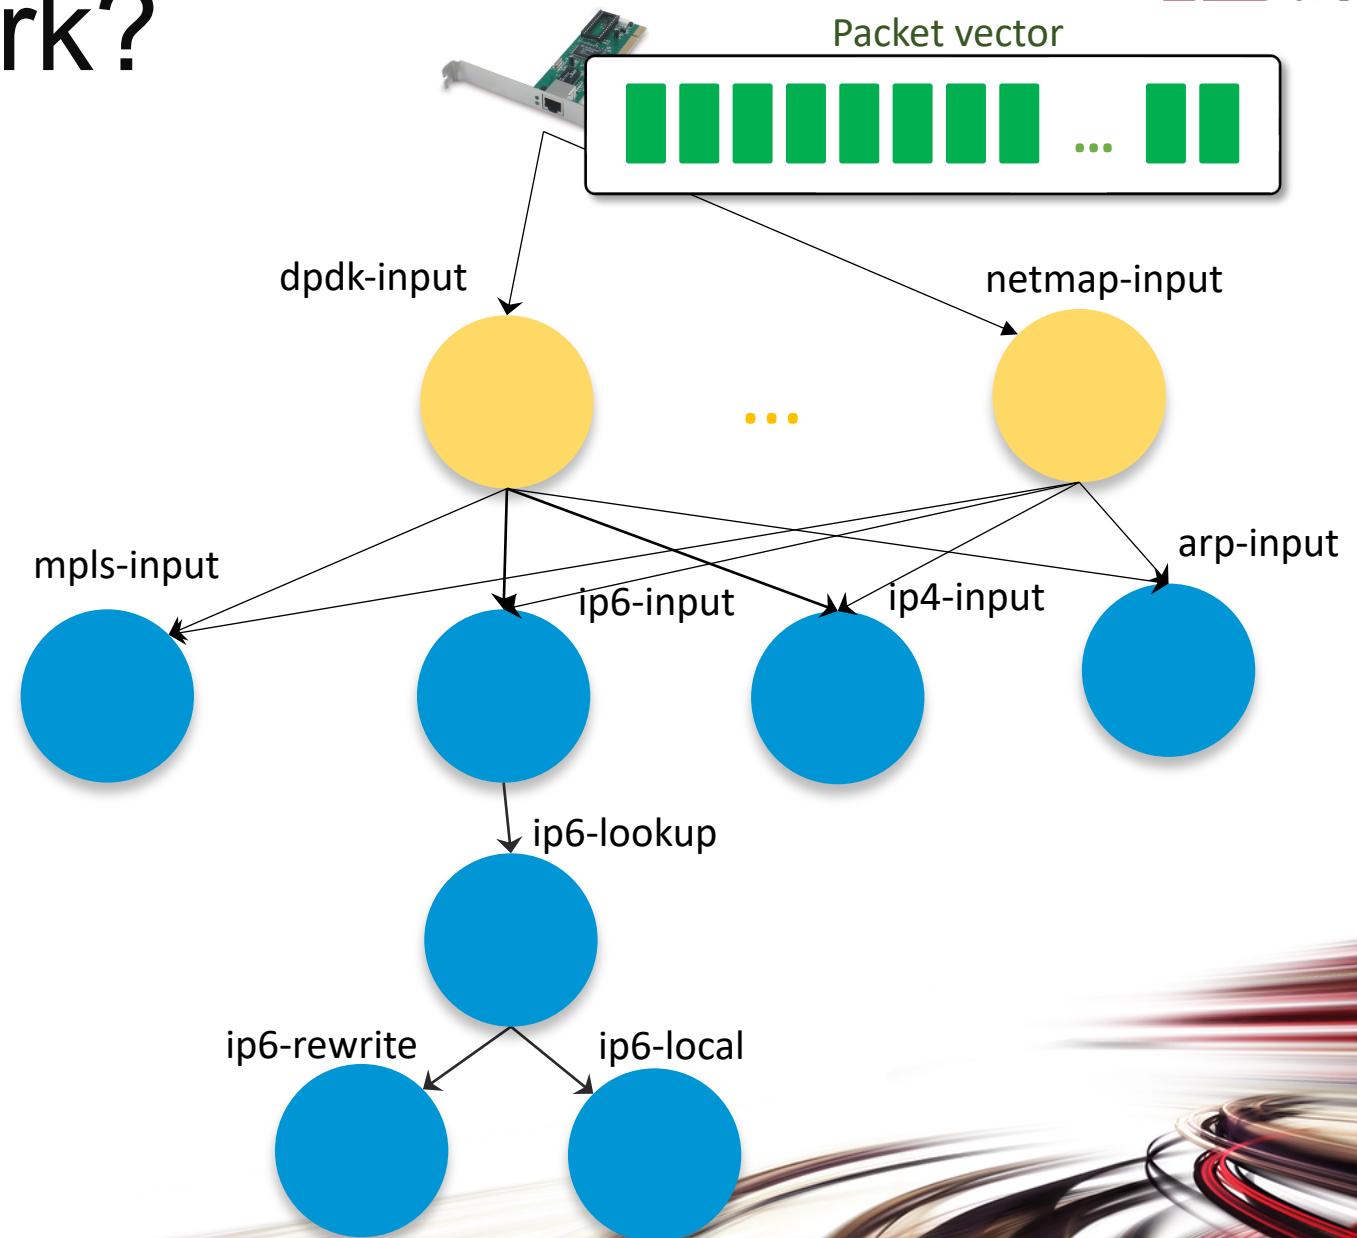
Alberto Compagno



Tutorial at ACM SIGCOMM ICN, Berlin, Germany  
26<sup>th</sup> of September 2017

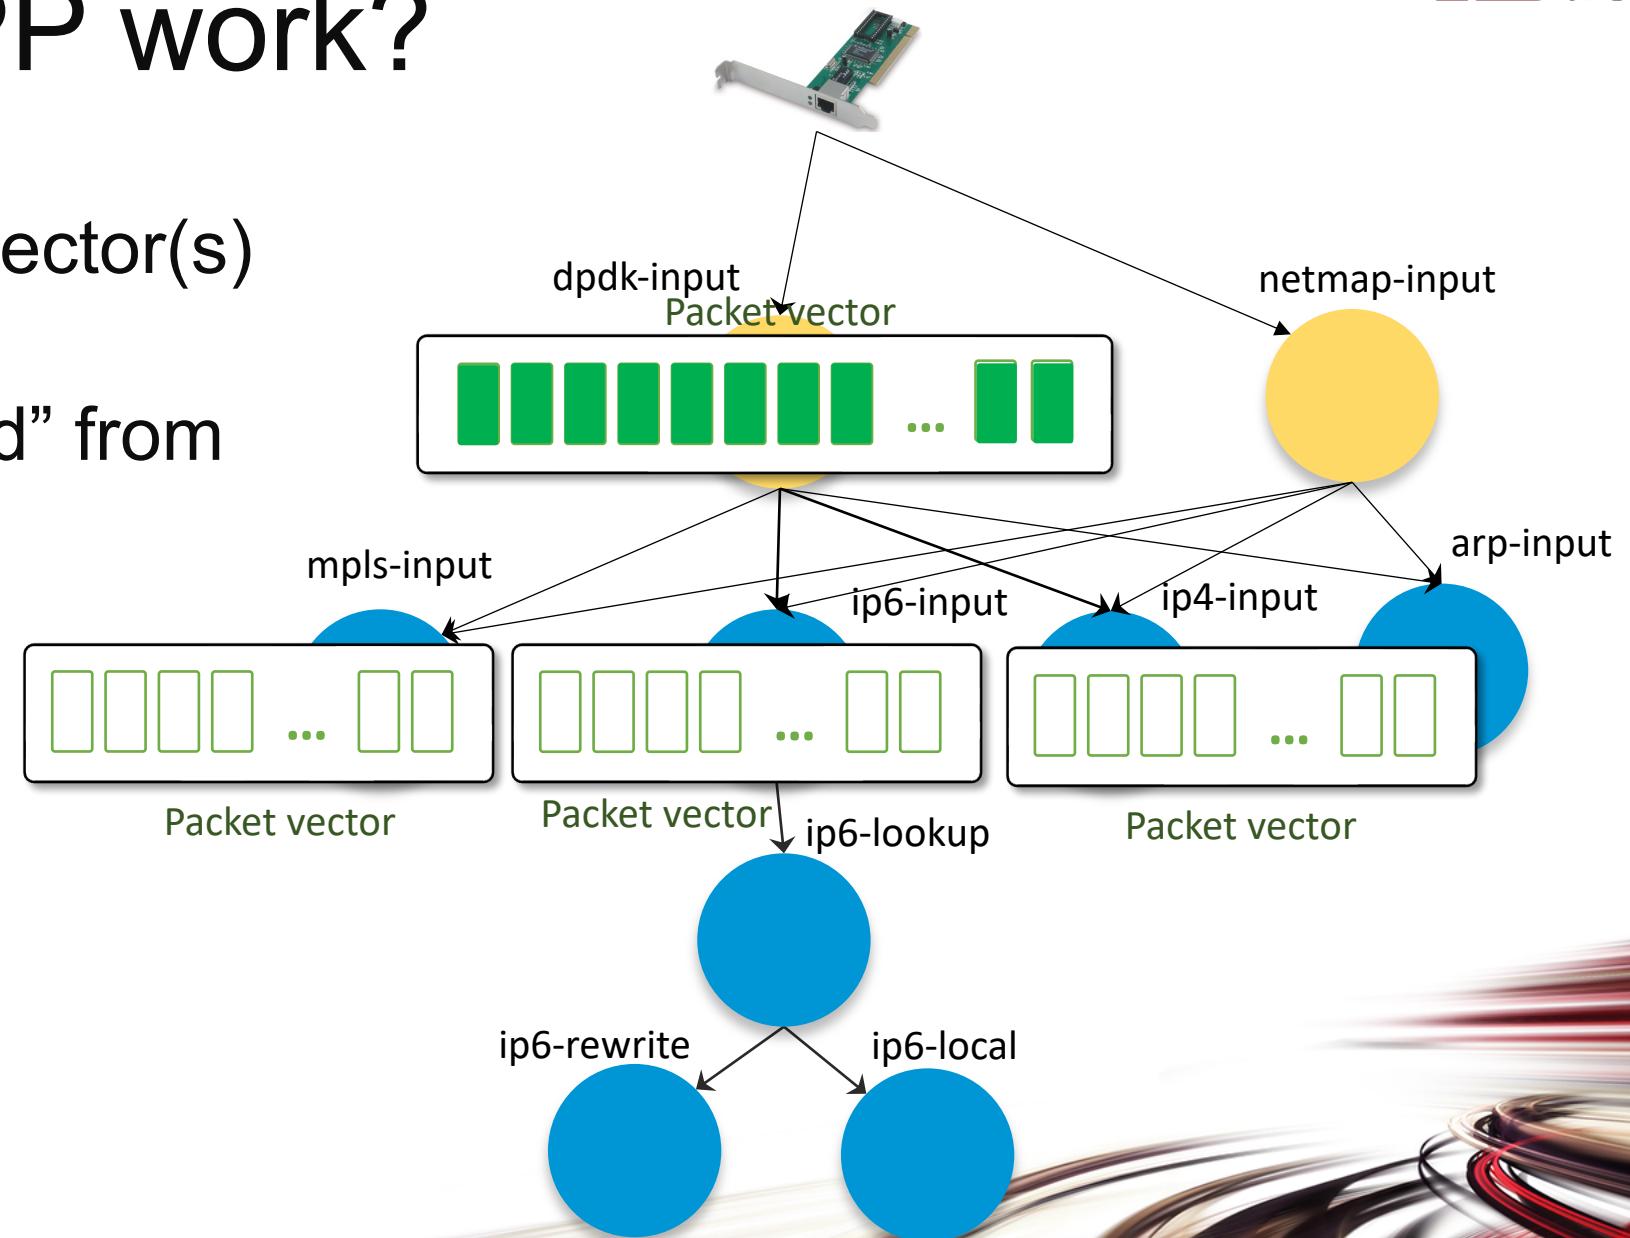
# How does VPP work?

- VPP is a ‘packet processing graph’
- Nodes are
  - Small
  - Loosely coupled
- VPP processes vectors of packets
  - Passed from node to node



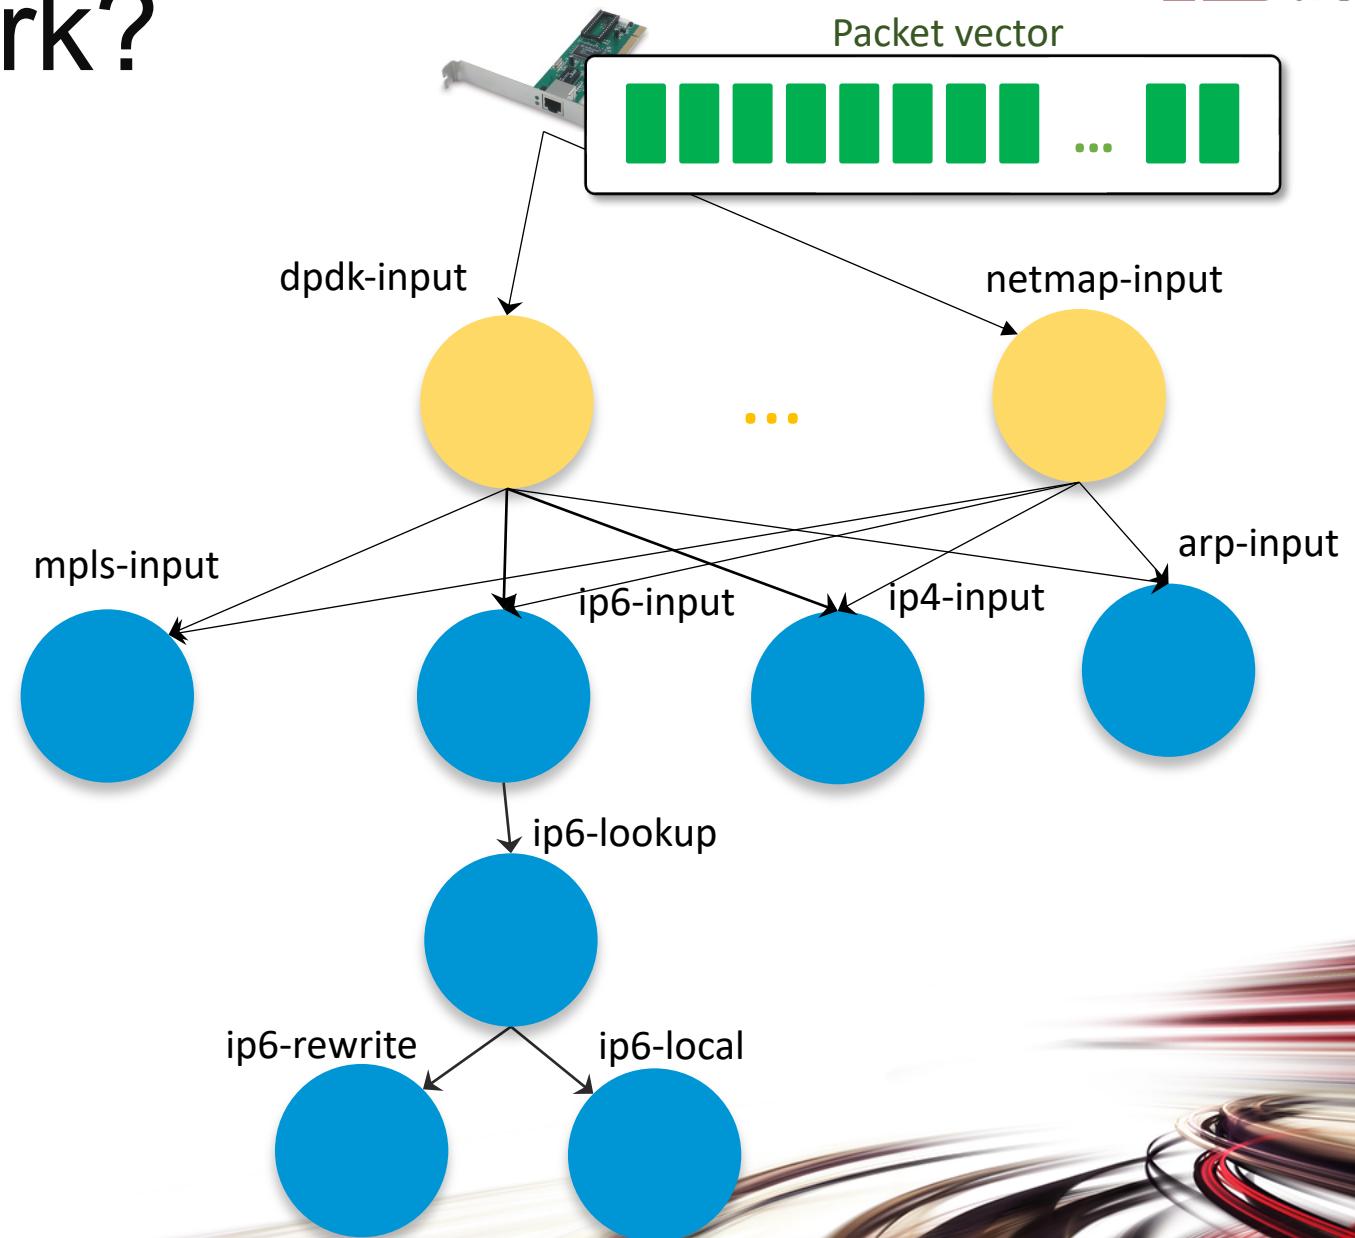
# How does VPP work?

- Each node has its vector(s)
- Packets are “passed” from vector to vector



# How does VPP work?

- Three types of nodes
  - Input
  - Internal
  - Process



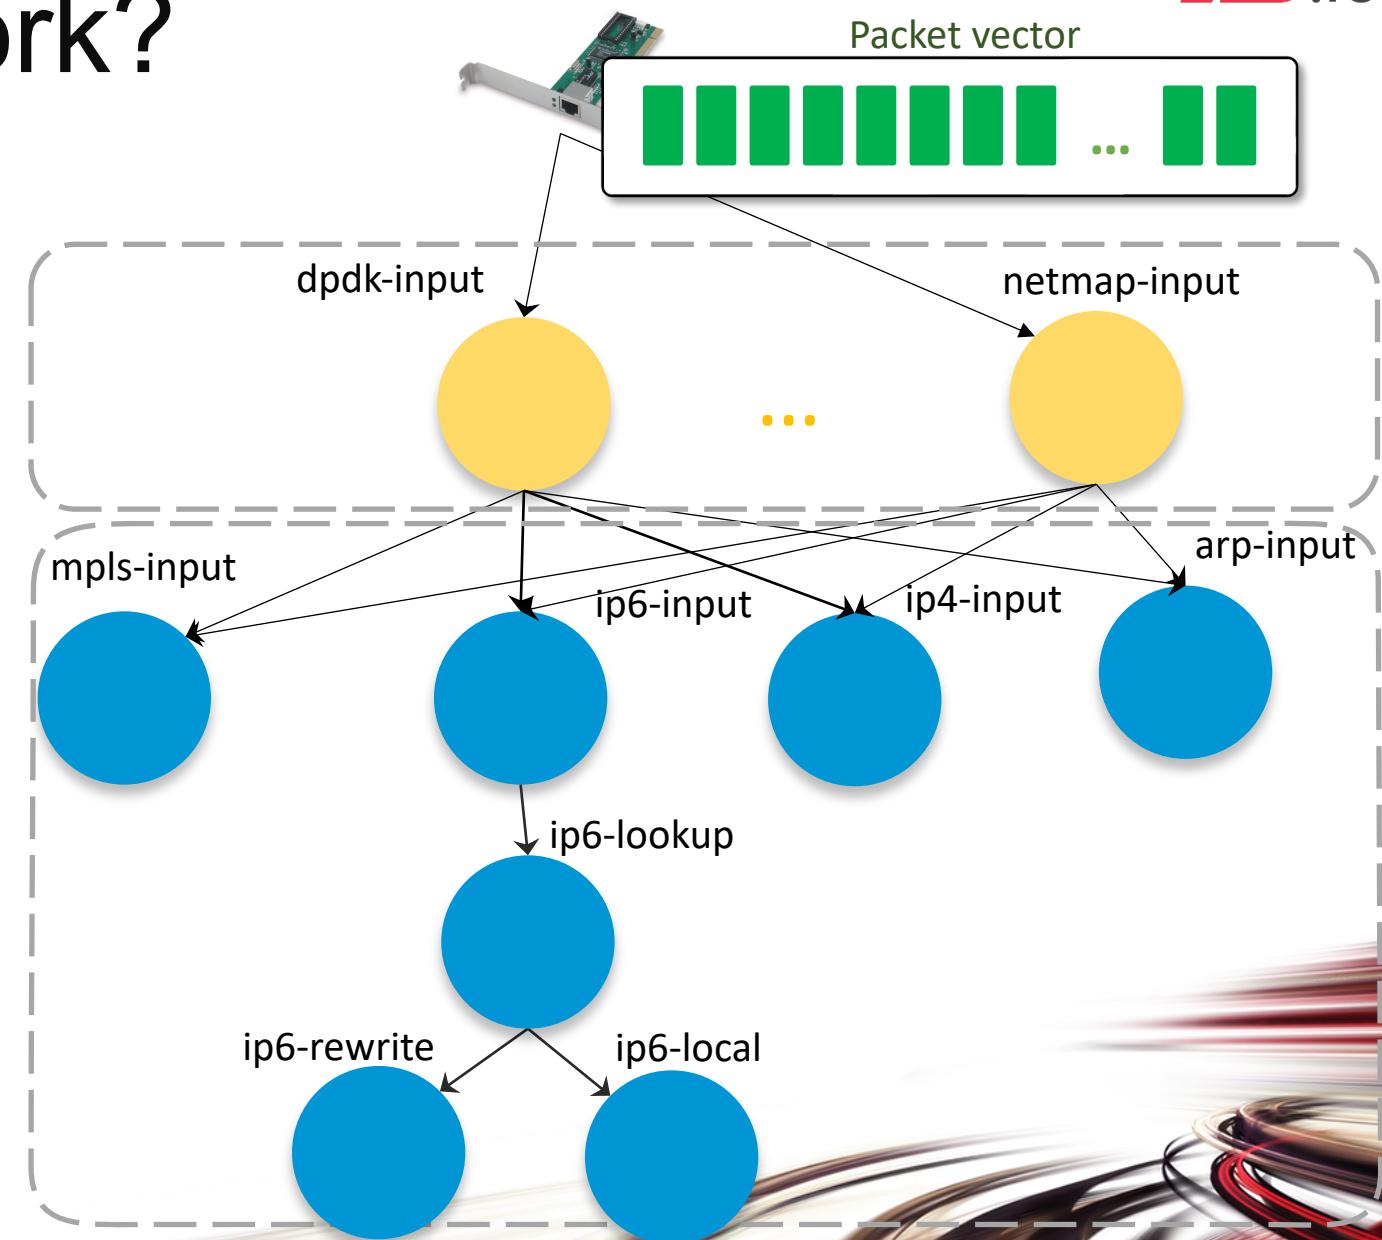
# How does VPP work?

## Input nodes

- Read packets from RX buffer
- Create the packet vector

## Internal nodes

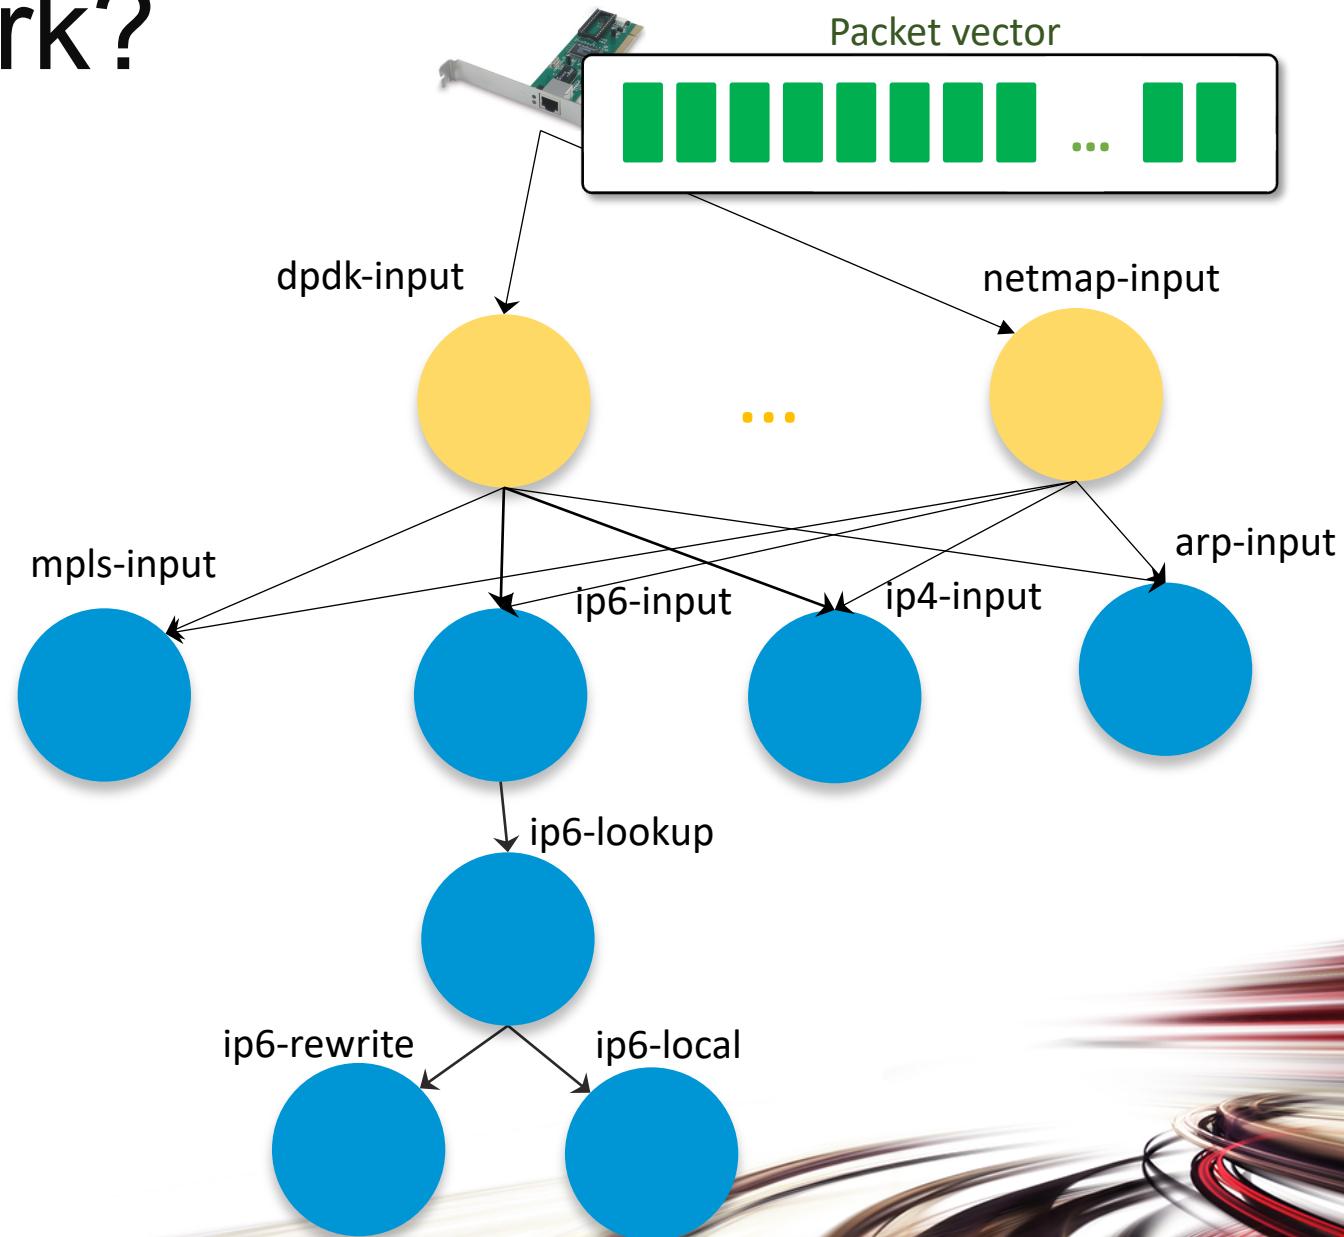
- Process packets
- Called from other nodes
- Can be leaf (drop or TX)



# How does VPP work?

## Process nodes

- Not part of the processing graph
- Run in background
- React to timer/event

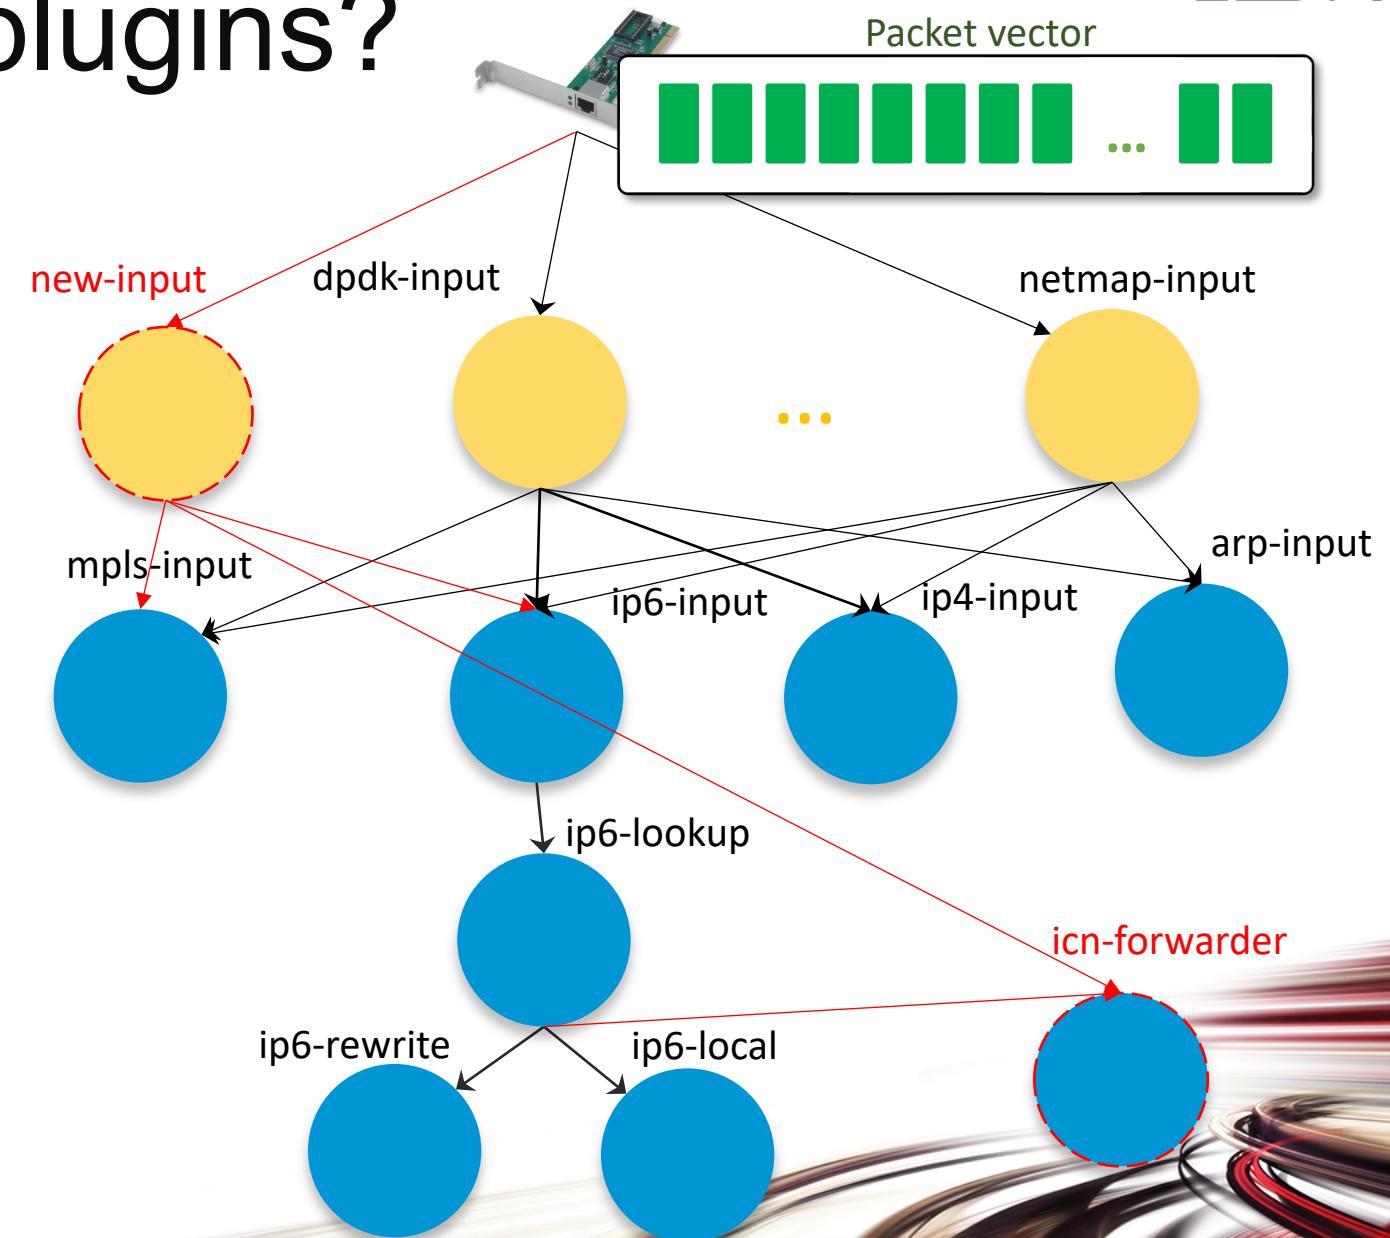


# Extend VPP with plugins?

Plugins are first class citizen

They can:

- Add nodes
- Add api
- Rearrange the graph



# How does VPP accelerate packet processing?





# Accelerating packet processing

- Kernel bypass
- Code Design (Multi-loop, Branch prediction, Function flattening, Lock-free structures, Numa aware)
- Reduce cache misses

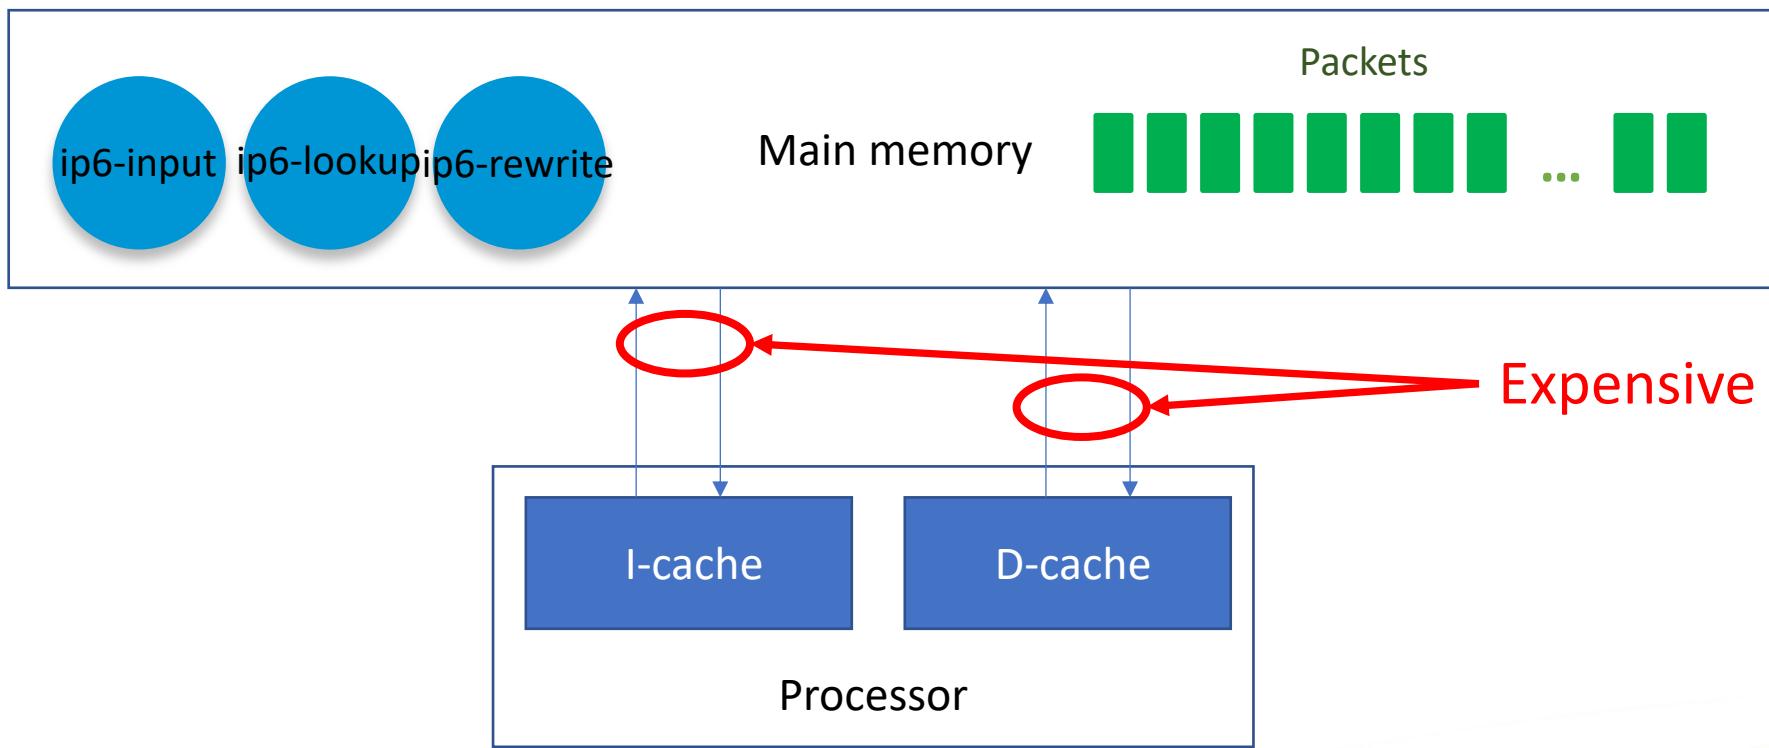


# Reduce cache misses – Why?

- 14 Mpps on 3.5GHz CPU = 250 cycles/packet
- Cache hit:
  - ~2-30 cycles
- Cache miss (main memory)
  - ~140 cycles

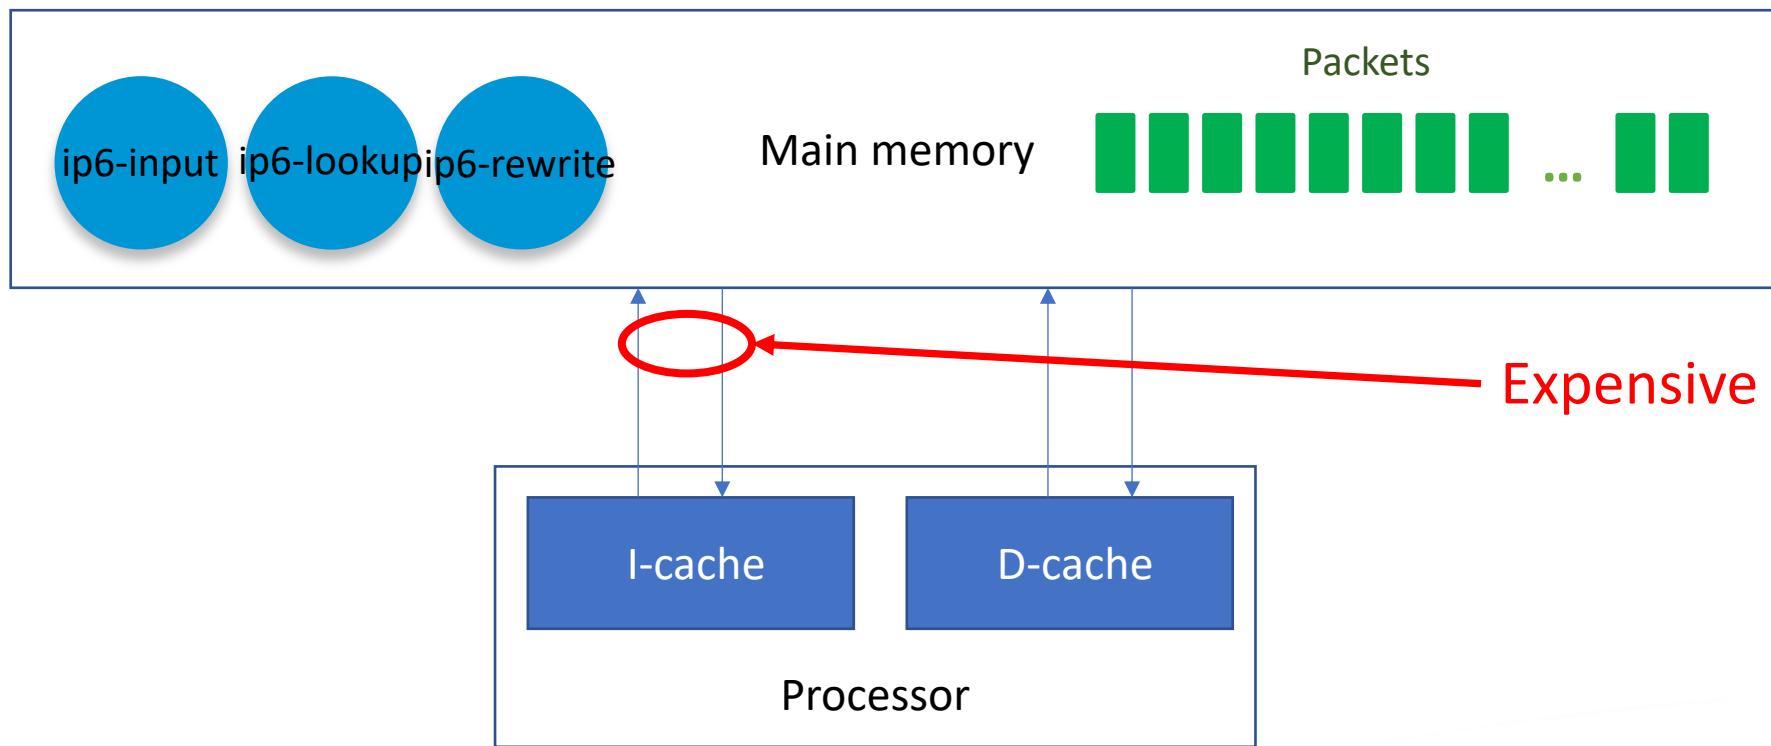


# Reduce cache misses



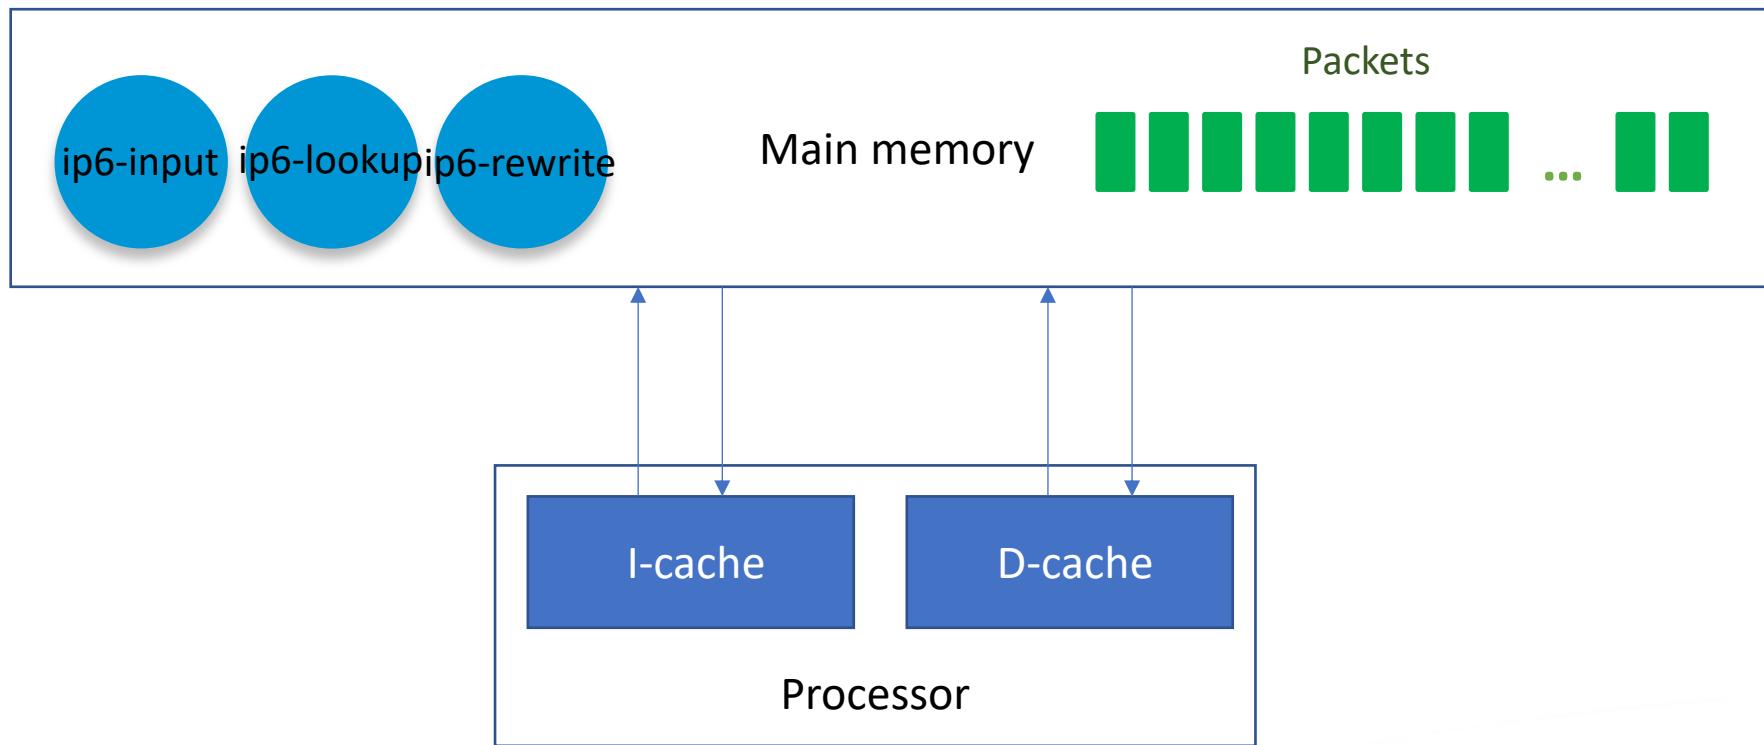
# Reduce cache misses – I-cache

Let's compare scalar processing with vector processing



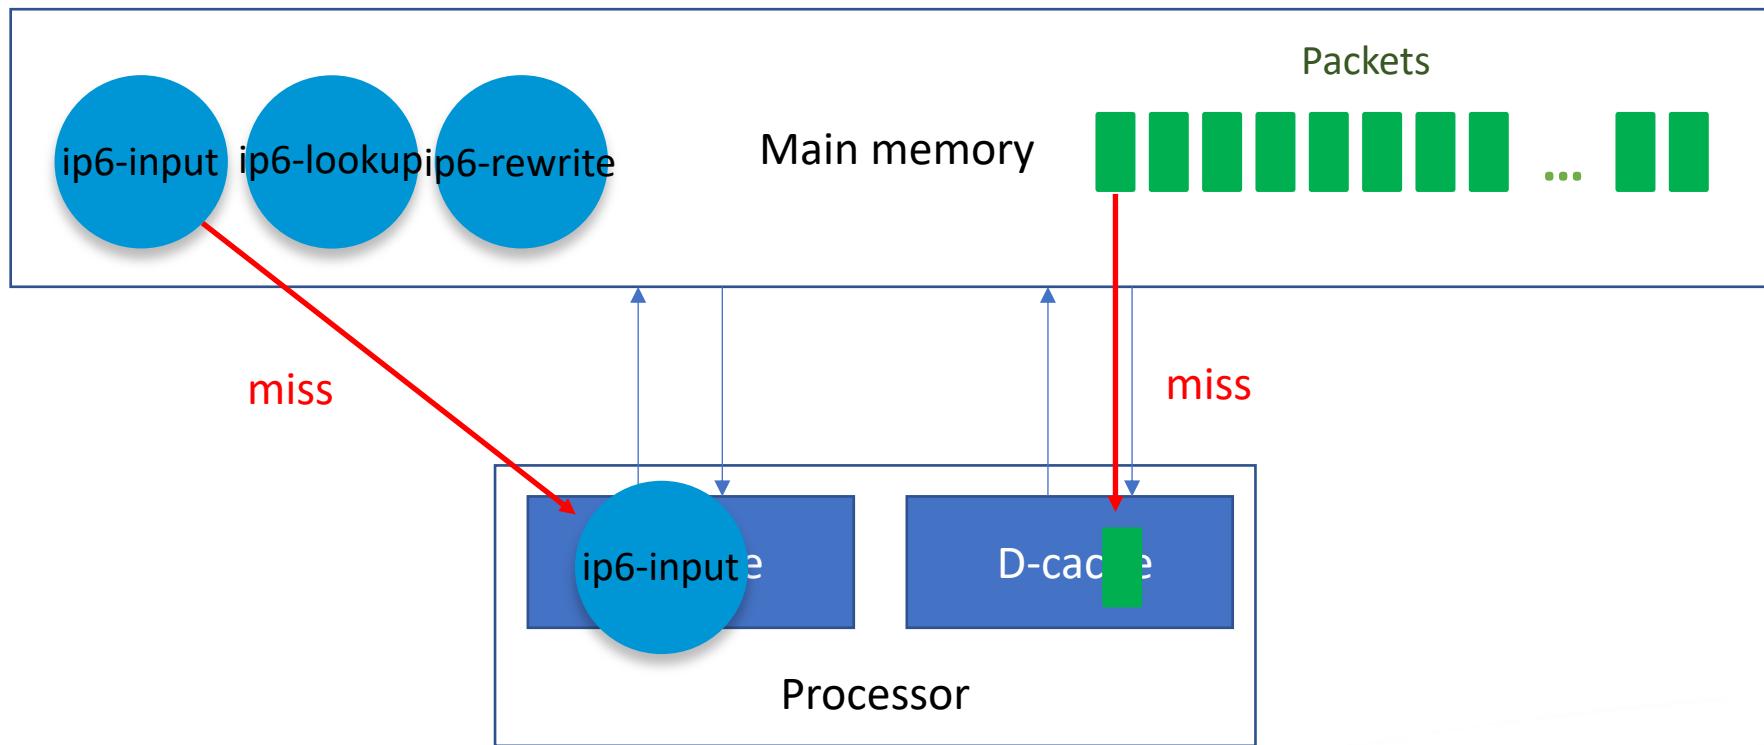
# Scalar Packet Processing

- Process one packet at a time



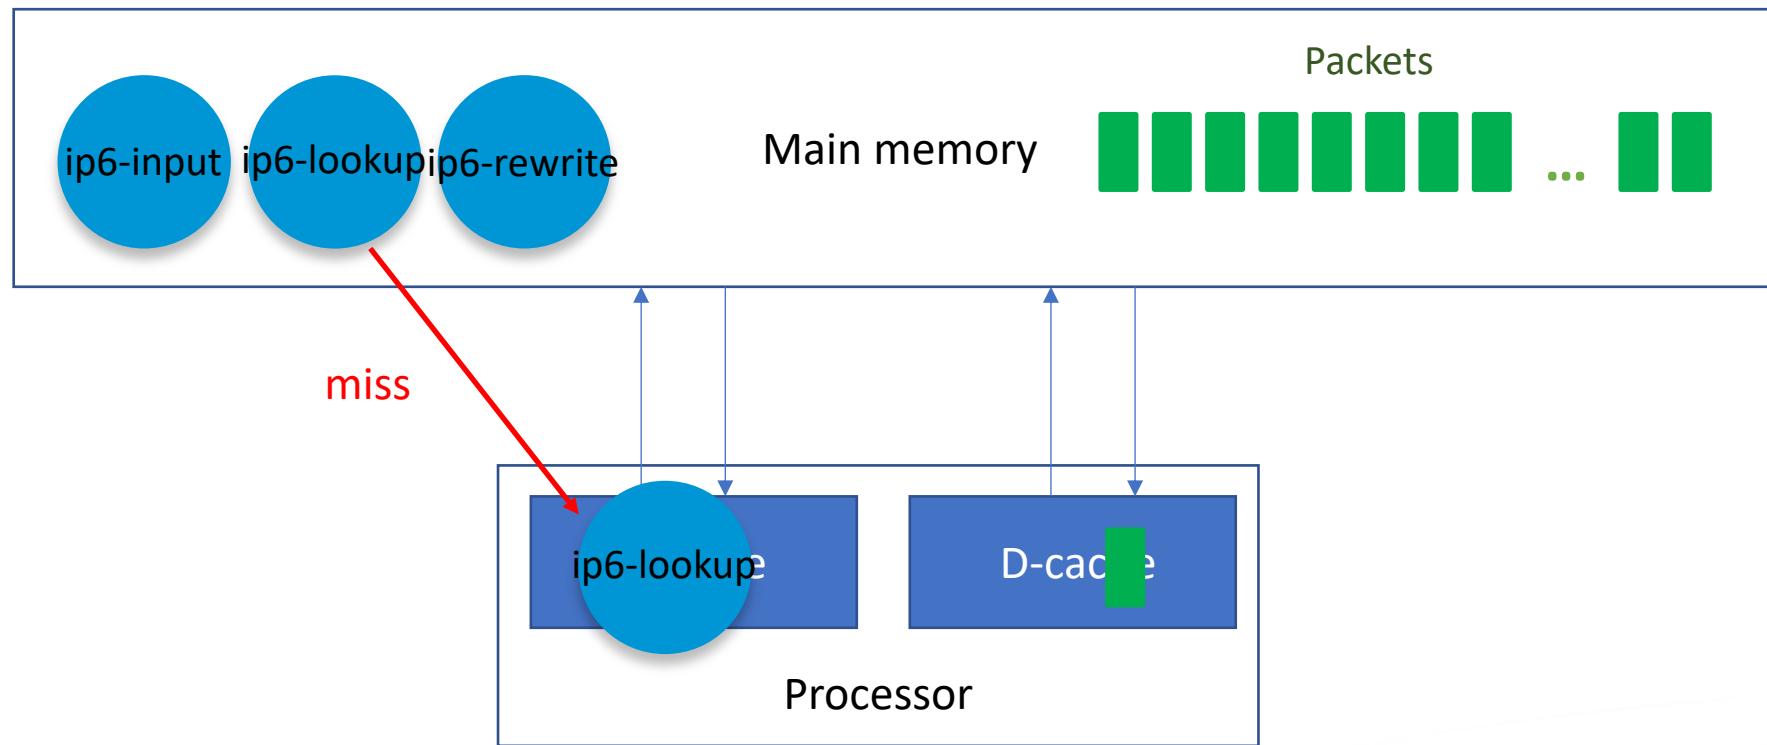
# Scalar Packet Processing

- Process one packet at a time



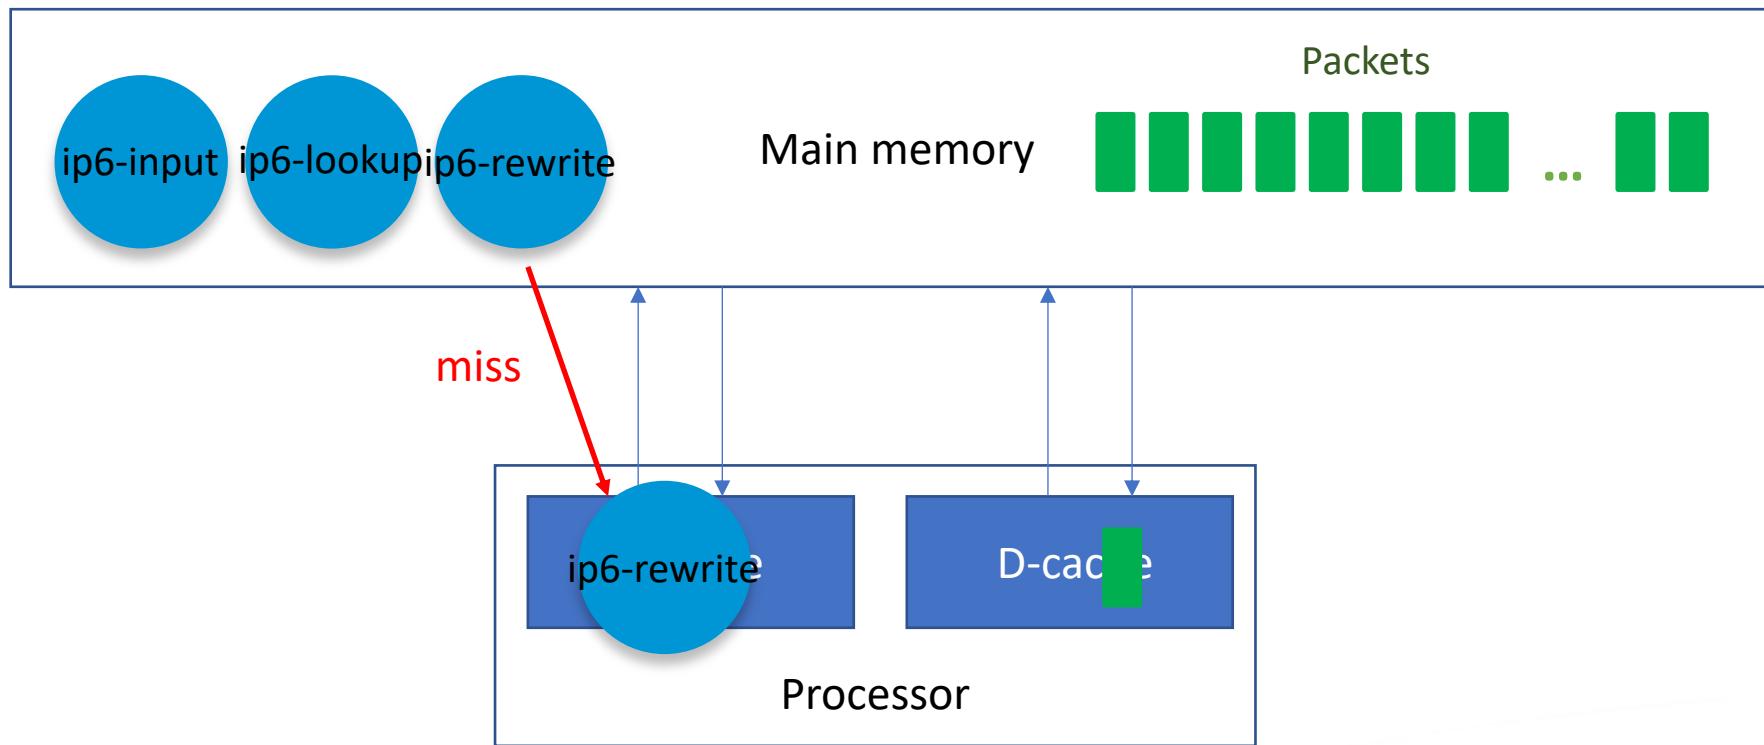
# Scalar Packet Processing

- Process one packet at a time



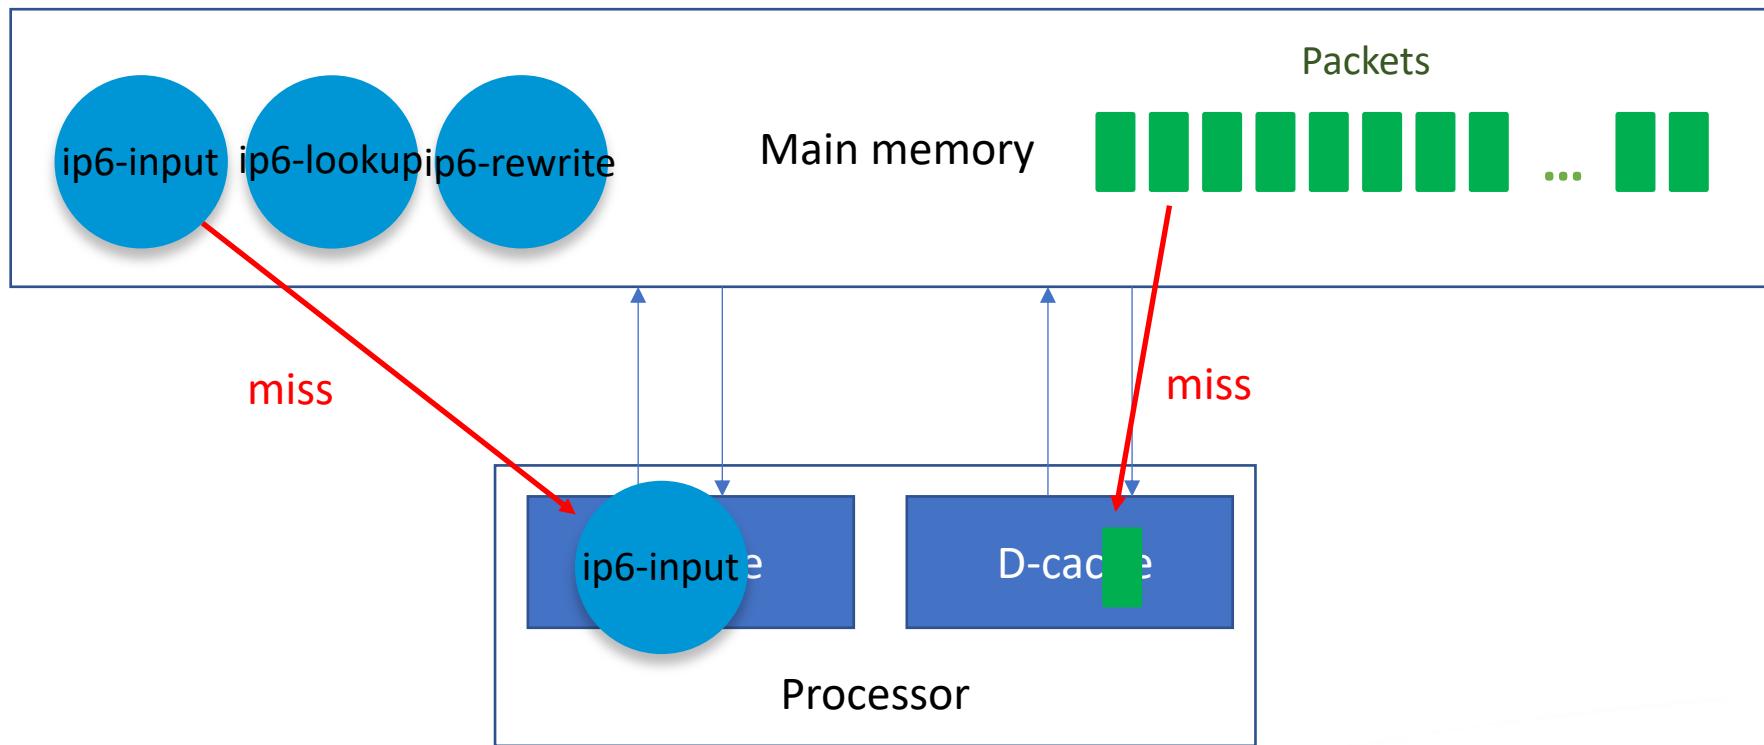
# Scalar Packet Processing

- Process one packet at a time



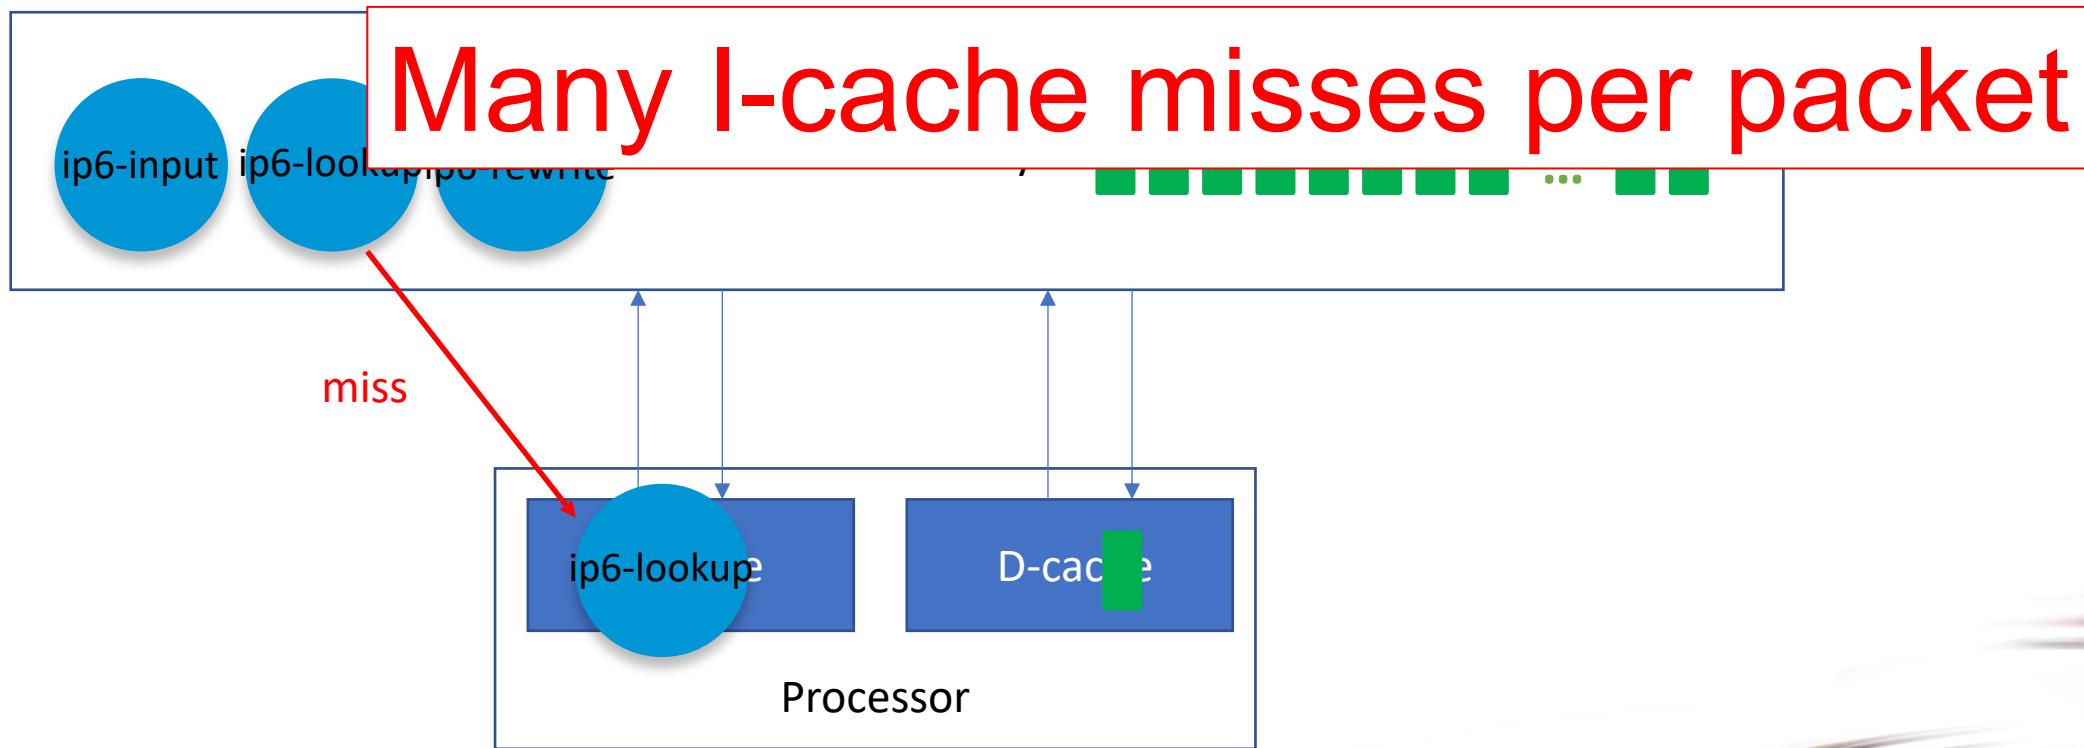
# Scalar Packet Processing

- Process one packet at a time



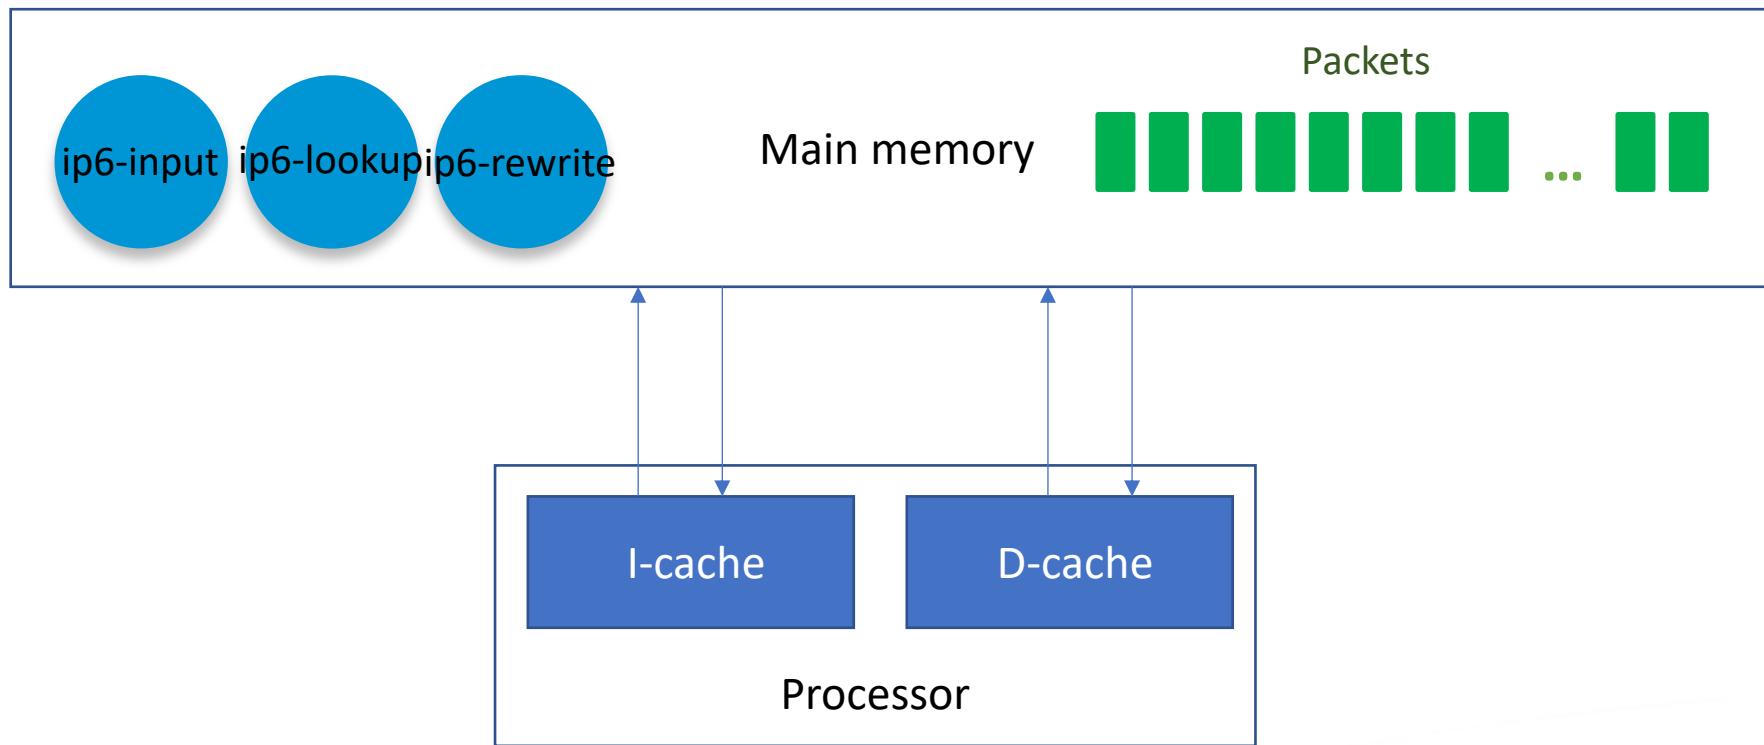
# Scalar Packet Processing

- Process one packet at a time



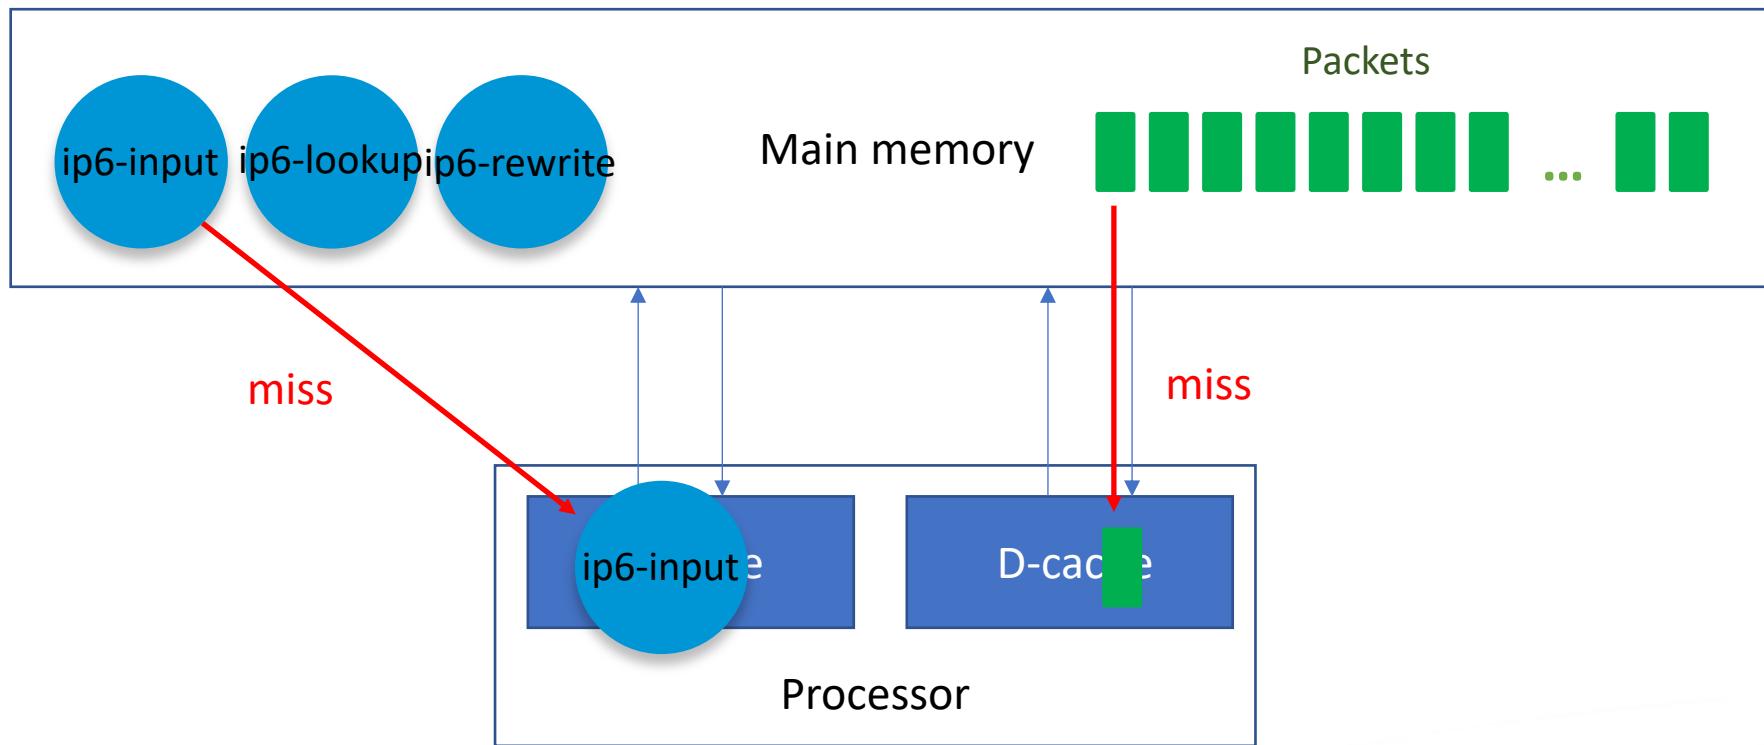
# Vector Packet Processing

- Every node process the full packet vector



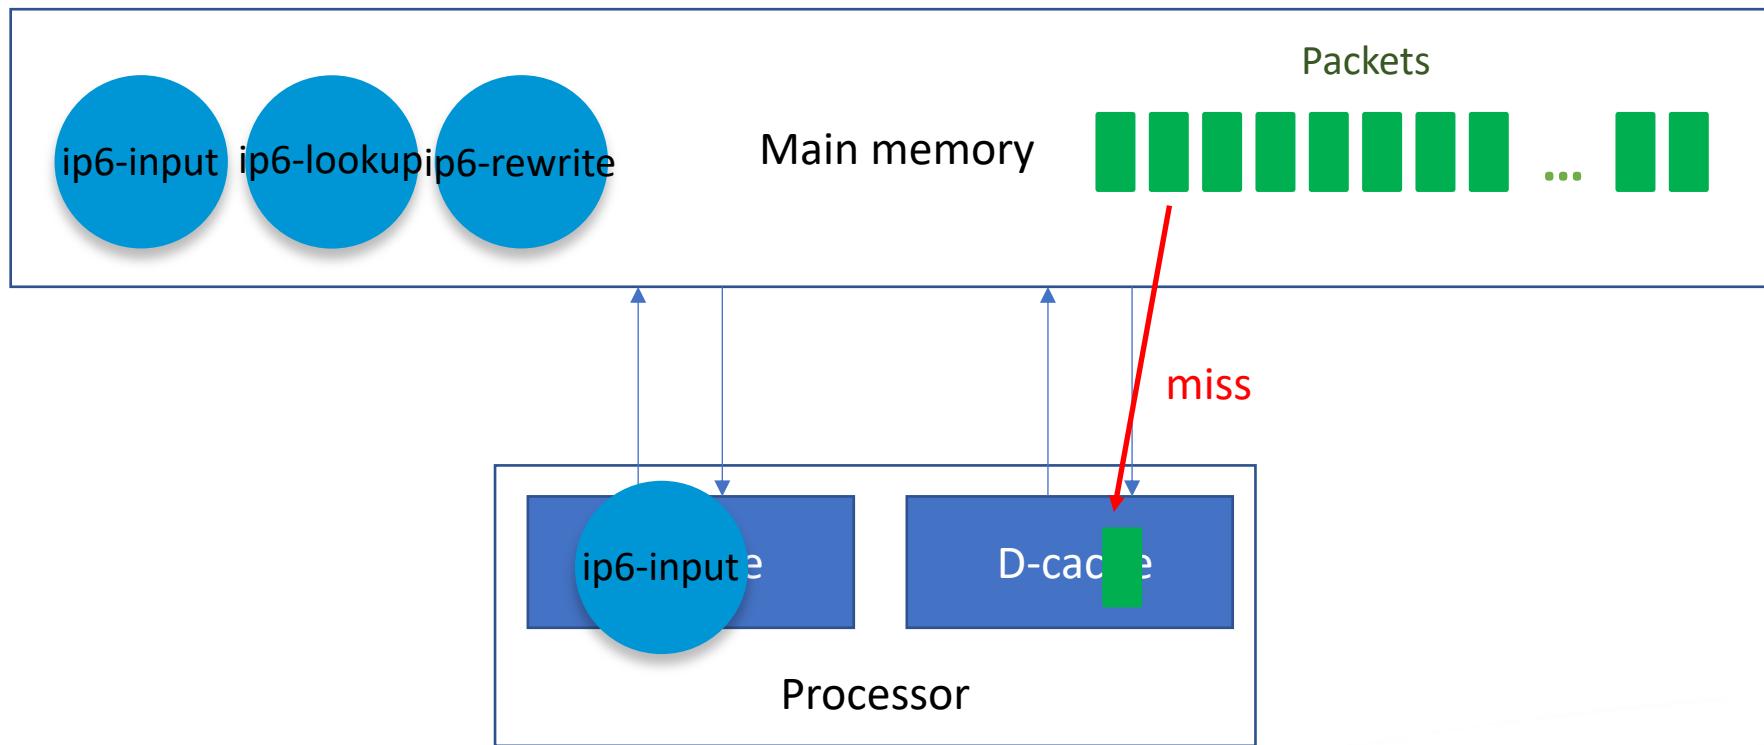
# Vector Packet Processing

- Every node process the full packet vector



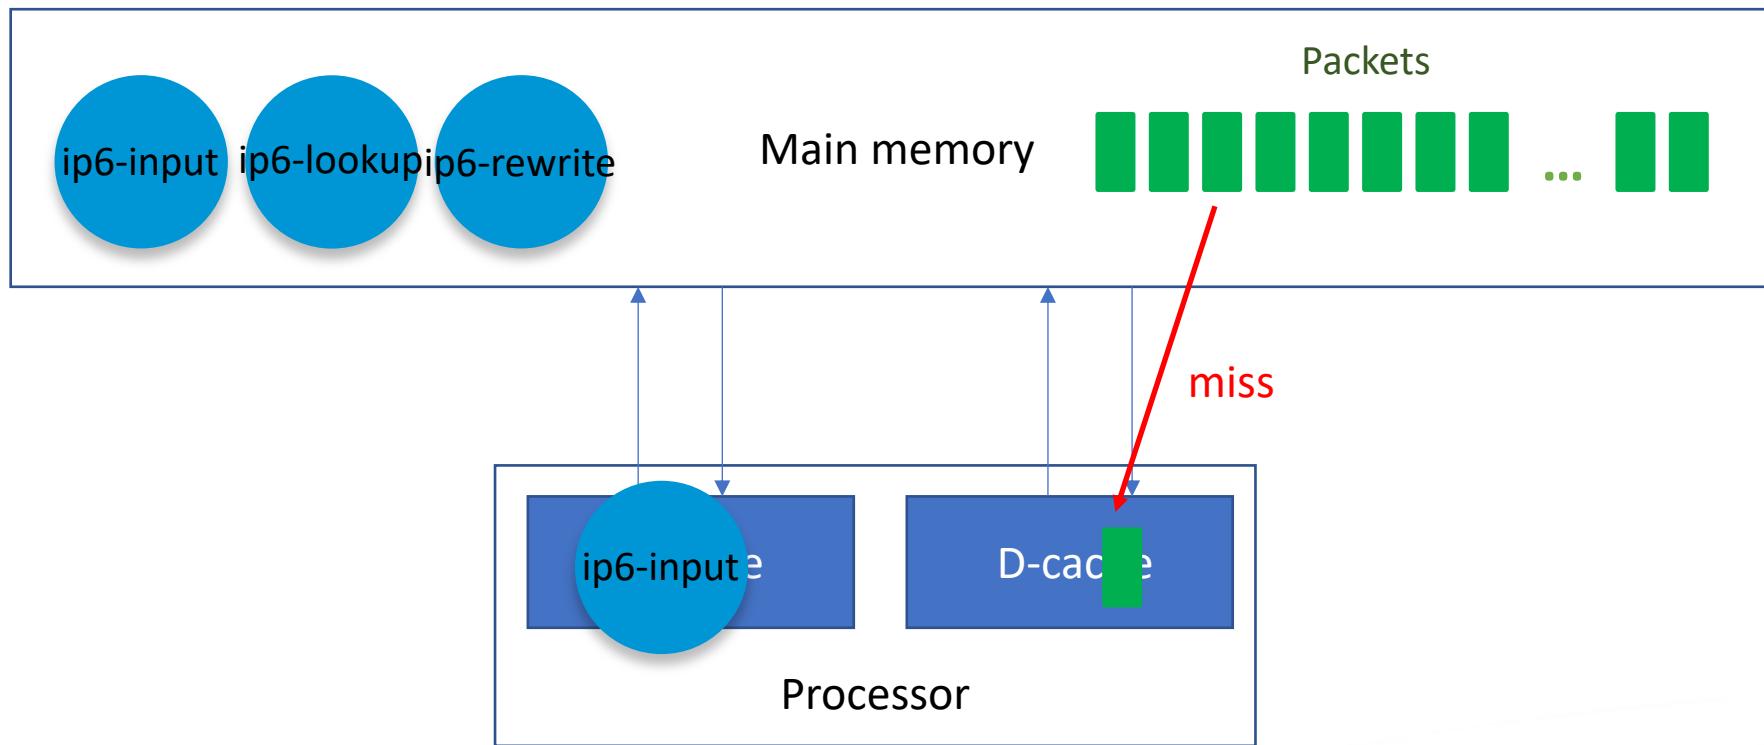
# Vector Packet Processing

- Every node process the full packet vector



# Vector Packet Processing

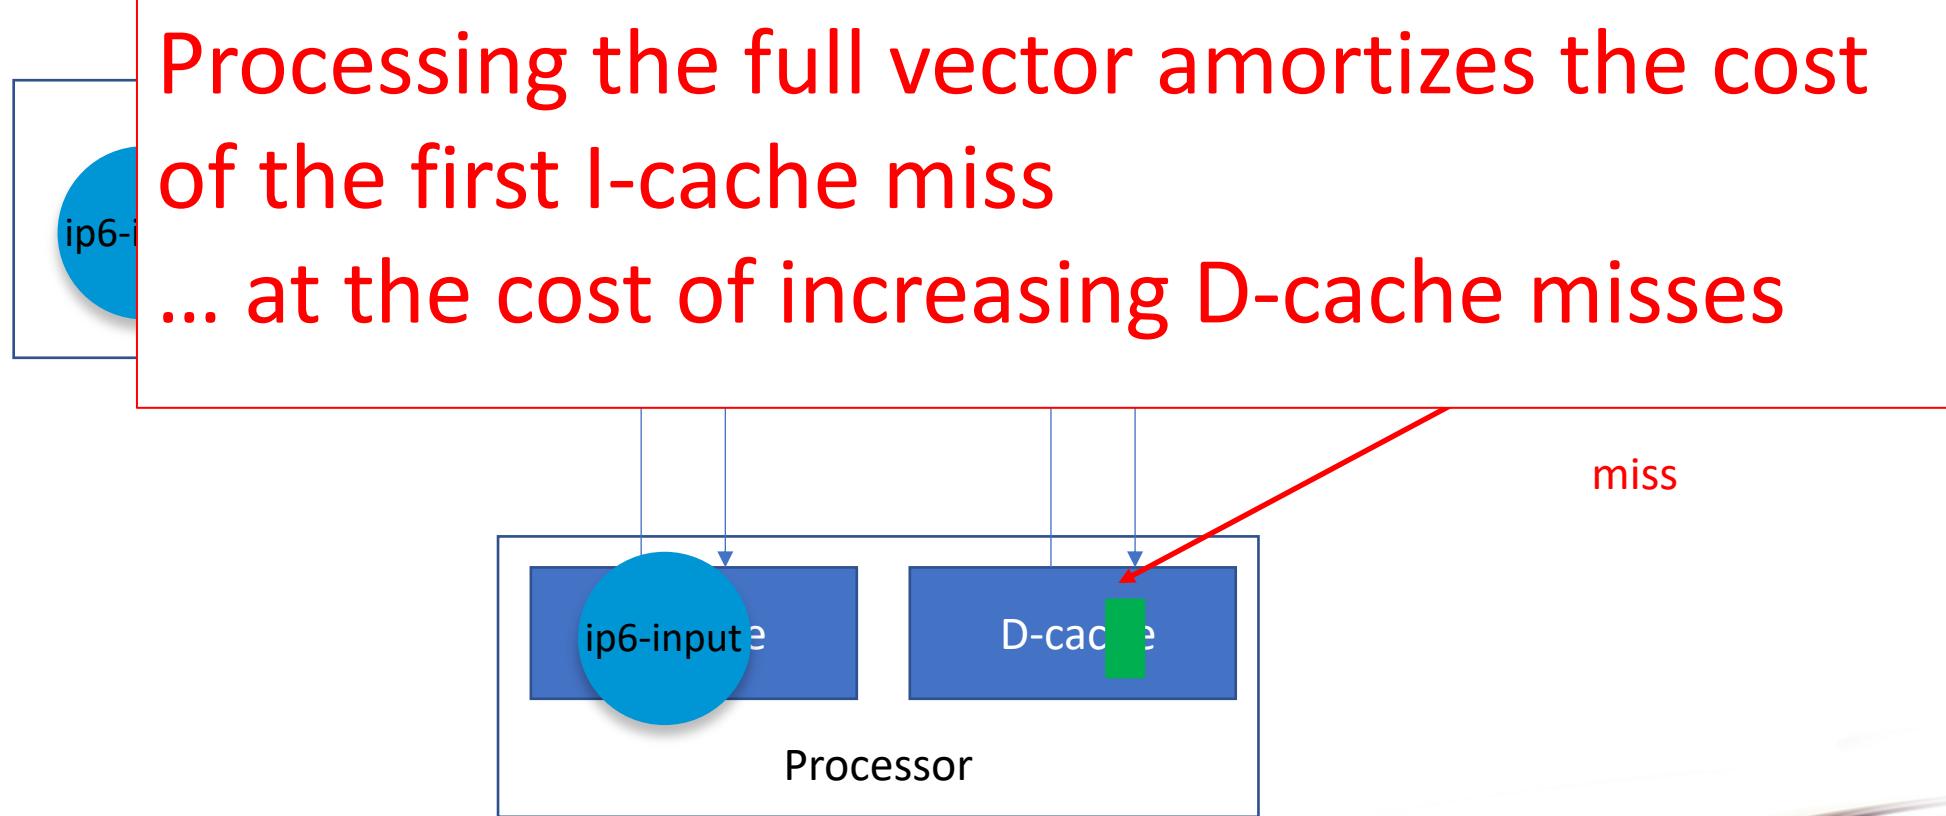
- Every node process the full packet vector



# Vector Packet Processing

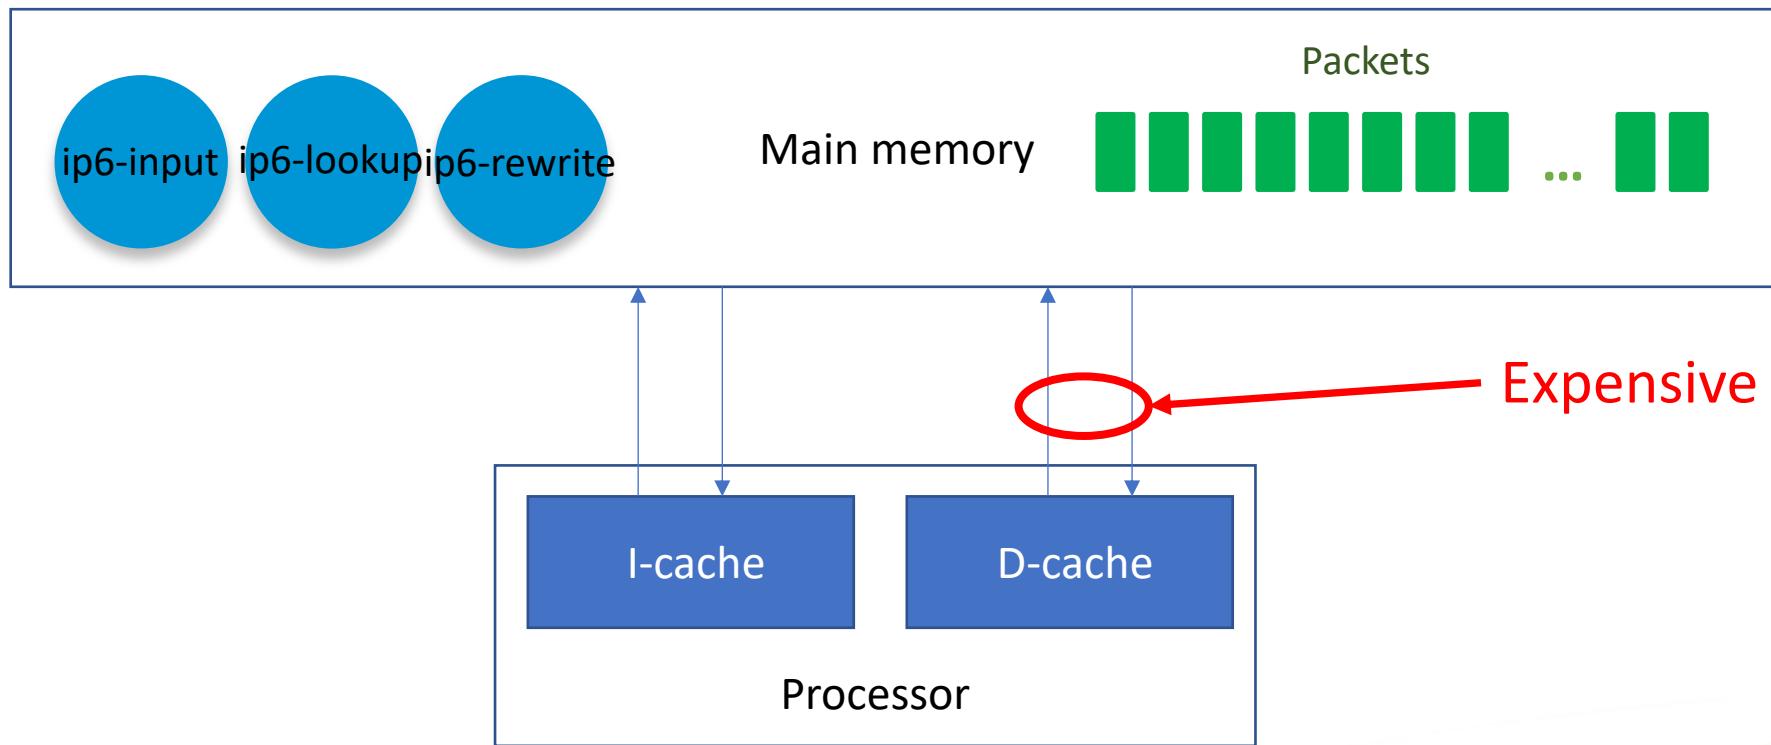
- Every node process the full packet vector

Processing the full vector amortizes the cost  
of the first I-cache miss  
... at the cost of increasing D-cache misses



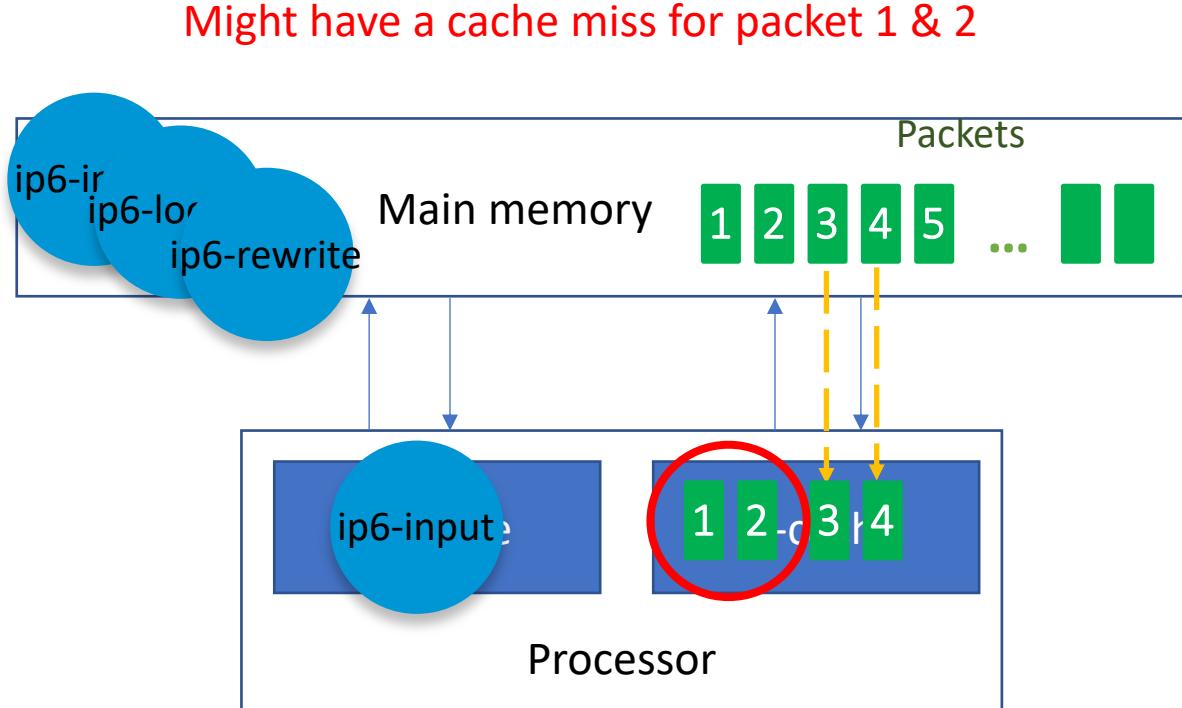
# Reduce cache miss – D-cache

VPP pre-fetches data into D-cache



# Reduce cache miss – D-cache

Example: Processing packet 1 & 2



VPP node pseudocode

**while** packets in vector

**while** 4 or more packets

PREFETCH #3 and #4

PROCESS #1 and #2

**while** any packets

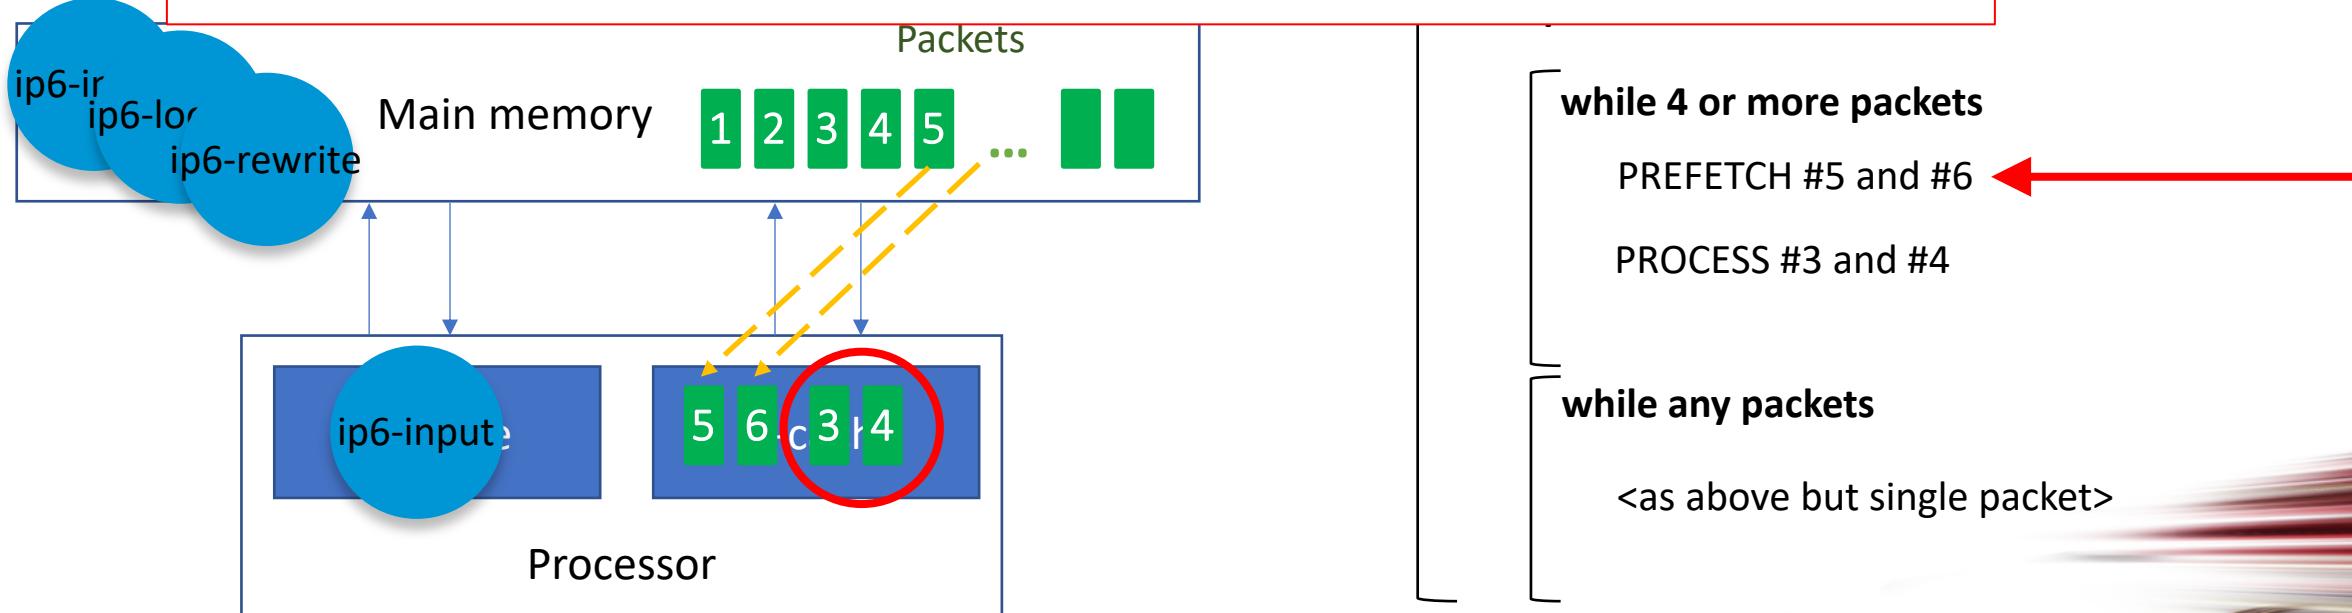
<as above but single packet>



# Reduce cache miss – D-cache

Example: Processing packet 3 & 4

The cost of the first D-cache miss is amortized by the subsequent D-cache hits.



# Hands on VPP!





# VPP documentation

- Wiki

<https://wiki.fd.io/view/VPP>

- Doxygen

<https://docs.fd.io/vpp/17.04/>



# Download VPP (v17.04)

- Clone the source code from git

git clone <https://gerrit.fd.io/r/vpp>

- Or install it from .deb pkg (rpm for Centos available too)

... see wiki



# Configure and Start VPP

- VPP configuration file

```
# emacs /etc/vpp/startup.conf
```

- Start vpp

```
# sudo vpp -c /etc/vpp/startup.conf
```



# VPP Command Line Interface

- To start a shell:

```
# vppctl
```

- To run one command:

```
# vppctl <command>
```



# VPP Command Line Interface

- A bunch of useful commands:

- ?
- show
- set

# Create your own plugin



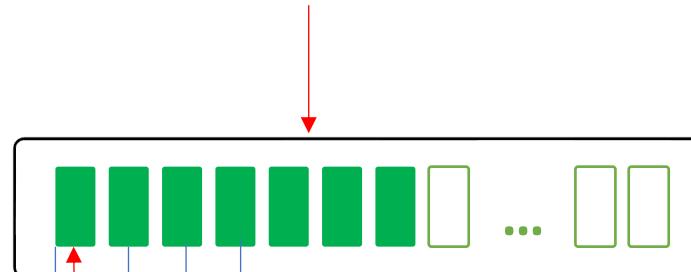


# Outline

- VPP structures
- Design & Implement your node(s)
- Insert your node(s) in the vlib\_graph
- Compile and install your plugin

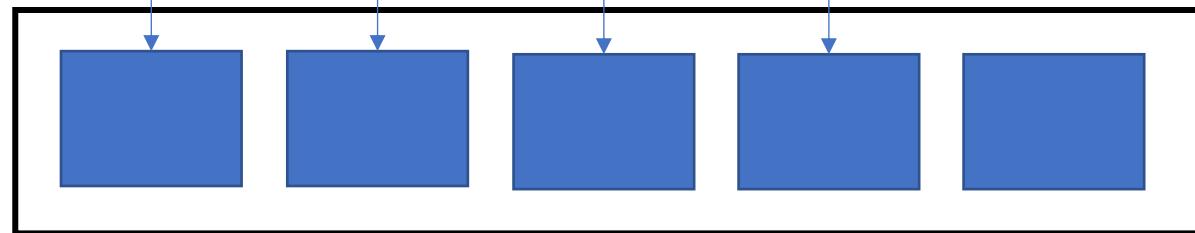
# VPP structures

The vector of packets is called **FRAME**



Each element is called **VECTOR**

A vector is an index to a **vlib\_buffer\_t**



Memory holding `vlib_buffer_t` objects

`vlib_buffer_t`

....

u8 data[0];

Pointer to packet data  
(DMA memory)



# Outline

- VPP structures
- Design & Implement your node(s)
- Inserting your node(s) in the vlib\_graph
- Compiling and installing your plugin



# Design & Implement your node(s)

- Your node should follow VPP style
  - Multi-loop, Branch prediction, Function flattening, Lock-free structures
- A node must implement a processing function that
  - “Moves vectors” from your node’s frame to the next node’s frame
  - Processes packets as YOU want
- Add whatever else you need
  - Supporting Functions, macros, variables, etc.. (C code)



# Register your node(s) to VPP

- Each node must be registered to VPP through `VLIB_REGISTER_NODE` macro

```
#define VLIB_REGISTER_NODE ( x,  
                           ...  
                         )
```

**Value:**

```
__VA_ARGS__ vlib_node_registration_t x;  
static void __vlib_add_node_registration_##x (void)  
  __attribute__((__constructor__));  
static void __vlib_add_node_registration_##x (void)  
{  
    vlib_main_t * vm = vlib_get_main();  
    x.next_registration = vm->node_main.node_registrations;  
    vm->node_main.node_registrations = &x;  
}  
__VA_ARGS__ vlib_node_registration_t x
```

Definition at line 143 of file `node.h`.

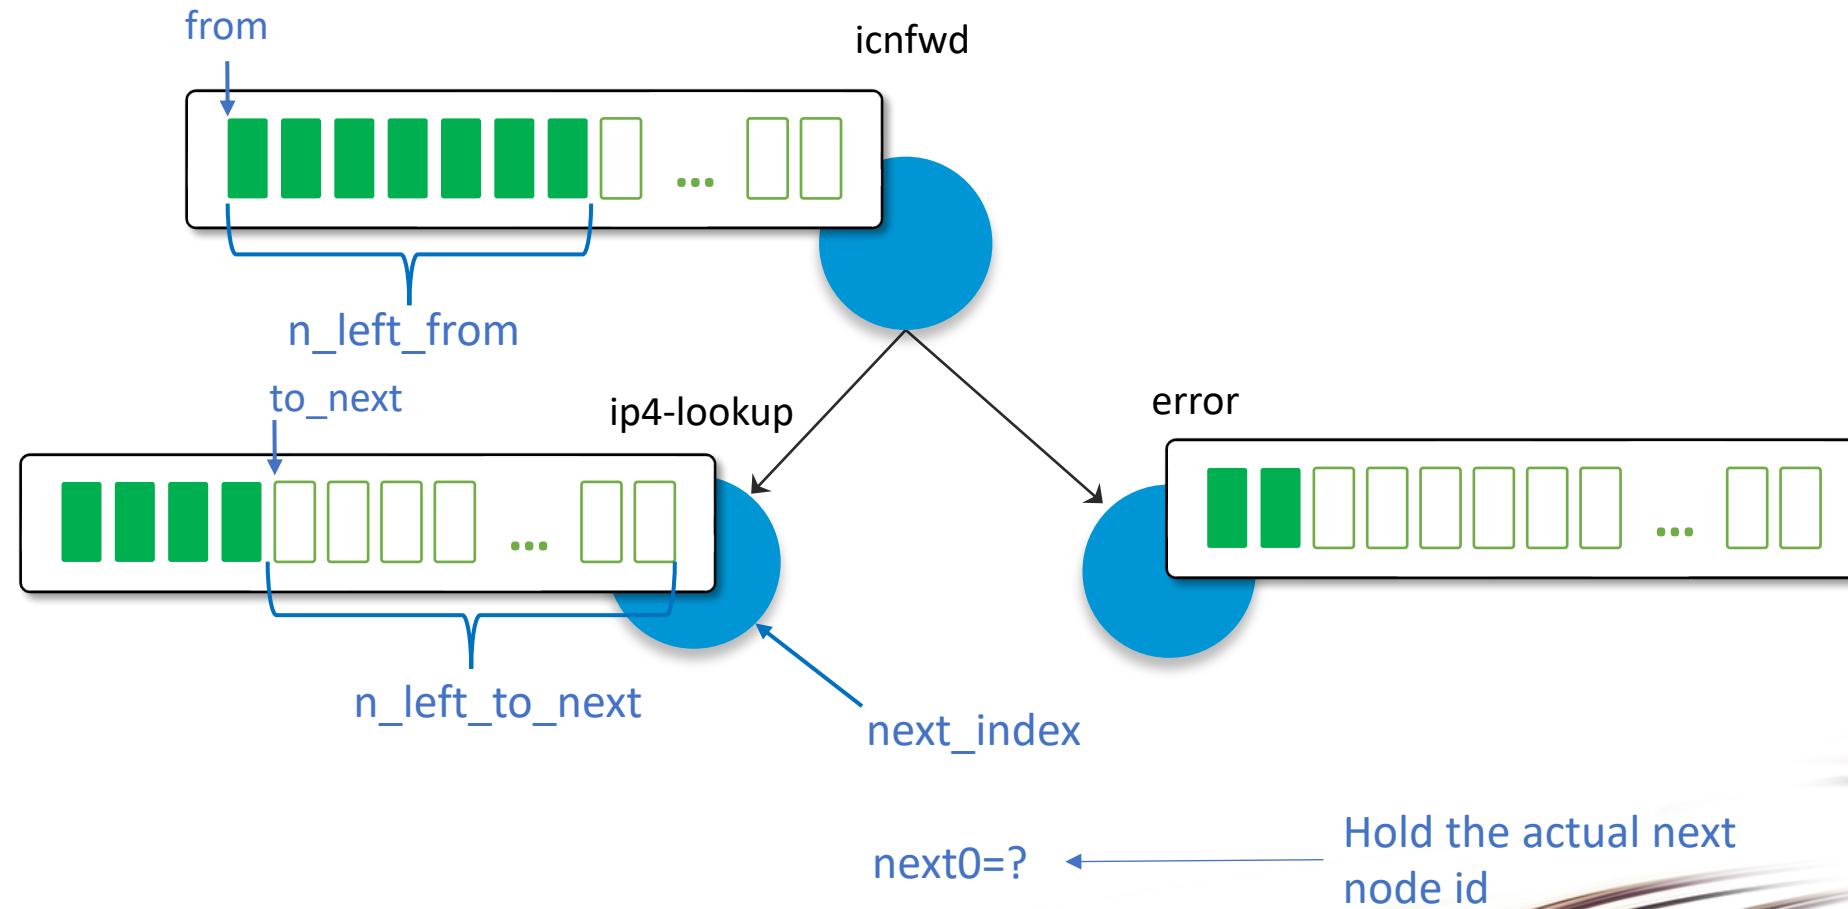
```
typedef struct _vlib_node_registration  
{  
    /* Vector processing function for this node. */  
    vlib_node_function_t *function;  
  
    /* Node name. */  
    char *name;  
  
    /* Name of sibling (if applicable). */  
    char *sibling_of;  
  
    /* Node index filled in by registration. */  
    u32 index;  
  
    /* Type of this node. */  
    vlib_node_type_t type;  
  
    /* Error strings indexed by error code for this node. */  
    char **error_strings;
```

# Example: Cicn plugin

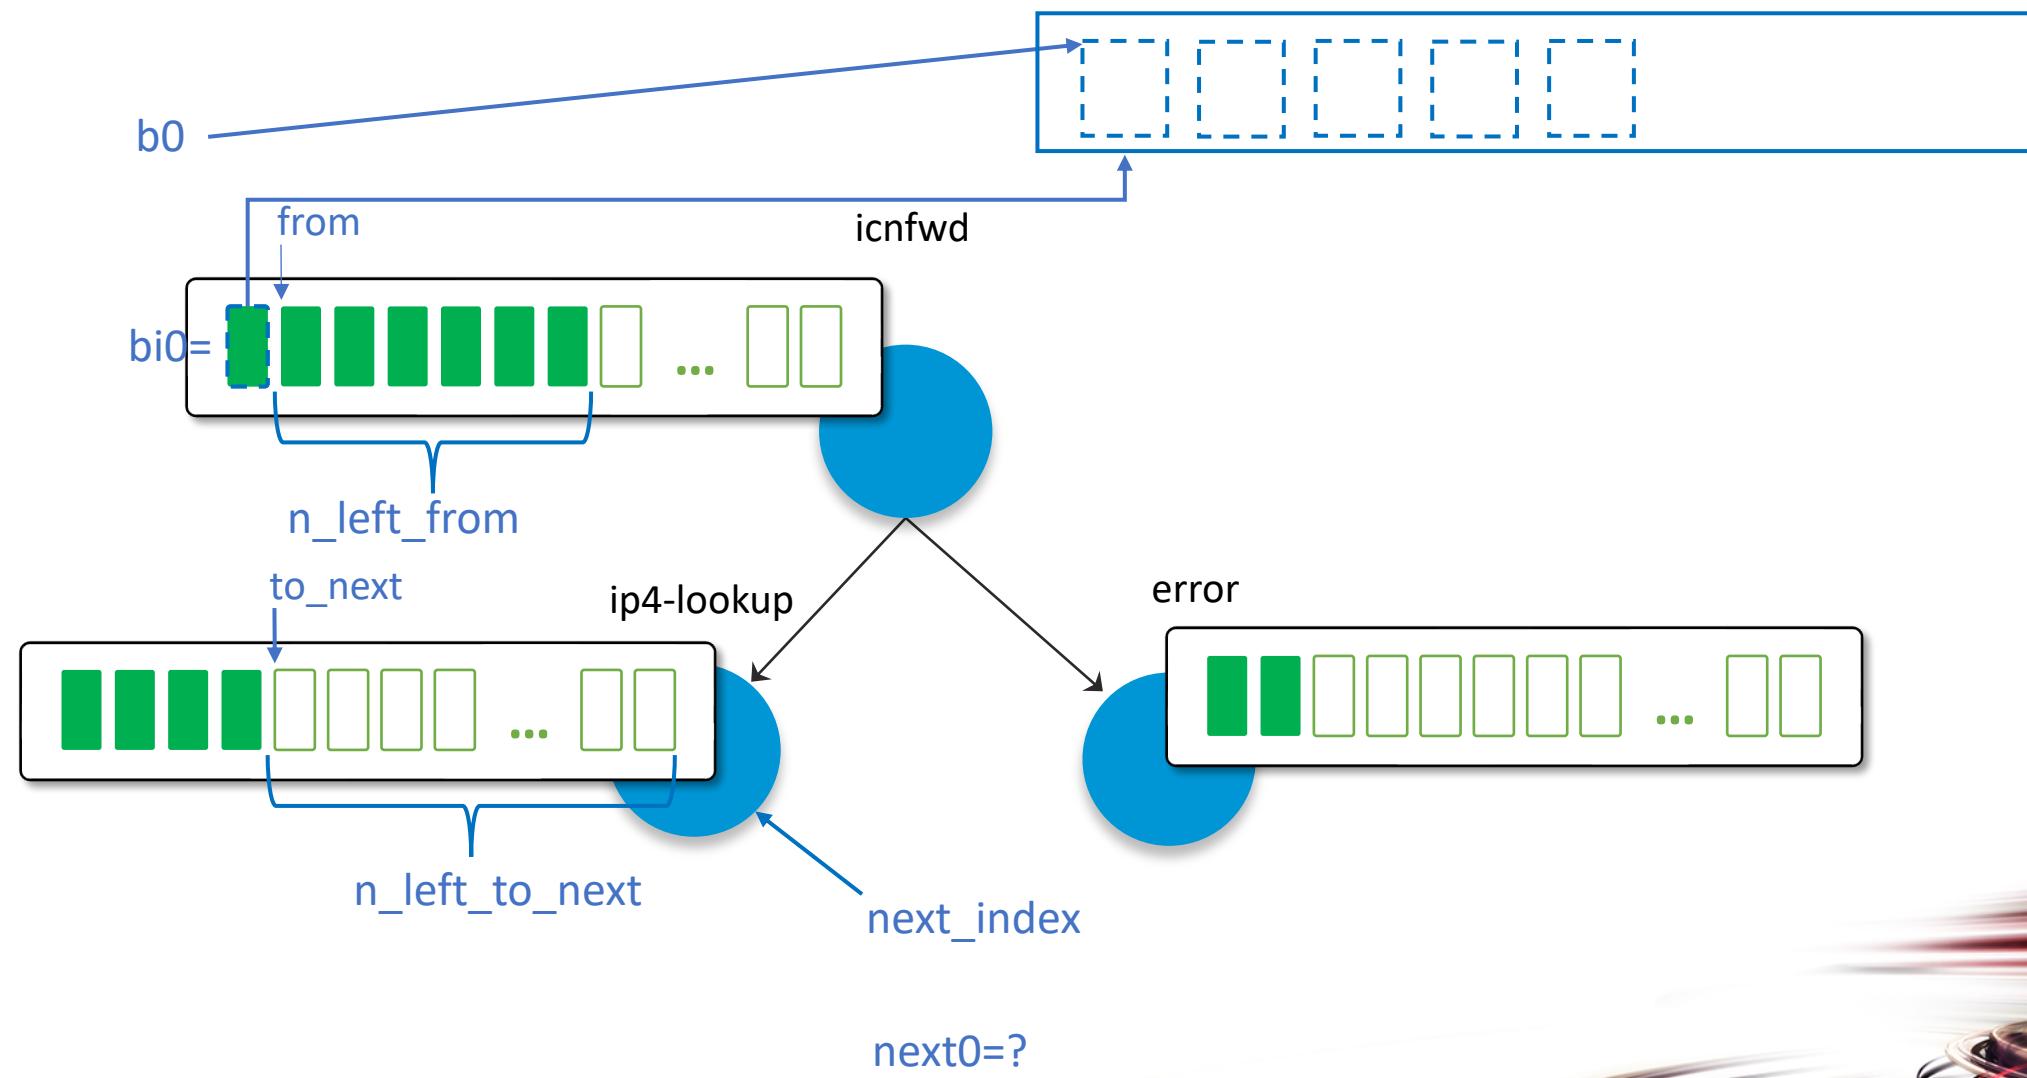
```
VLIB_REGISTER_NODE (icnfwd_node) =  
{  
    .function = icnfwd_node_fn,           ← Node processing function  
    .name = "icnfwd",                   ← Name of the node  
    .vector_size = sizeof (u32),  
    .runtime_data_bytes = sizeof (icnfwd_runtime_t), ← Runtime structure  
    .format_trace = icnfwd_format_trace,  
    .type = VLIB_NODE_TYPE_INTERNAL,     ← You need to initialize it by yourself  
    .n_errors = ARRAY_LEN (icnfwd_error_strings),  
    .error_strings = icnfwd_error_strings,  
    .n_next_nodes = ICNFWD_N_NEXT,  
    .next_nodes = {  
        [ICNFWD_NEXT_LOOKUP] = "ip4-lookup",  
        [ICNFWD_NEXT_ERROR_DROP] = "error-drop",  
    }  
};
```

Let's take a look to icnfwd\_node\_fn

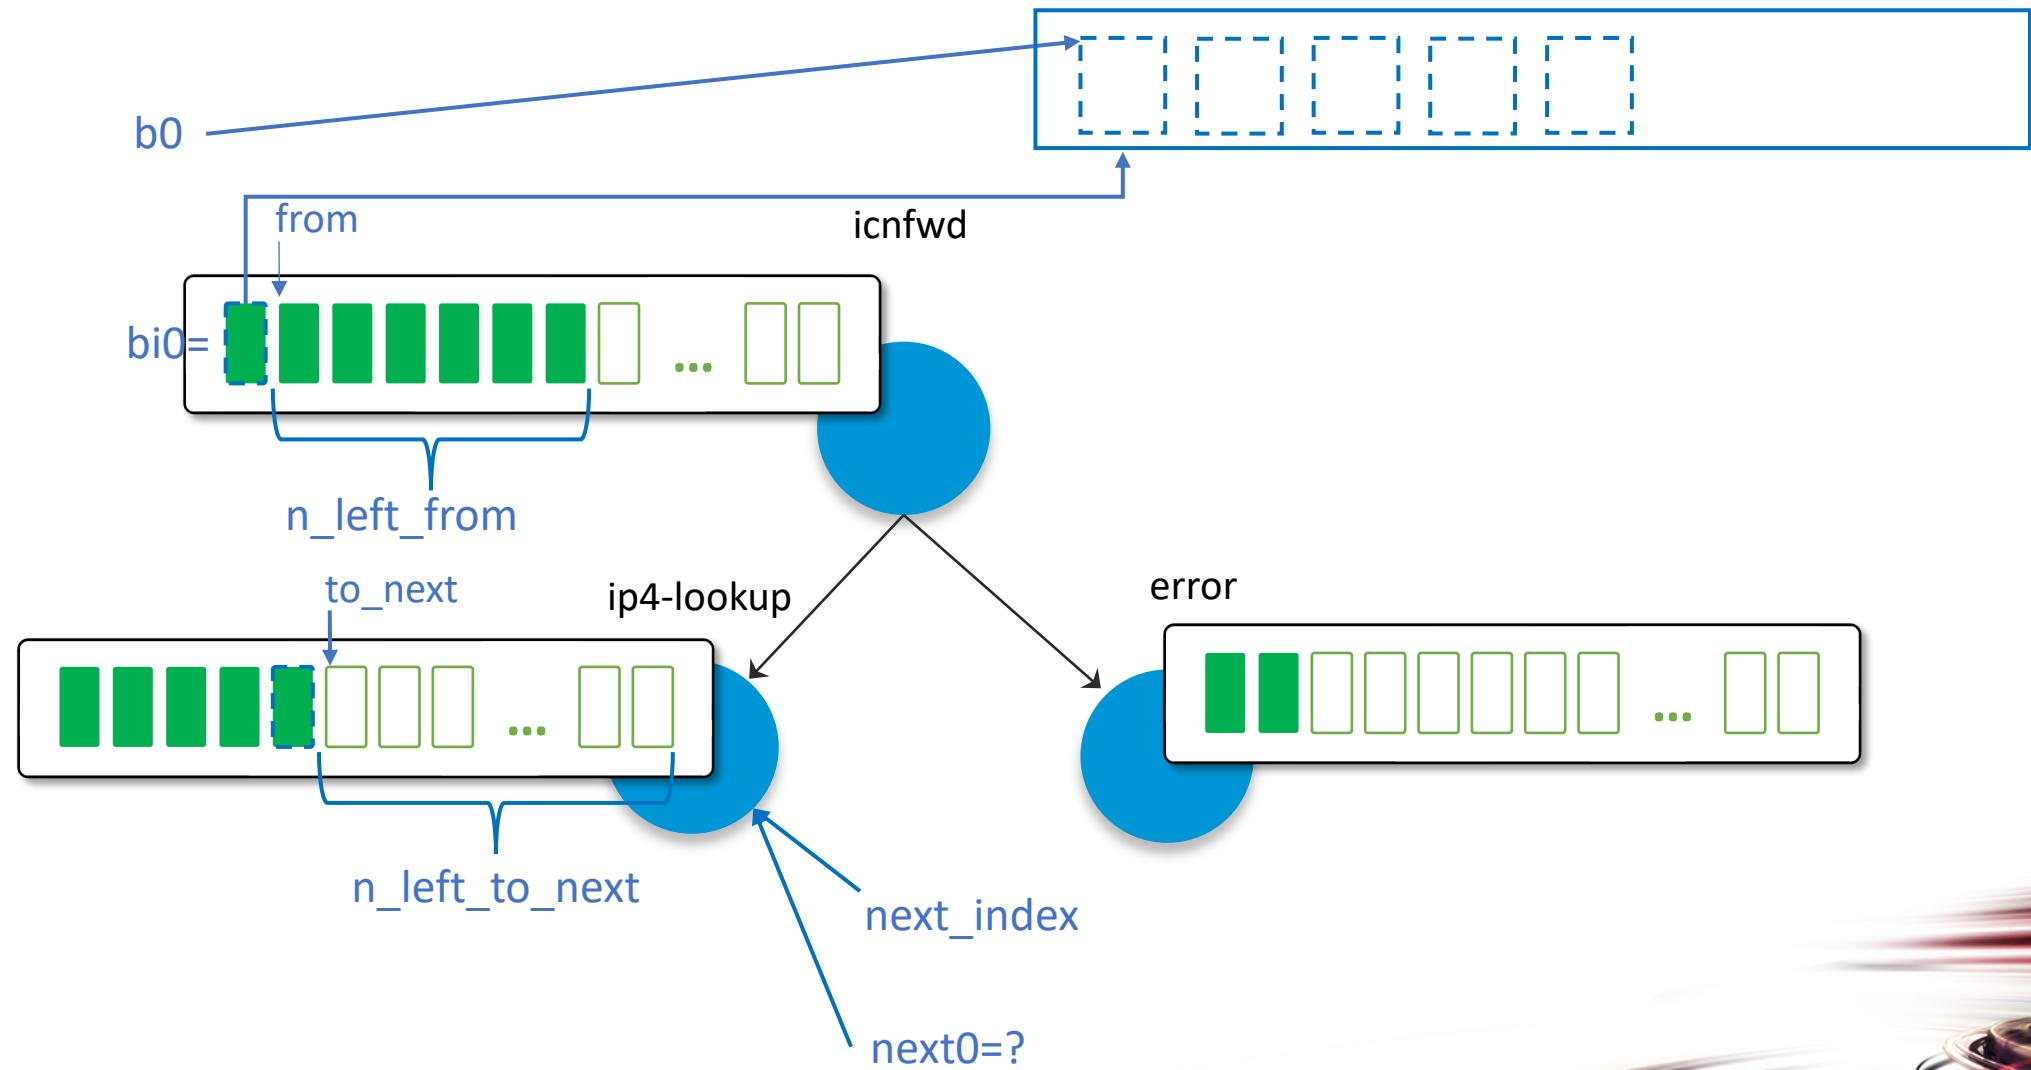
# icnfwd node



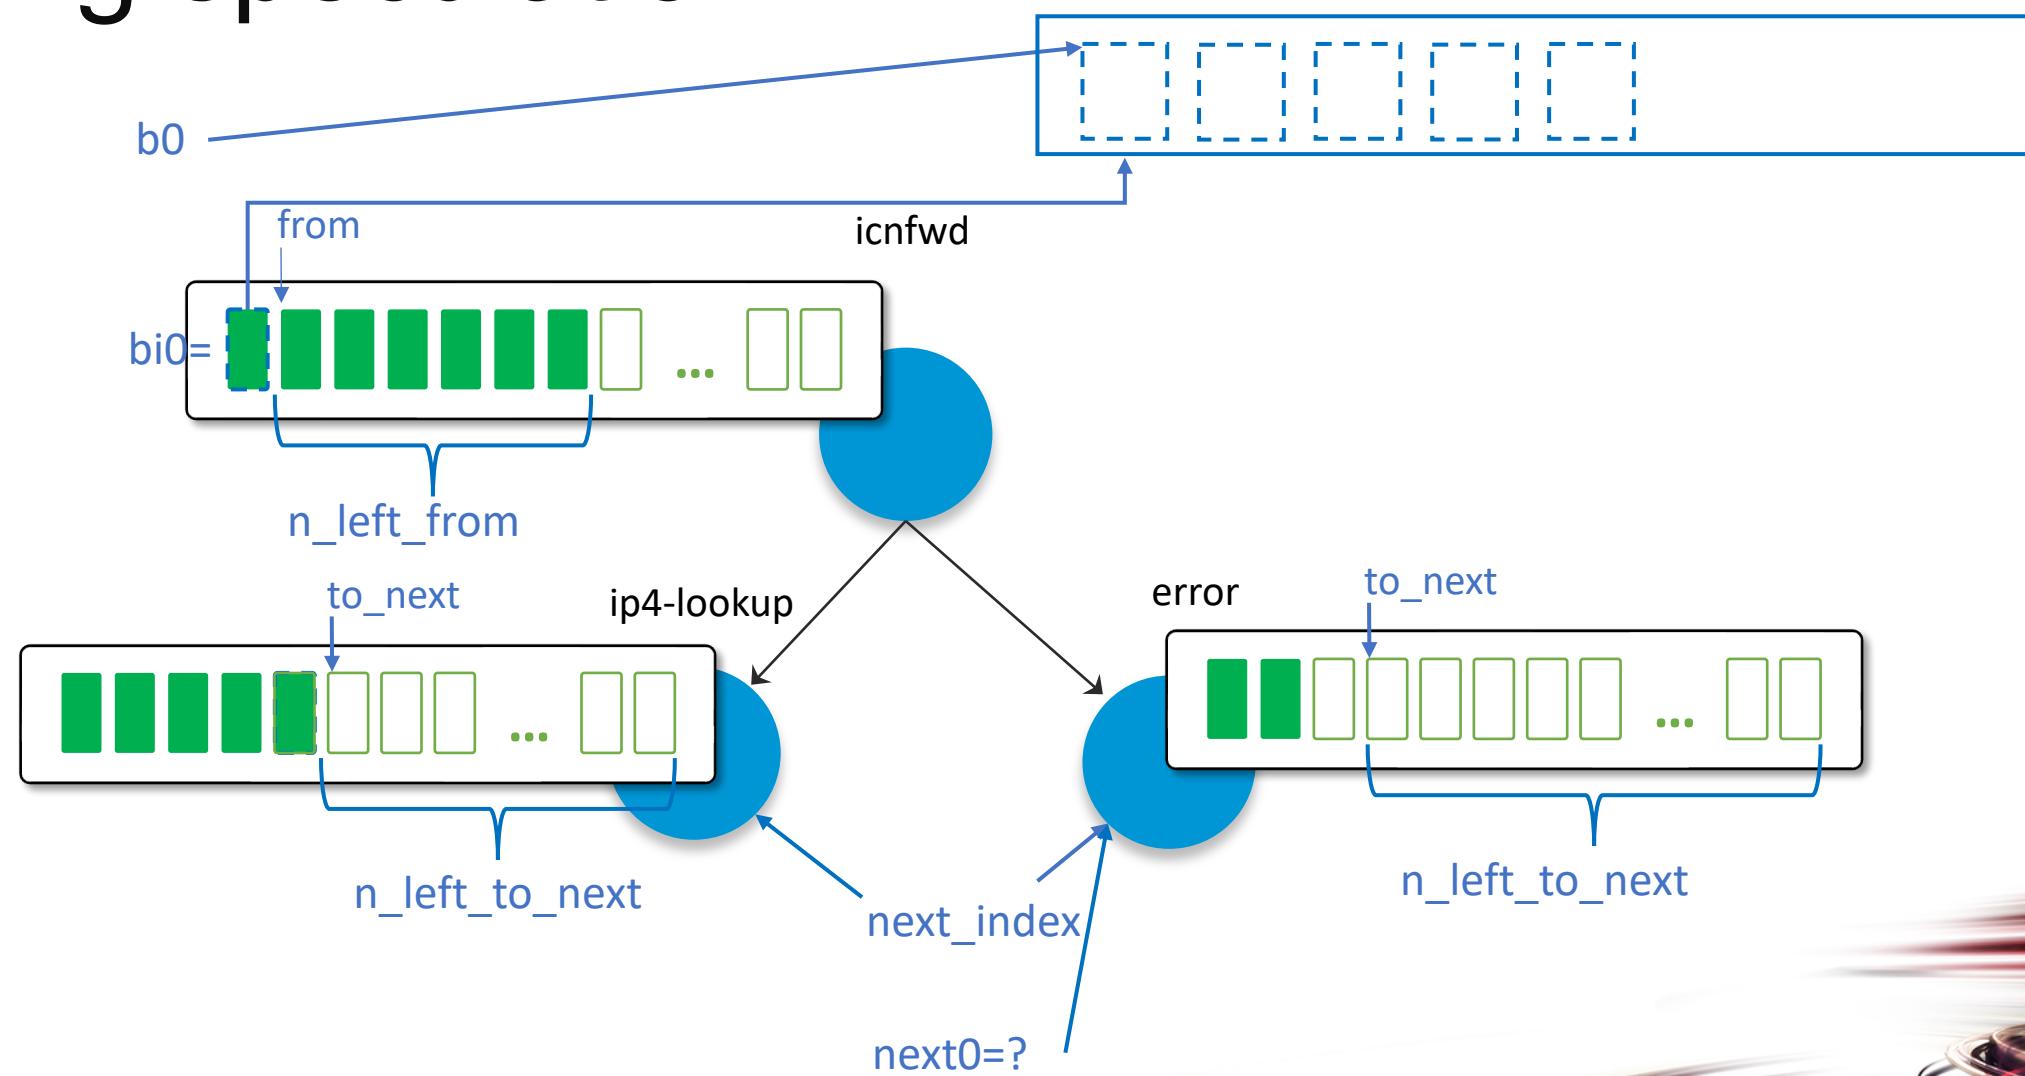
# icnfwd node



# icnfwd node



# Wrong speculation



# Example: Cicn plugin

```
VLIB_REGISTER_NODE (icnfwd_node) =  
{  
    .function = icnfwd_node_fn,  
    .name = "icnfwd",  
    .vector_size = sizeof (u32),  
    .runtime_data_bytes = sizeof (icnfwd_runtime_t),  
    .format_trace = icnfwd_format_trace,  
    .type = VLIB_NODE_TYPE_INTERNAL,  
    .n_errors = ARRAY_LEN (icnfwd_error_strings),  
    .error_strings = icnfwd_error_strings,  
    .n_next_nodes = ICNFWD_N_NEXT,  
    .next_nodes = {  
        [ICNFWD_NEXT_LOOKUP] = "ip4-lookup",  
        [ICNFWD_NEXT_ERROR_DROP] = "error-drop",  
    }  
,};
```

} ← Errors handling (counters)



# Other important macros

- **VPP\_INIT\_FUNCTION**
  - Function that is called during VPP initialization
- **VPP\_REGISTER\_PLUGIN**
  - Required to guarantee that your plugin is actually a VPP plugin  
...and not a library copied by mistake in /usr/lib/vpp\_plugins



# Outline

- VPP structures
- Design & Implement your node(s)
- Insert your node(s) in the vlib\_graph
- Compile and install your plugin



# Insert your node to VPP graph

1. direct all the packets from one interface
  - `vnet_hw_interface_rx_redirect_to_node (vnet_main, hw_if_index, my_graph_node.index /* redirect to my_graph_node */);`
2. capture packets with a particular ethertype
  - `ethernet_register_input_type (vm, ETHERNET_TYPE_CDP, cdp_input_node.index);`
3. for-us packet for new protocol on top of IP
  - `ip4_register_protocol (IP_PROTOCOL_GRE, gre_input_node.index);`



# Insert your node to VPP graph

4. ip-for-us packet sent to a specific UDP port
  - `udp_register_dst_port (vm, UDP_DST_PORT_vxlan,  
vxlan_input_node.index, 1 /* is_ip4 */);`
5. direct all packets from one ip prefix
  - Create your own Data Path Object (i.e. result of a FIB lookup)



# Outline

- VPP structures
- Design & Implement your node(s)
- Insert your node(s) in the vlib\_graph
- Compile and install your plugin

# Compiling your plugin

- VPP provides Automake/Autoconf examples
  - Install vpp-dev and move to /usr/share/doc/vpp/examples
- Adapting Makefile.am and sample.am is trivial
- Compile cicn-plugin:

```
$ cd cicn-plugin
$ autoreconf -i -f
$ mkdir -p build
$ cd build
$ ../configure --with-plugin-toolkit
OR, to omit UT code
$ ../configure --with-plugin-toolkit --without-cicn-test
$ make
$ sudo make install
```

# vICN: configuration, management and control of an virtual ICN network

Marcel Enguehard

ACM ICN Conference – CICN tutorial

September 26<sup>th</sup> 2017



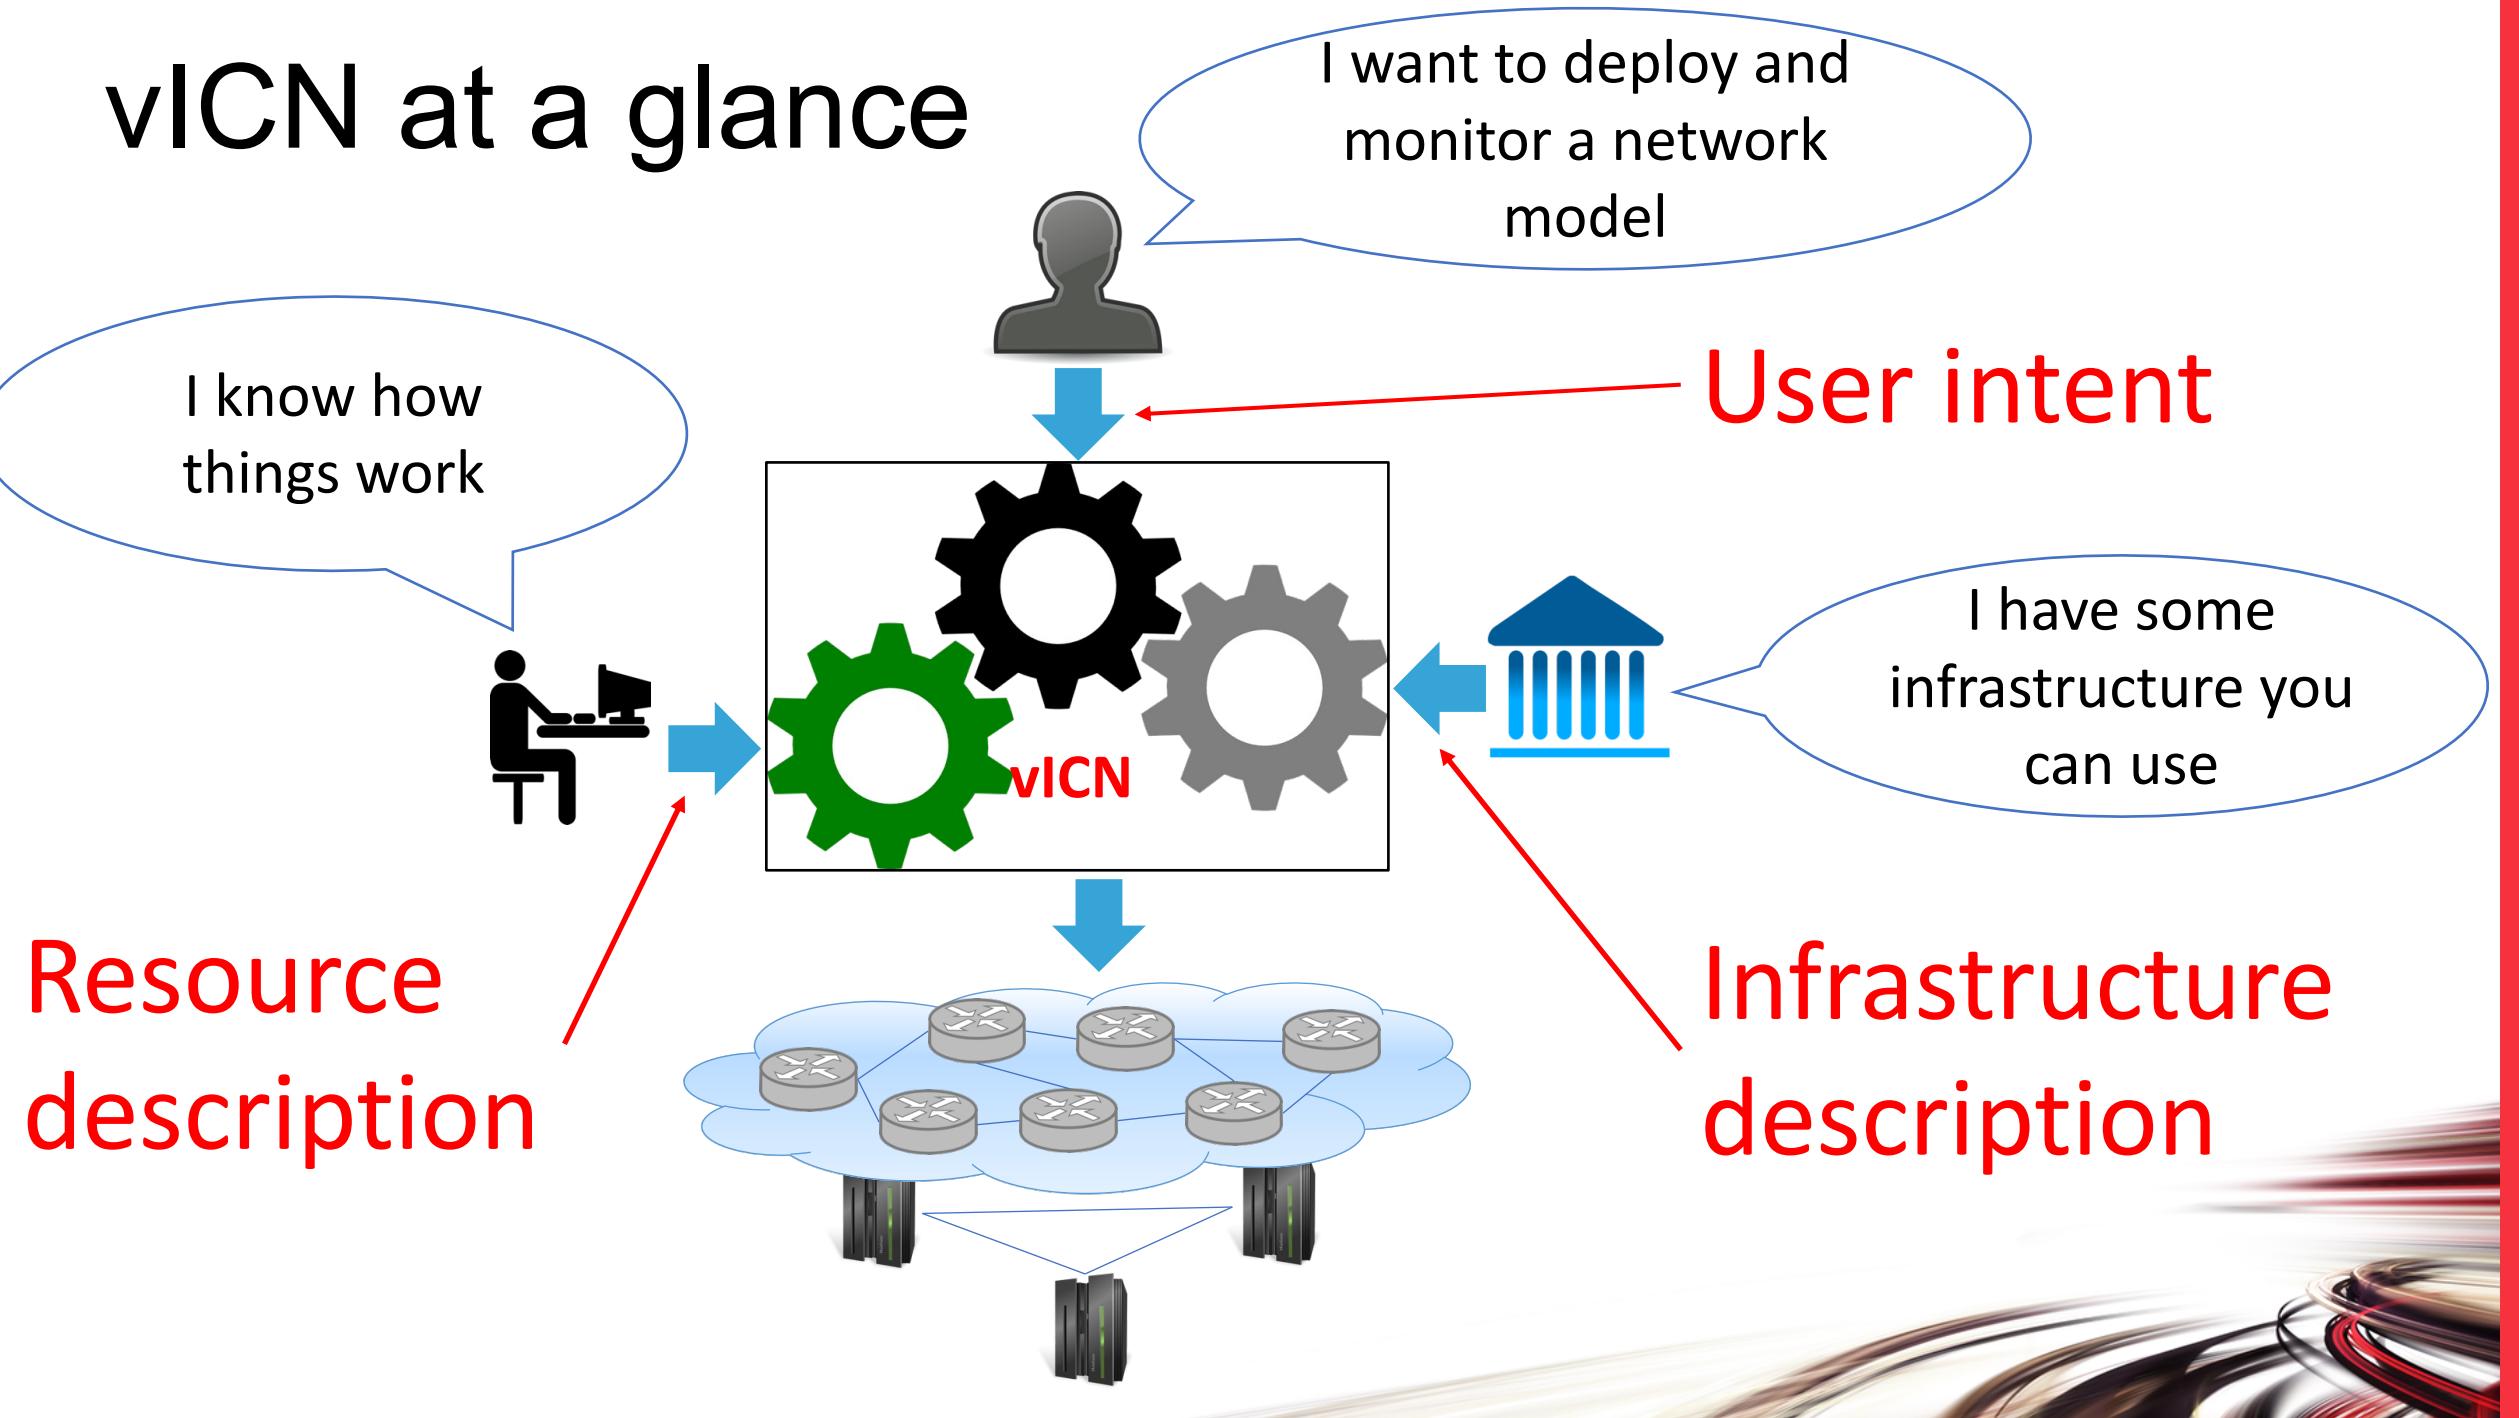


# What is vICN

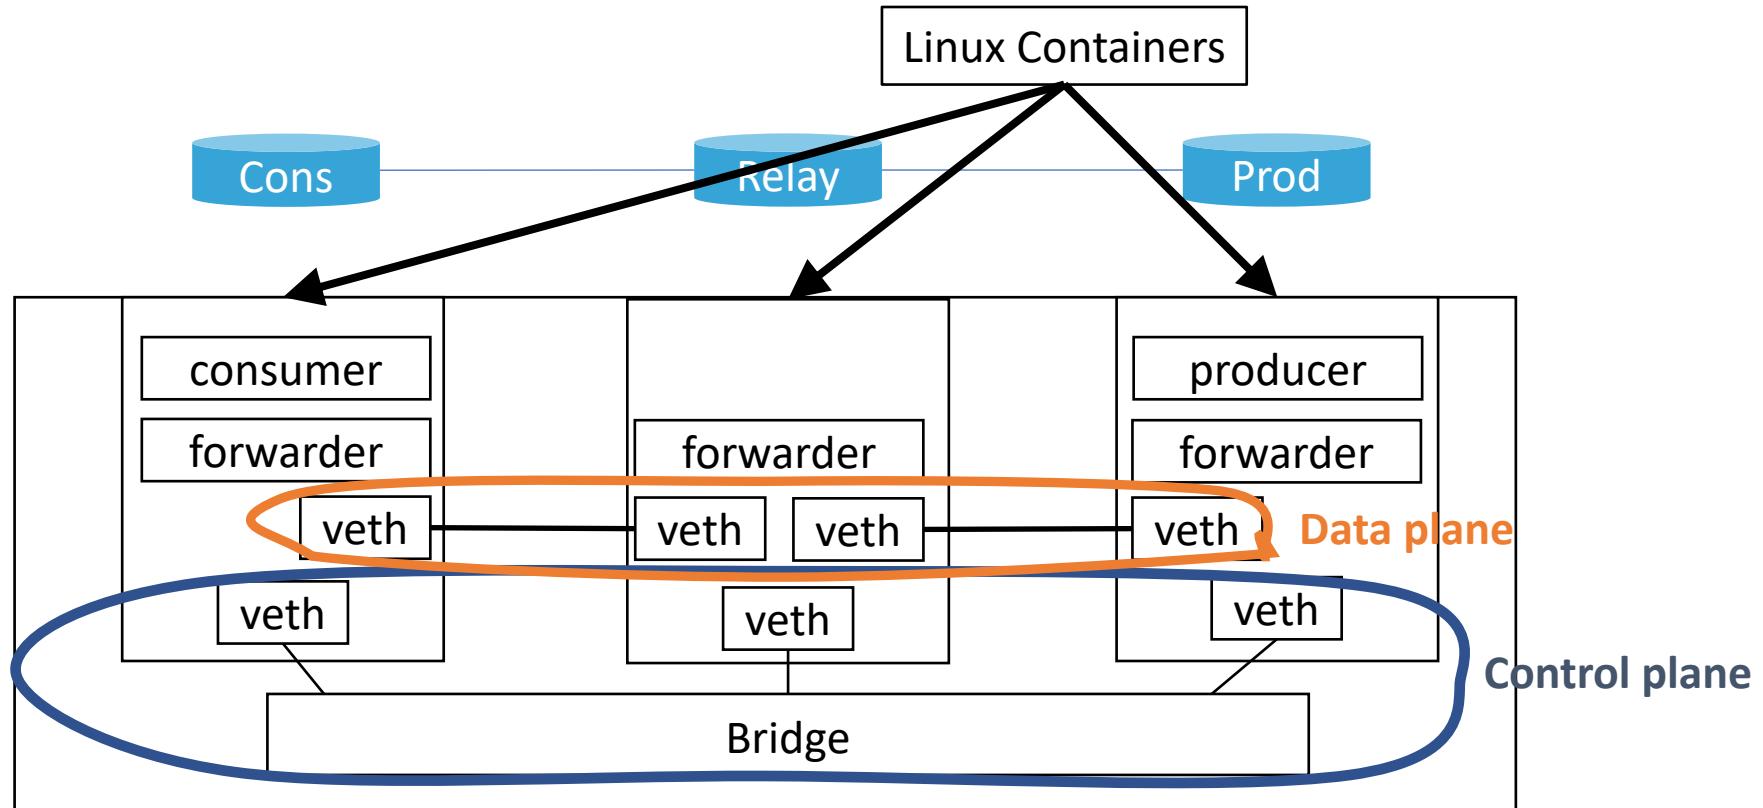
- Unified framework for network deployment, management and monitoring
- Integrates all the tools of the CICN fd.io suite
- Provides an API to easily bootstrap ICN deployments and get meaningful telemetry out of it



# vICN at a glance



# Example vICN topology



# vICN resources

Class

- Virtual representation of deployment element
- Node, forwarder, application, link, etc.
- Described by *attributes*

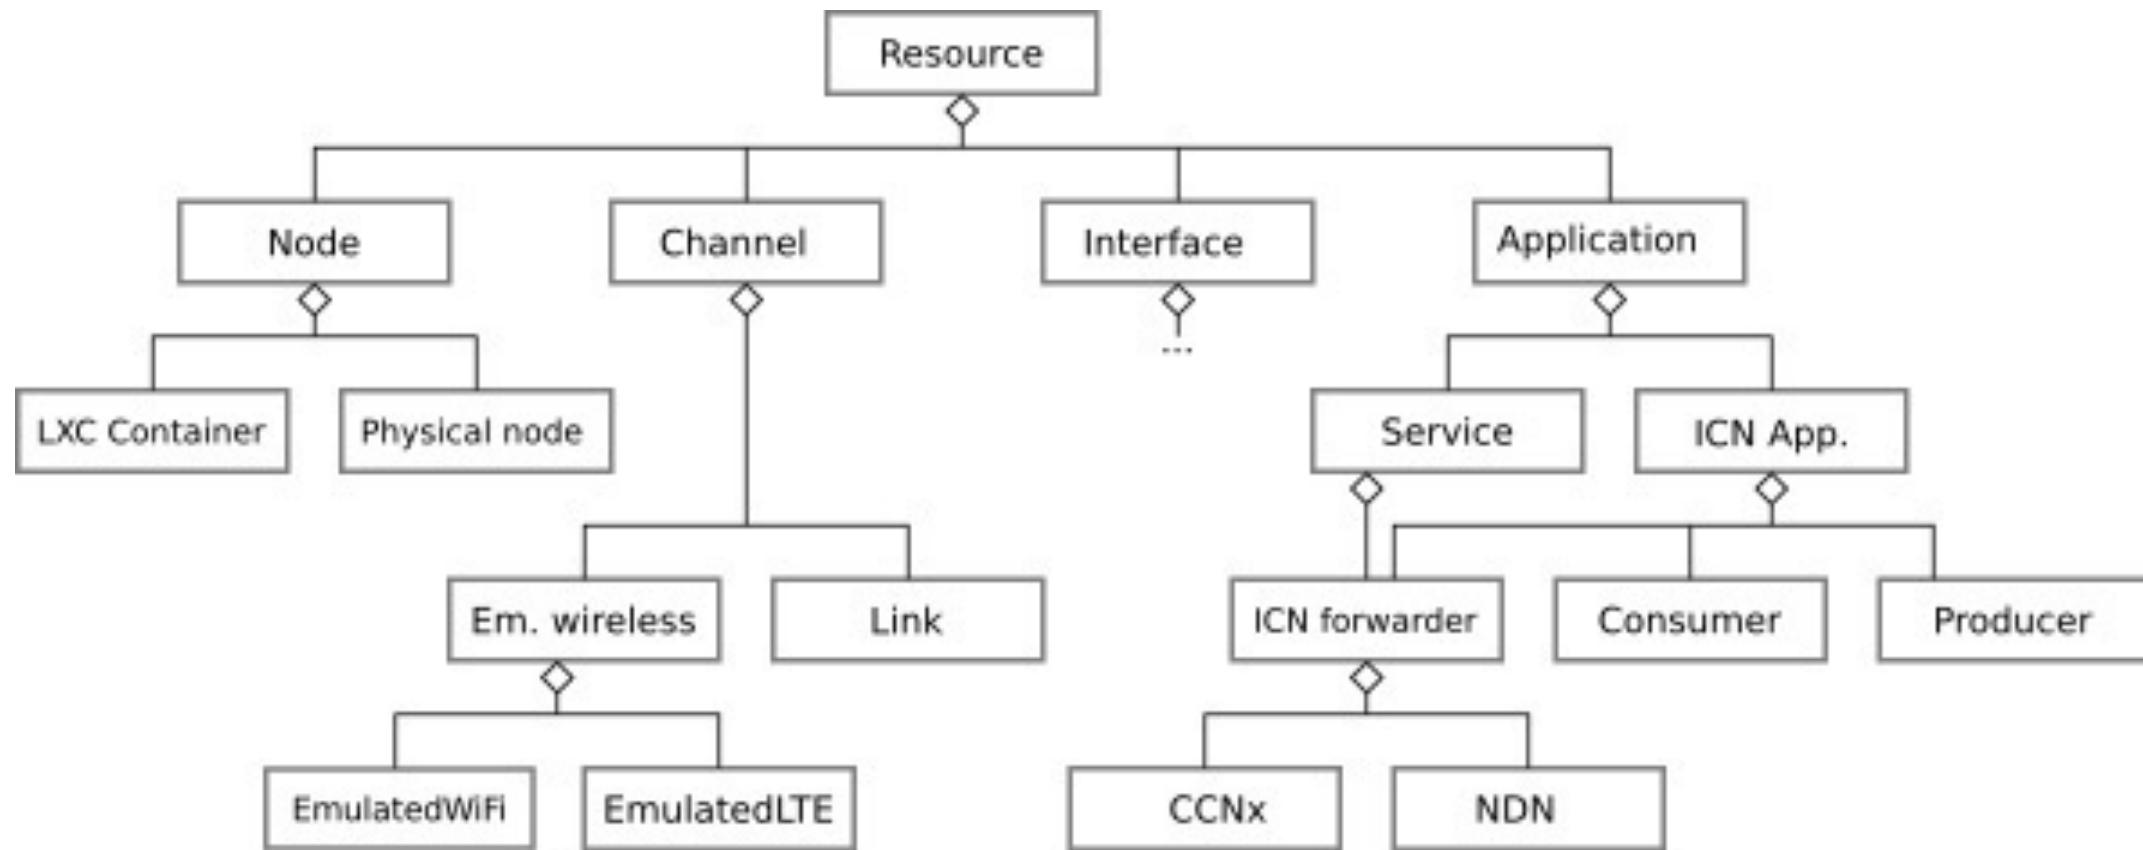
Members



# Example resource: forwarder

- Represents an ICN forwarder
- Attributes:
  - node
  - cache\_size
  - cache\_policy (e.g., LRU)
  - log\_file
  - etc.

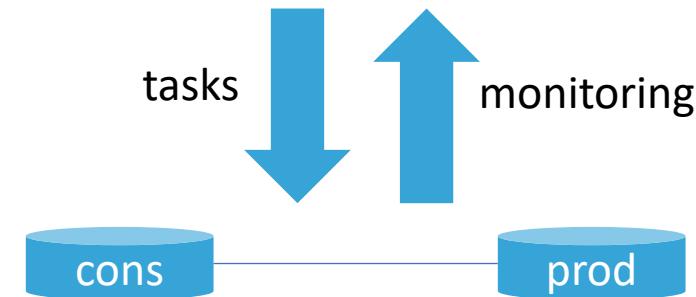
# Resource hierarchy



# How does it work?

- Intent based-framework
- Object-based model
- State reconciliation between model and deployment

```
cons = LxcContainer()  
prod = LxcContainer()  
link = Link(src=cons, dst=prod)
```

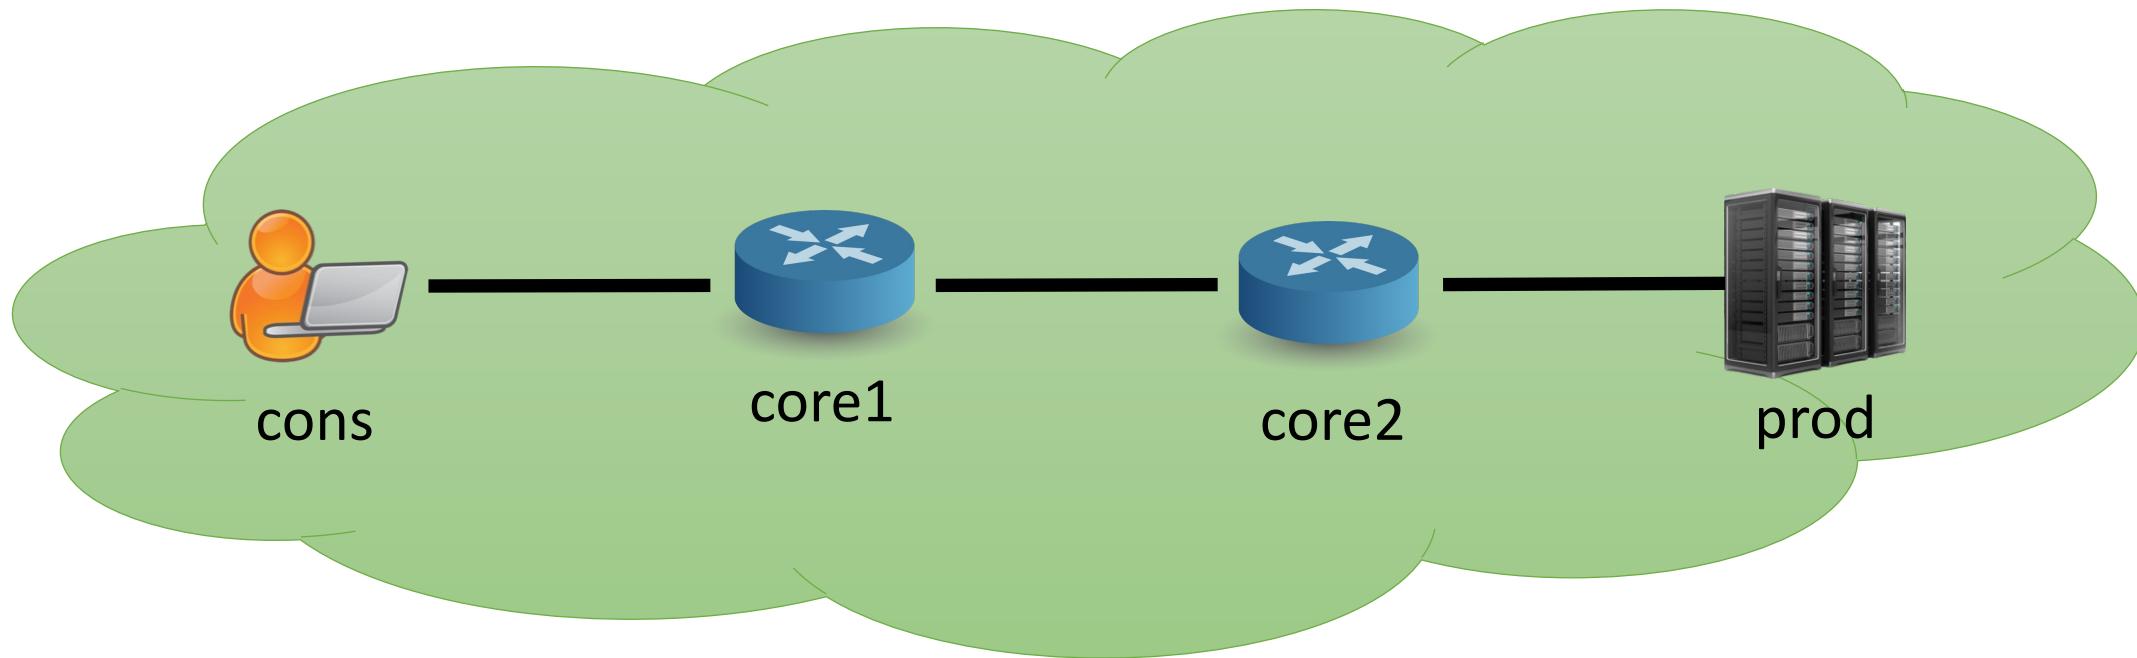




# vICN functionalities

- Multithreaded deployment of network models
- SDN controller for IPv4, IPv6, and ICN
- Wireless links emulation
- Connection of real devices
- Built-in monitoring through Python model

# Our example deployment



# Network model deployment





# Network model declaration

- JSON file containing list of resources
- Resources complemented with “key” attributes
- Intent-based declaration: descriptive approach (not imperative)

# Physical resources

```
{  
    "type": "Physical",  
    "name": "server",  
    "hostname": "localhost"  
},  
{  
    "type": "LxcImage",  
    "name": "cicn-image",  
    "node": "server",  
    "image": "ubuntu1604-cicnsuite-rc3"  
}
```

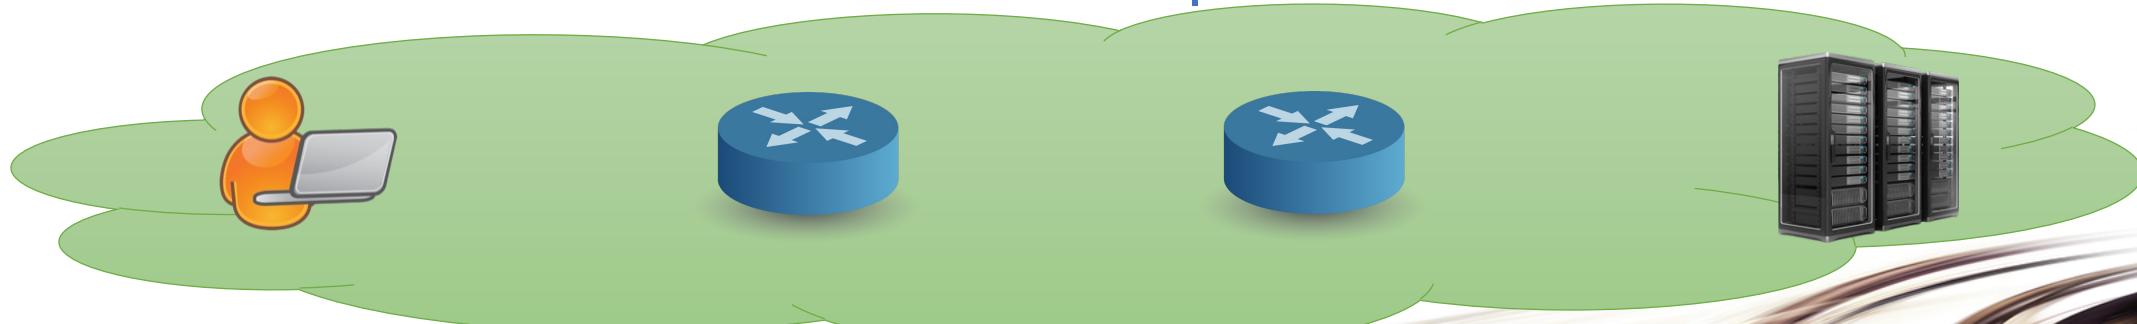


# Nodes

```
{  
    "type" : "LxcContainer",  
    "image": "cicn-image",  
    "name" : "cons",  
    "groups": [ "virtual" ],  
    "node" : "server"  
},  
{  
    "type" : "LxcContainer",  
    "image": "cicn-image",  
    "name" : "prod",  
    "groups": [ "virtual" ],  
    "node" : "server"  
},
```

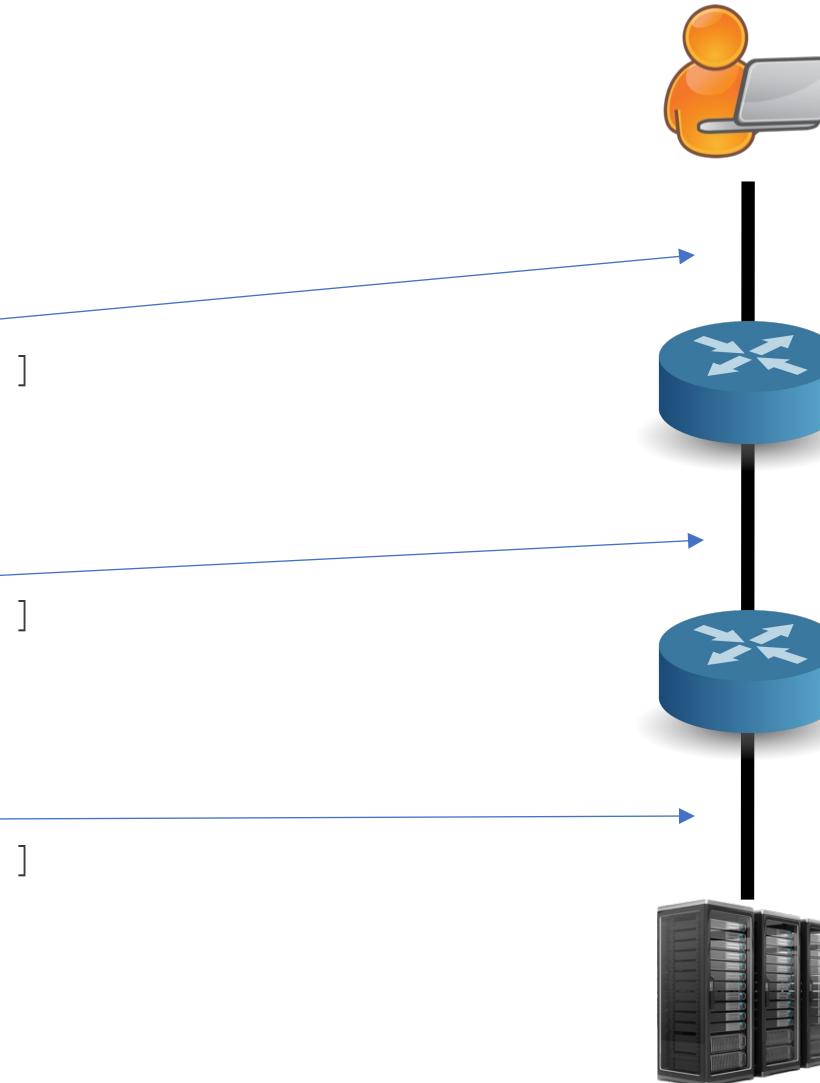
Refer to topology previous resources

```
{  
    "type" : "LxcContainer",  
    "image": "cicn-image",  
    "name" : "core1",  
    "groups": [ "virtual" ],  
    "node" : "server"  
},  
{  
    "type" : "LxcContainer",  
    "image": "cicn-image",  
    "name" : "core2",  
    "groups": [ "virtual" ],  
    "node" : "server"  
},
```



# Links

```
{  
    "type": "Link",  
    "src_node": "cons",  
    "dst_node": "core1",  
    "groups": [ "virtual" ]  
},  
{  
    "type": "Link",  
    "src_node": "core1",  
    "dst_node": "core2",  
    "groups": [ "virtual" ]  
},  
{  
    "type": "Link",  
    "src_node": "core2",  
    "dst_node": "prod",  
    "groups": [ "virtual" ]  
},
```



# IP networking on topology

```
{  
    "type": "CentralIP",  
    "ip4_data_prefix": "192.168.19.0/24",  
    "ip6_data_prefix": "9001::/16",  
    "ip_routing_strategy": "spt",  
    "groups": [  
        "virtual"  
    ]  
}
```

Defines objects  
on which to act



IPv6 addresses are  
attributed by /64

CentralIP is similar to an SDN controller that assigns addresses and sets up the routing in the network:

CentralIP = (Ipv4Assignment | Ipv6Assignment) > IPRoutes

# ICN forwarders

```
{  
    "type": "MetisForwarder",  
    "cache_size": 0,  
    "node": "cons"  
},  
{  
    "type": "MetisForwarder",  
    "cache_size": 2000,  
    "node": "core1"  
},  
{  
    "type": "MetisForwarder",  
    "cache_size": 0,  
    "node": "core2"  
},  
{  
    "type": "MetisForwarder",  
    "cache_size": 0,  
    "node": "prod"  
},
```

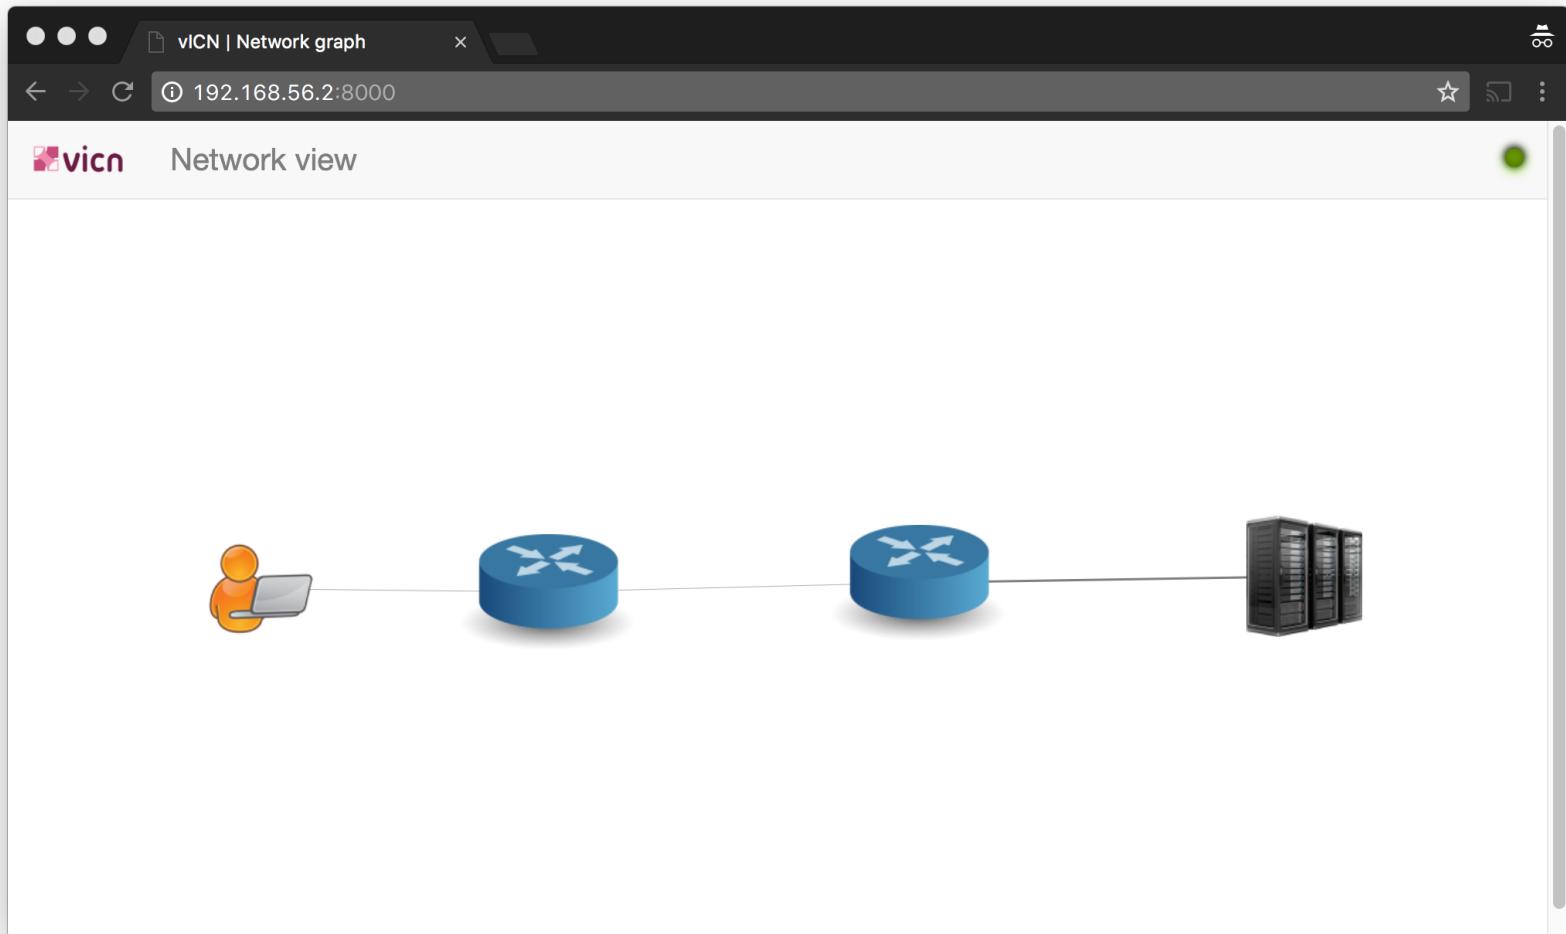
```
{  
    "type": "WebServer",  
    "prefixes": [  
        "/webserver"  
    ],  
    "node": "prod"  
},  
{  
    "type": "CentralICN",  
    "groups": [ "virtual" ],  
    "face_protocol": "udp4"  
}  
}
```

Like CentralIP

ether, udp4, udp6,  
tcp4, tcp6

# GUI

```
{  
    "type": "GUI",  
    "groups": ["virtual"]  
},
```



# Launching vicn

```
cicn@cicn-VirtualBox:~/vicn$ sudo vicn/bin/vicn.py -s  
examples/tutorial/tutorial06-acm-icn17.json
```

[...]

```
2017-09-21 17:48:15,023 - vicn.core.task - INFO - Scheduling task <Task[apy]  
partial<_task_resource_update>> for resource <UUID MetisForwarder-MPDRB>  
2017-09-21 17:48:15,024 - vicn.core.resource_mgr - INFO - Resource <UUID  
MetisForwarder-MPDRB> is marked as CLEAN (99/104)  
2017-09-21 17:48:15,146 - vicn.core.task - INFO - Scheduling task <Task[apy]  
partial<_task_resource_update>> for resource <UUID MetisForwarder-NC33W>  
2017-09-21 17:48:15,148 - vicn.core.resource_mgr - INFO - Resource <UUID  
MetisForwarder-NC33W> is marked as CLEAN (100/104)
```



# Traffic creation

## Producer setup:

- **producer-test**

```
producer-test -D ccnx:/webserver
```

- **Webserver**

```
http-server -p $server_folder -l  
http://webserver
```

## Consumer setup

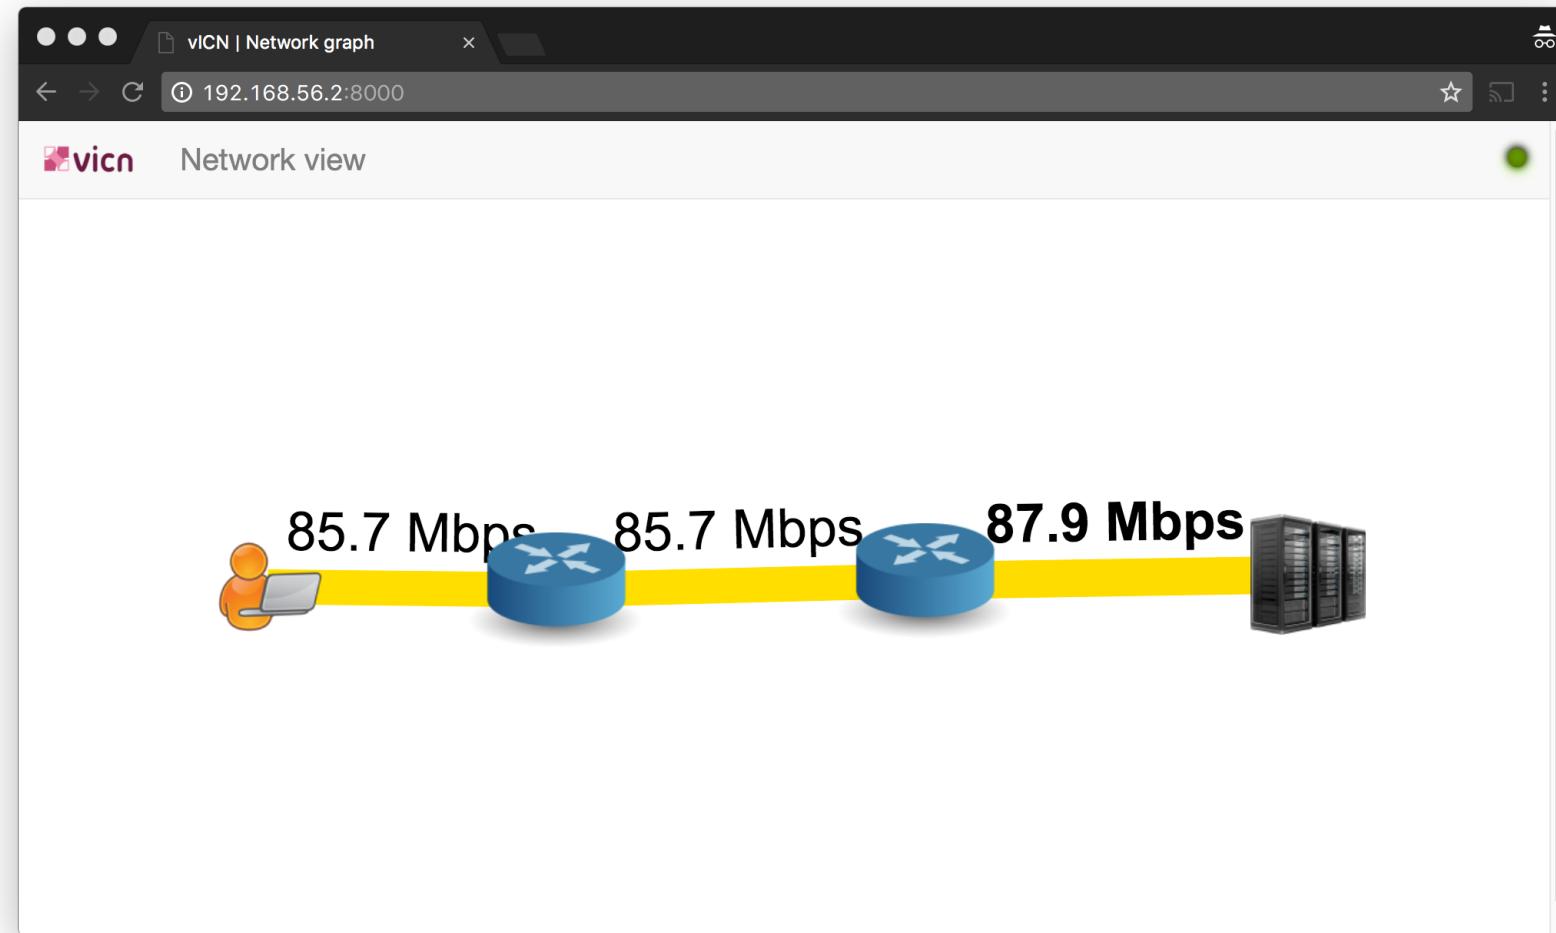
- **consumer-test**

```
consumer-test -D ccnx:/webserver
```

- **iget**

```
iget http://webserver/$filename
```

# Traffic visualization on the GUI



# Network teardown

```
cicn@cicn-VirtualBox:~/vicn$ sudo ./scripts/topo_cleanup.sh  
examples/tutorial/tutorial06-acm-icn17.json
```

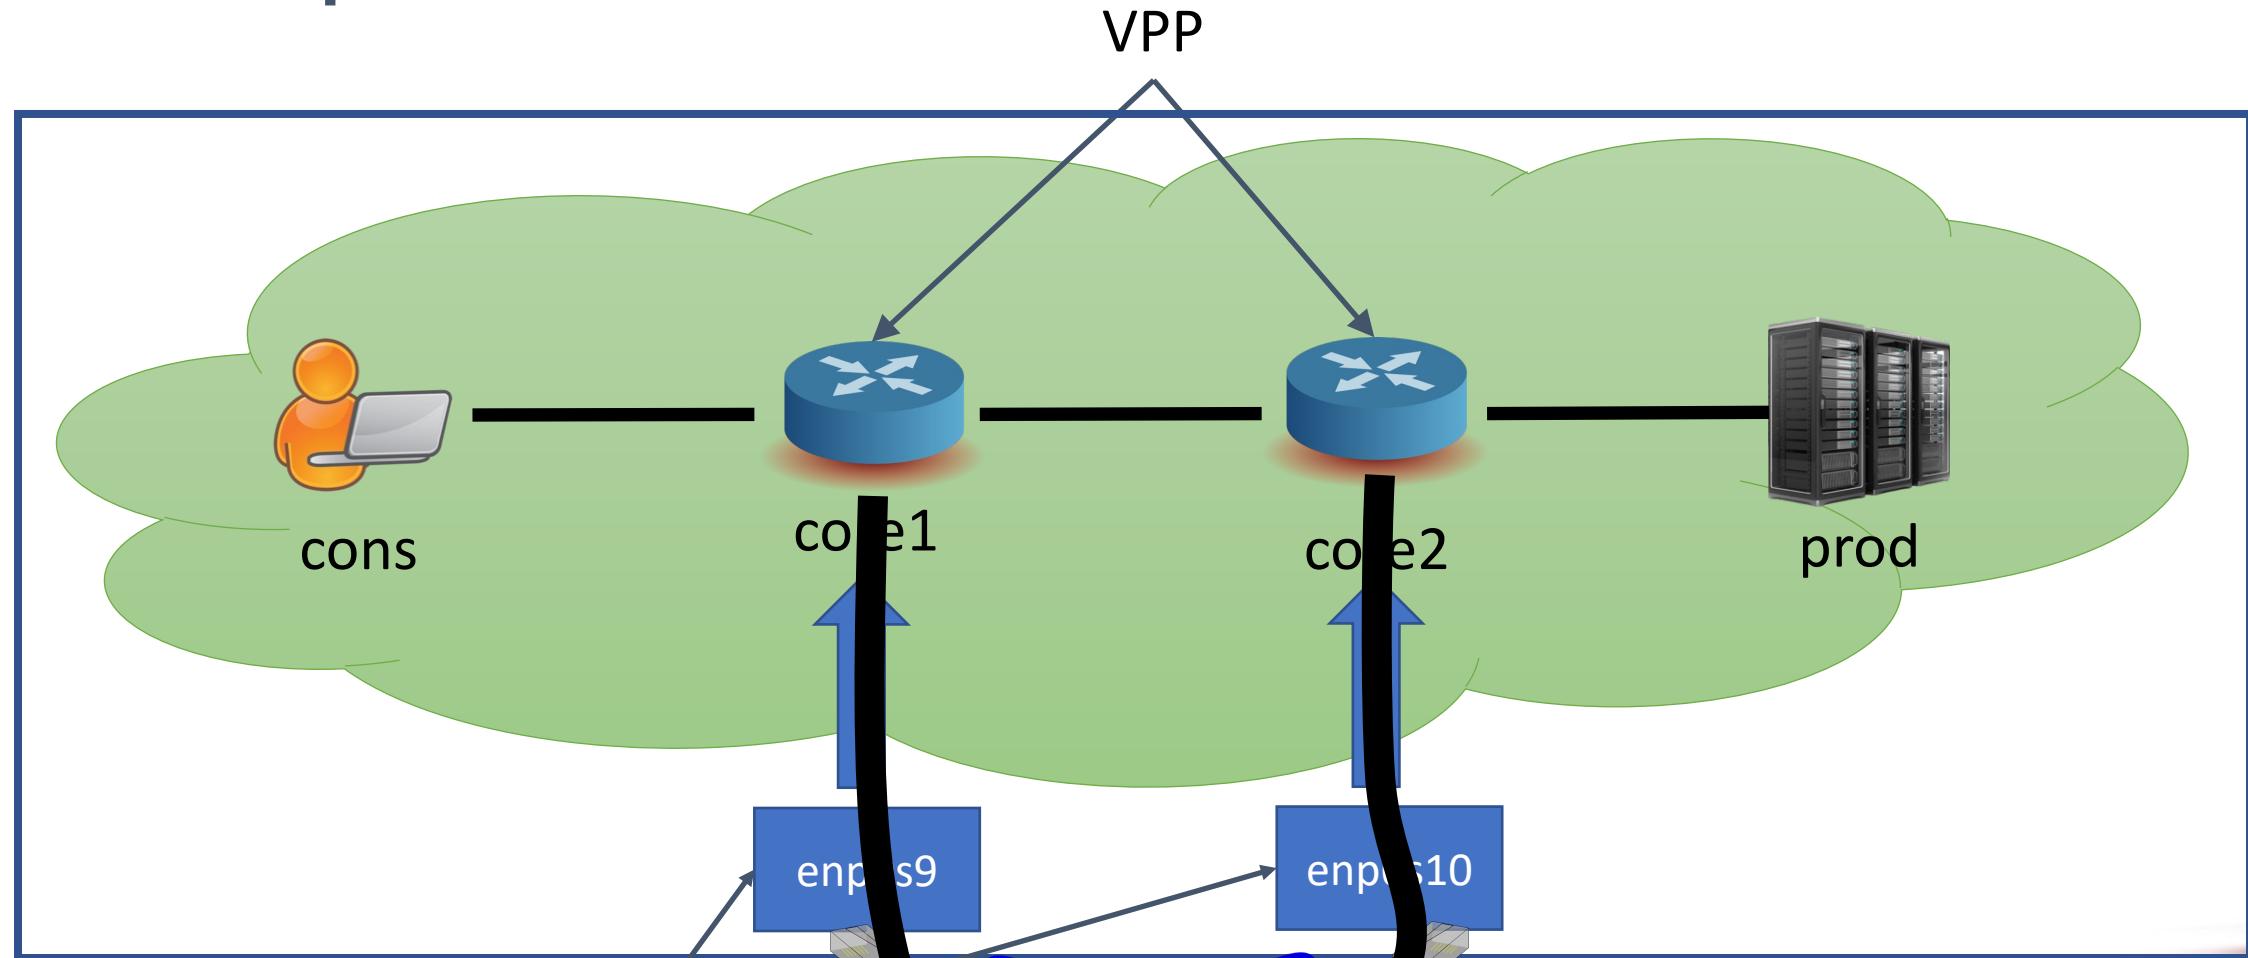
```
wifi_emulator: no process found  
lte_emulator: no process found  
kill: usage: kill [-s sigspec | -n signum | -sigspec] pid | jobspec ...  
or kill -l [sigspec]  
Removing bridge...  
Removing interface...  
Removing stale routes
```

# VPP in vICN

- Objective: learn to setup vICN to use your Intel interfaces
- VPP running in container
- Uses DPDK and zC-forwarding



# Setup



DPDK-compatible  
interfaces

# Identifying the DPDK interfaces

Compare:

```
sudo lshw -c network -businfo
```

with <http://dpdk.org/doc/nics>

```
cicn@cicn-VirtualBox:~$ sudo lshw -c network -businfo
  Bus info          Device      Class      Description
  =====
  pci@0000:00:03.0  enp0s3    network    82540EM Gigabit Ethernet Controller
  pci@0000:00:08.0  enp0s8    network    82540EM Gigabit Ethernet Controller
  pci@0000:00:09.0  enp0s9    network    82545EM Gigabit Ethernet Controller (Copper)
  pci@0000:00:0a.0  enp0s10   network    82545EM Gigabit Ethernet Controller (Copper)
```

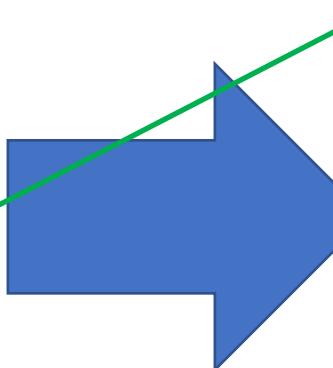
# Declaring the DPDK Interfaces

```
{  
    "type": "DpdkDevice",  
    "name": "core1-dpdk1",  
    "mac_address": "08:00:27:44:9a:38",  
    "node": "core1",  
    "device_name": "enp0s9",  
    "pci_address": "0000:00:09.0"  
},  
{  
    "type": "DpdkDevice",  
    "name": "core2-dpdk1",  
    "mac_address": "08:00:27:18:42:f2",  
    "node": "core2",  
    "device_name": "enp0s10",  
    "pci_address": "0000:00:0a.0"  
},
```

```
[cicn@cicn-VirtualBox:~$ sudo lshw -c network -businfo  
Bus info          Device      Class      Description  
=====  
pci@0000:00:03.0  enp0s3      network    82540EM Gigabit Ethernet Controller  
pci@0000:00:08.0  enp0s8      network    82540EM Gigabit Ethernet Controller  
pci@0000:00:09.0  enp0s9      network    82545EM Gigabit Ethernet Controller (Copper)  
pci@0000:00:0a.0  enp0s10     network    82545EM Gigabit Ethernet Controller (Copper)
```

# Changes to resources

```
{  
    "type": "Link",  
    "src_node": "core1",  
    "dst_node": "core2",  
    "groups": [ "virtual" ]  
},  
{  
    "type": "MetisForwarder",  
    "cache_size": 2000,  
    "node": "core1"  
},  
{  
    "type": "MetisForwarder",  
    "cache_size": 0,  
    "node": "core2"  
},
```



```
{  
    "type": "PhyLink",  
    "src": "core1-dpdk1",  
    "dst": "core2-dpdk1",  
    "groups": [ "virtual" ]  
},  
{  
    "type": "VPP",  
    "node": "core1",  
    "name": "vpp_core1"  
},  
{  
    "type": "CICNPlugin",  
    "node": "core1",  
    "name": "vpp-fwd"  
},  
{  
    "type": "VPP",  
    "node": "core2",  
    "name": "vpp_core2"  
},  
{  
    "type": "CICNPlugin",  
    "node": "core1",  
    "name": "vpp-fwd"  
},
```

# What is vICN actually doing?

- **VPP-ready host**
  - Install (if necessary) the DPDK driver and load it in the host kernel
  - Change driver for DPDK-compatible devices
  - Change number of hugepages for VPP
- **VPP-ready container**
  - Create a privileged container by changing its apparmor profile
  - Add DPDK-enabled interfaces to the container



# What is vICN actually doing? (cont'd)

- Start VPP on the container
  - Create configuration file for VPP in the container
  - Start VPP
  - Set up IP forwarding
- Start CICN plugin in VPP
  - Enable CICN plugin
  - Set up ICN faces and routes



# Launching vicn

```
cicn@cicn-VirtualBox:~/vicn$ sudo vicn/bin/vicn.py -s  
examples/tutorial/tutorial06-acm-icn17-vpp.json
```

[...]

```
2017-09-21 17:48:15,023 - vicn.core.task - INFO - Scheduling task <Task[apy]  
partial<_task_resource_update>> for resource <UUID MetisForwarder-MPDRB>  
2017-09-21 17:48:15,024 - vicn.core.resource_mgr - INFO - Resource <UUID  
MetisForwarder-MPDRB> is marked as CLEAN (99/104)  
2017-09-21 17:48:15,146 - vicn.core.task - INFO - Scheduling task <Task[apy]  
partial<_task_resource_update>> for resource <UUID MetisForwarder-NC33W>  
2017-09-21 17:48:15,148 - vicn.core.resource_mgr - INFO - Resource <UUID  
MetisForwarder-NC33W> is marked as CLEAN (100/104)
```



# Traffic creation

## Producer setup:

- **producer-test**

```
producer-test -D ccnx:/webserver
```

- **Webserver**

```
http-server -p $server_folder -l  
http://webserver
```

## Consumer setup

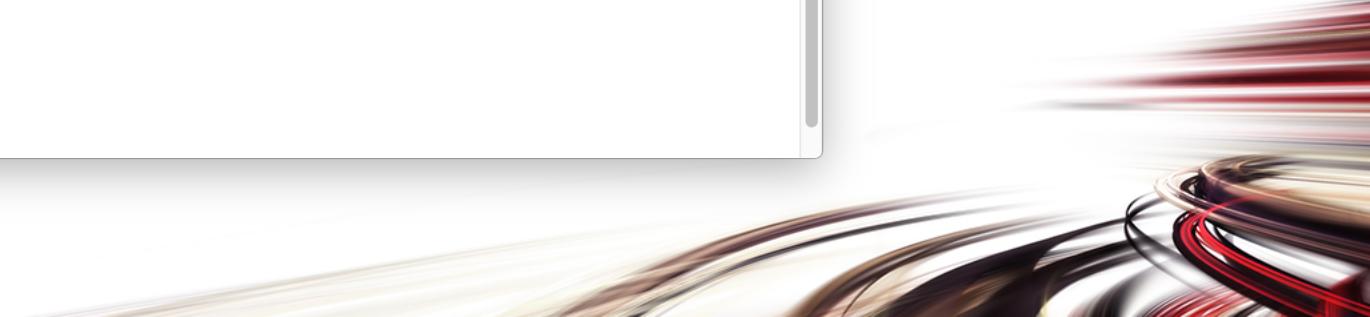
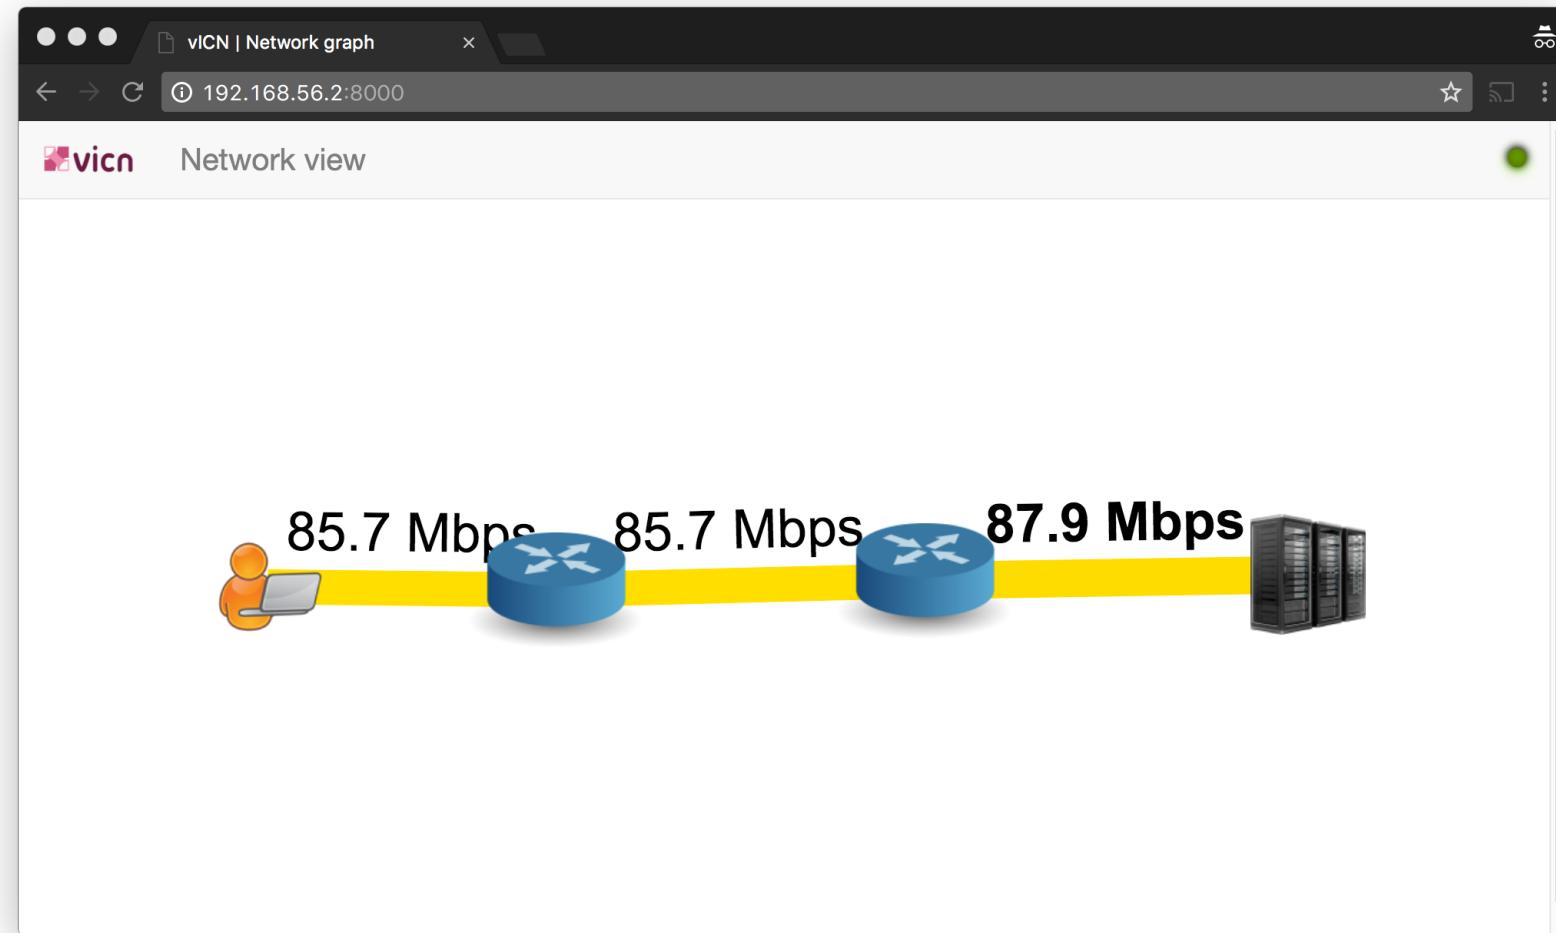
- **consumer-test**

```
consumer-test -D ccnx:/webserver
```

- **iget**

```
iget http://webserver/$filename
```

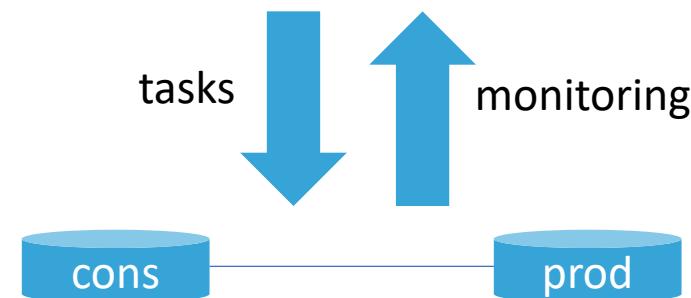
# Traffic visualization on the GUI



# Toward a new Python API

Use python objects instead of static JSON file

```
cons = LxcContainer()  
prod = LxcContainer()  
link = Link(src=cons, dst=prod)
```



# More on vICN

- Demonstration session: new dynamic python API
- Thursday 10:50am: vICN paper presentation



# Available tutorials

In examples/tutorial/:

- tutorial01.json → Simple topology
- tutorial02-dumbell → VPP
- tutorial03-hetnets.json → Wireless emulators
- tutorial06-acm-icn17\* → Today's tutorial (soon)

[https://wiki.fd.io/view/Vicn#Tutorials\\_overview](https://wiki.fd.io/view/Vicn#Tutorials_overview)

# References

vICN wiki: <https://wiki.fd.io/View/Vicn>

vICN paper: <http://conferences.sigcomm.org/acm-icn/2017/proceedings/icn17-26.pdf>

vICN code: git clone -b vicn/master <https://gerrit.fd.io/r/cicn> vicn

# Libicnet: transport layer library for ICN

Mauro Sardara



Tutorial at ACM SIGCOMM ICN, Berlin, Germany  
26<sup>th</sup> of September 2017



# What is Libicnet?

- Library implementing a transport layer and exposing **socket API** to applications willing to communicate through an ICN protocol stack
- Relieves applications from the task of managing layer 4 problems, such as **segmentation and congestion control**
- Enhances the **separation** between Application Data Unit (ADU) and Protocol Data Unit (PDU) processing

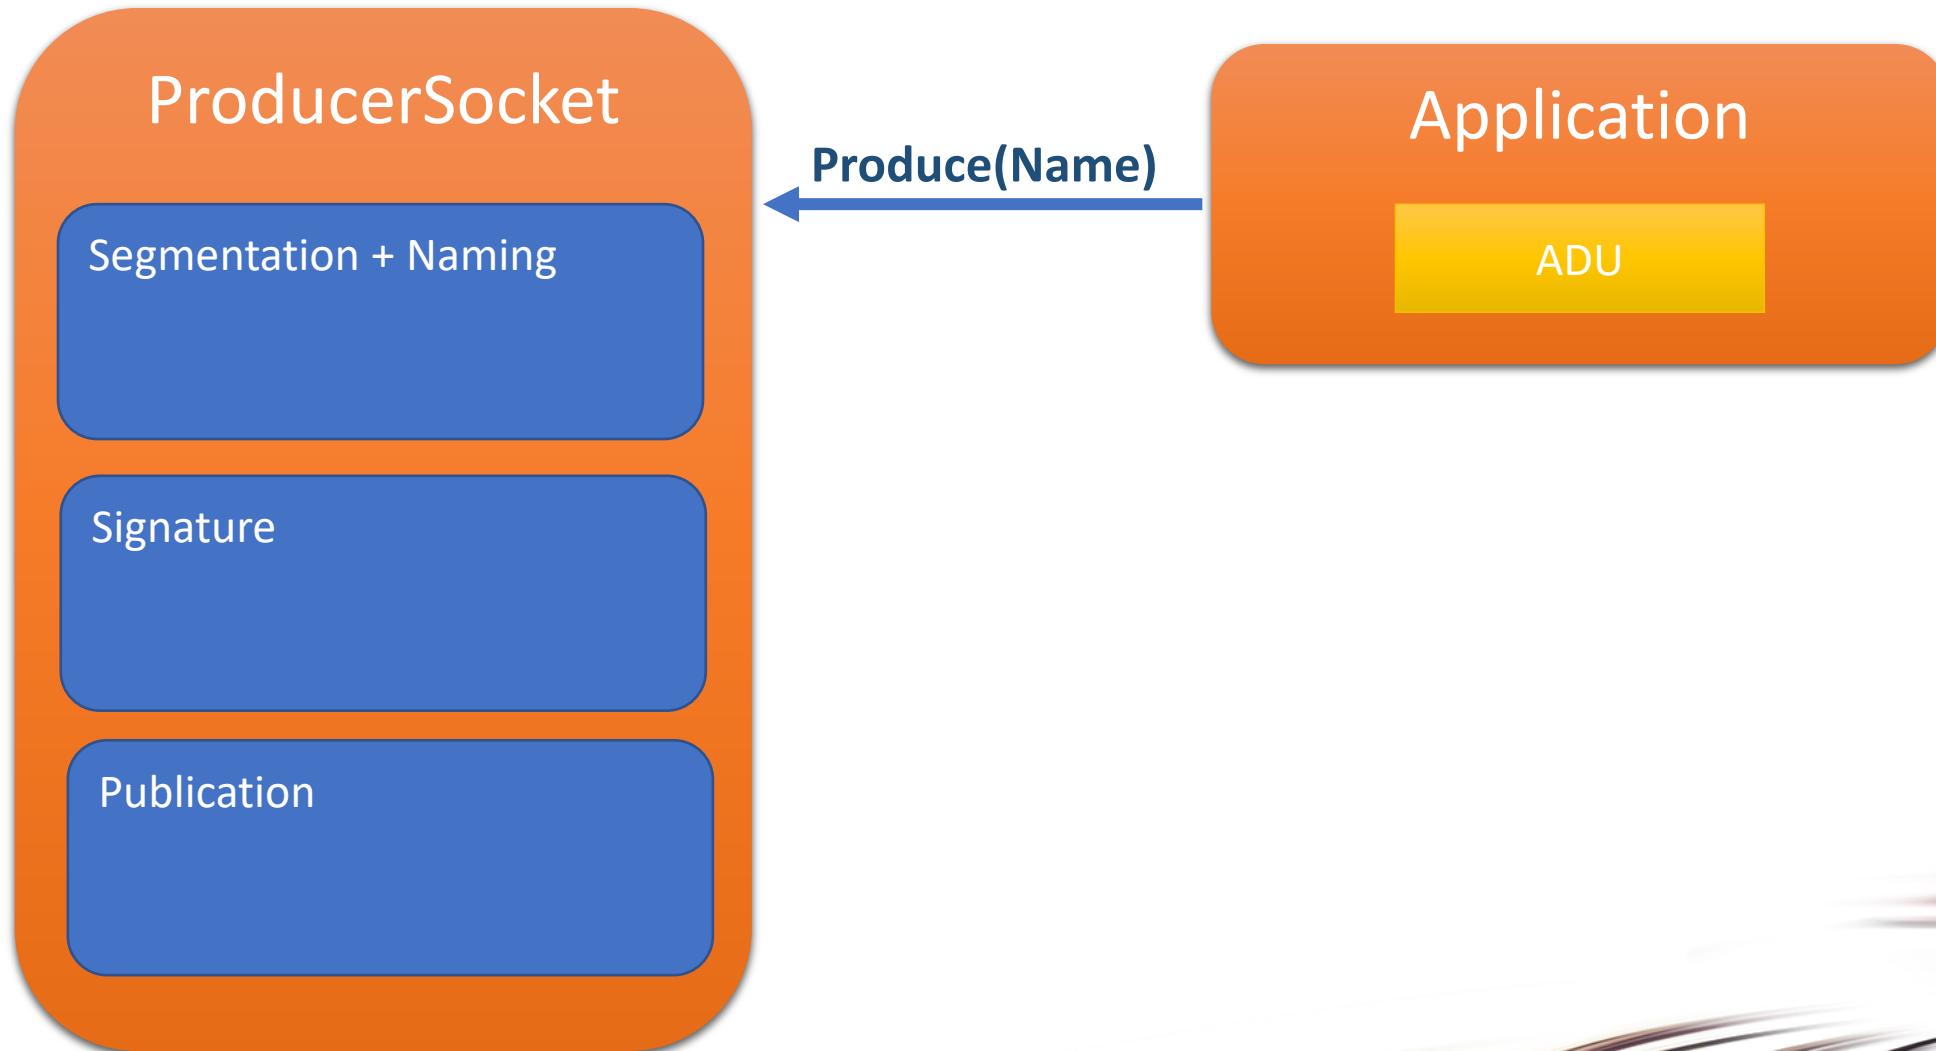


# Core Elements

- ProducerSocket
  - ADU Segmentation and Naming → Layer 4 PDU (ICN Content Object)
  - L4 PDU Signature
  - L4 PDU Publication
- ConsumerSocket
  - Congestion control
  - L4 PDU Fetching
  - Signature verification
  - L4 PDU reassembly → ADU

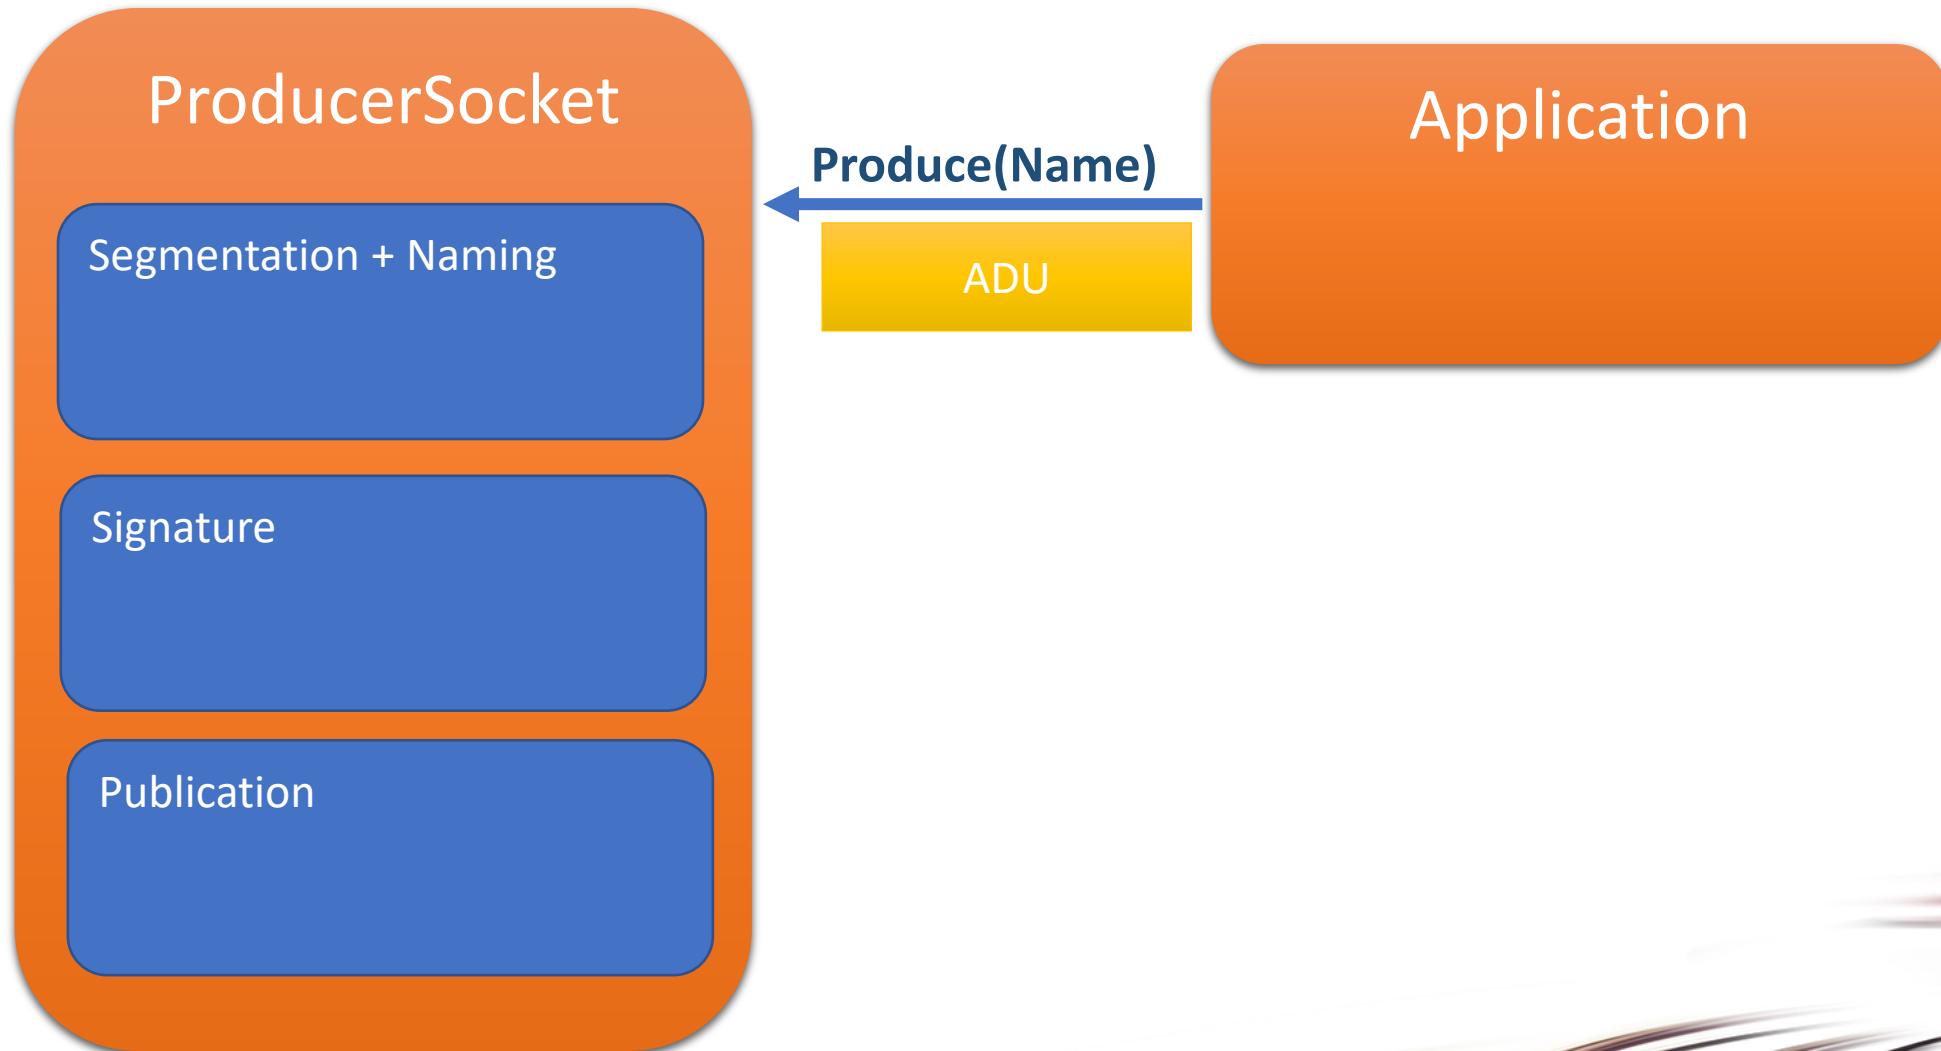


# ProducerSocket

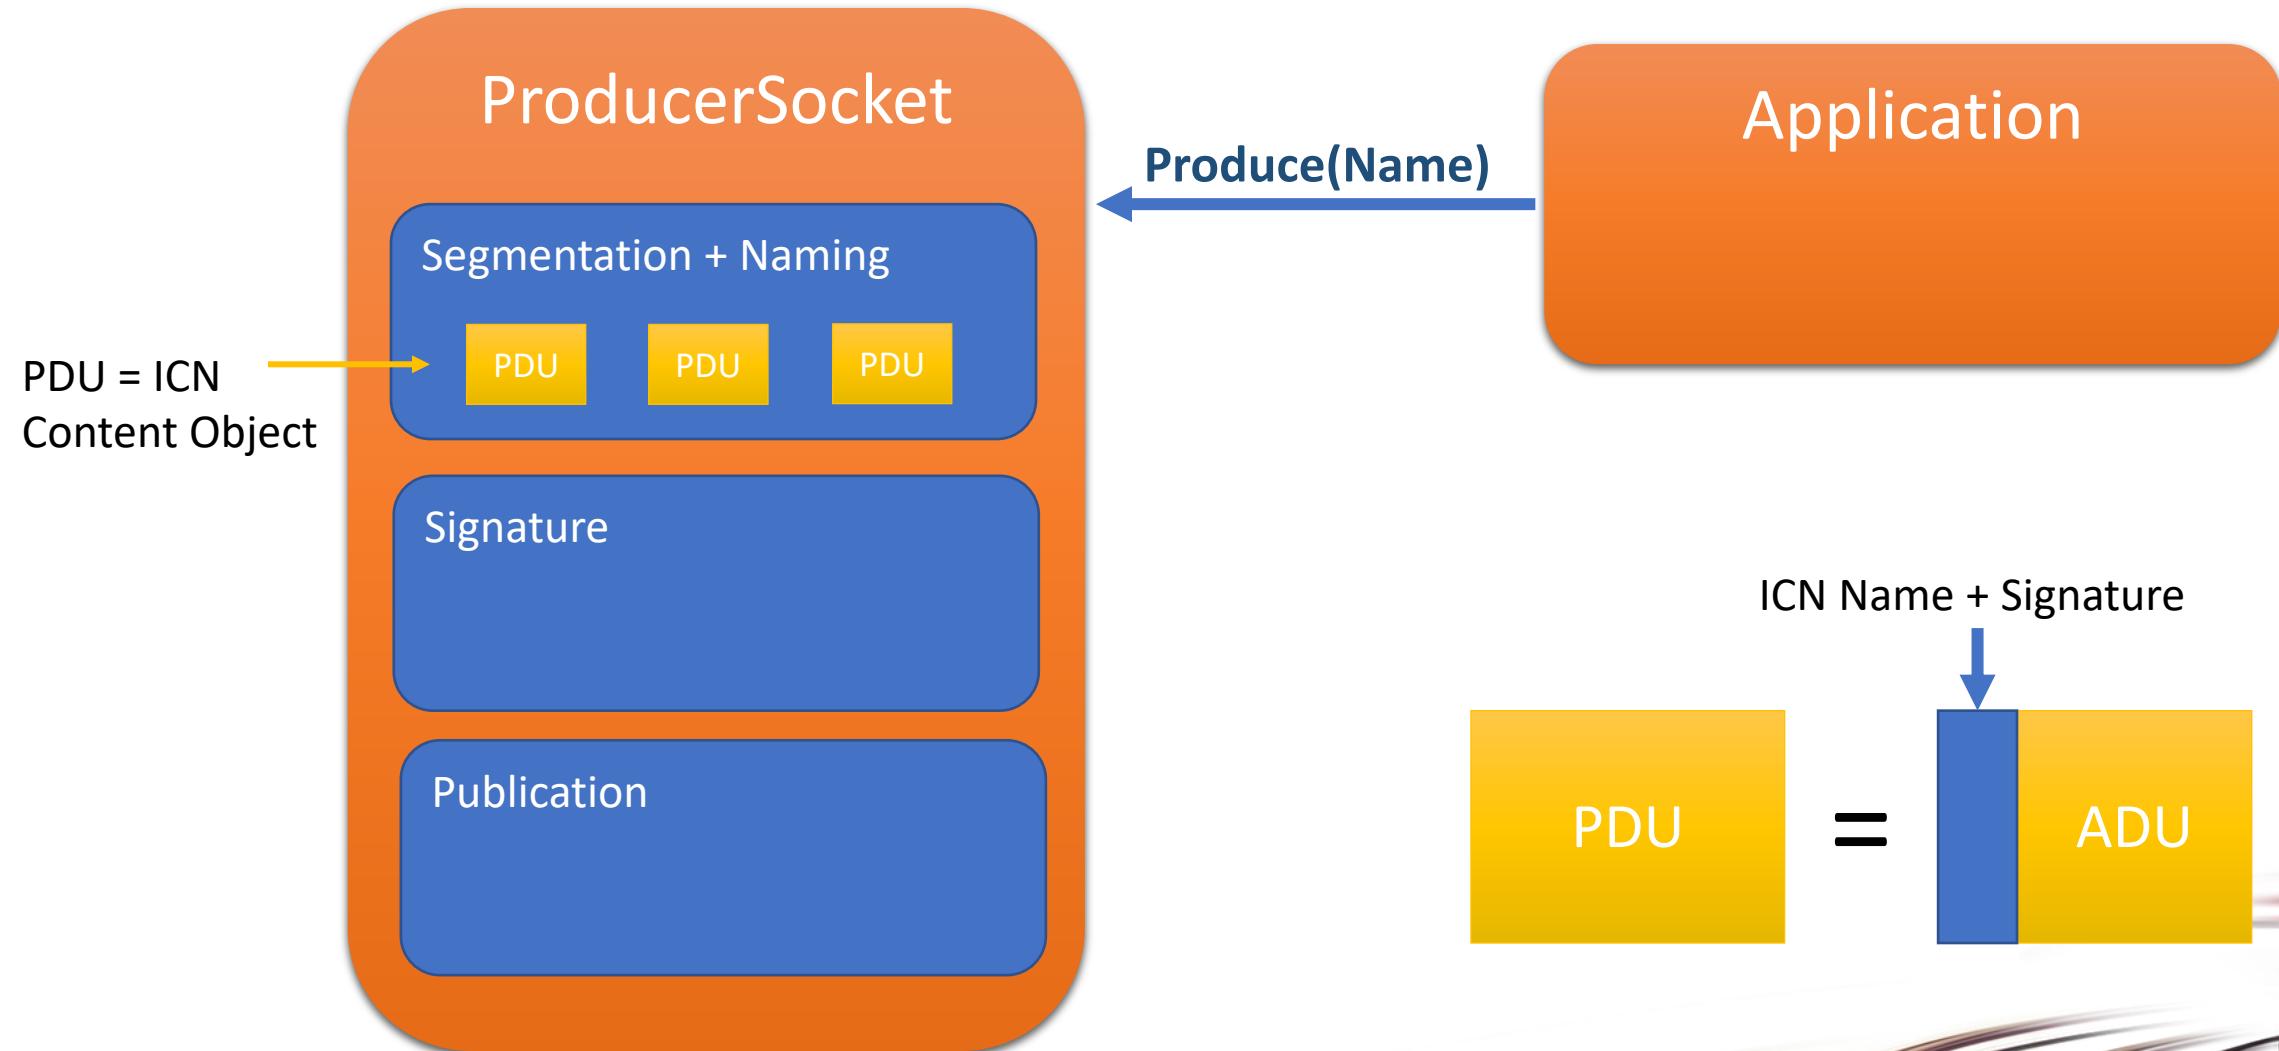




# ProducerSocket

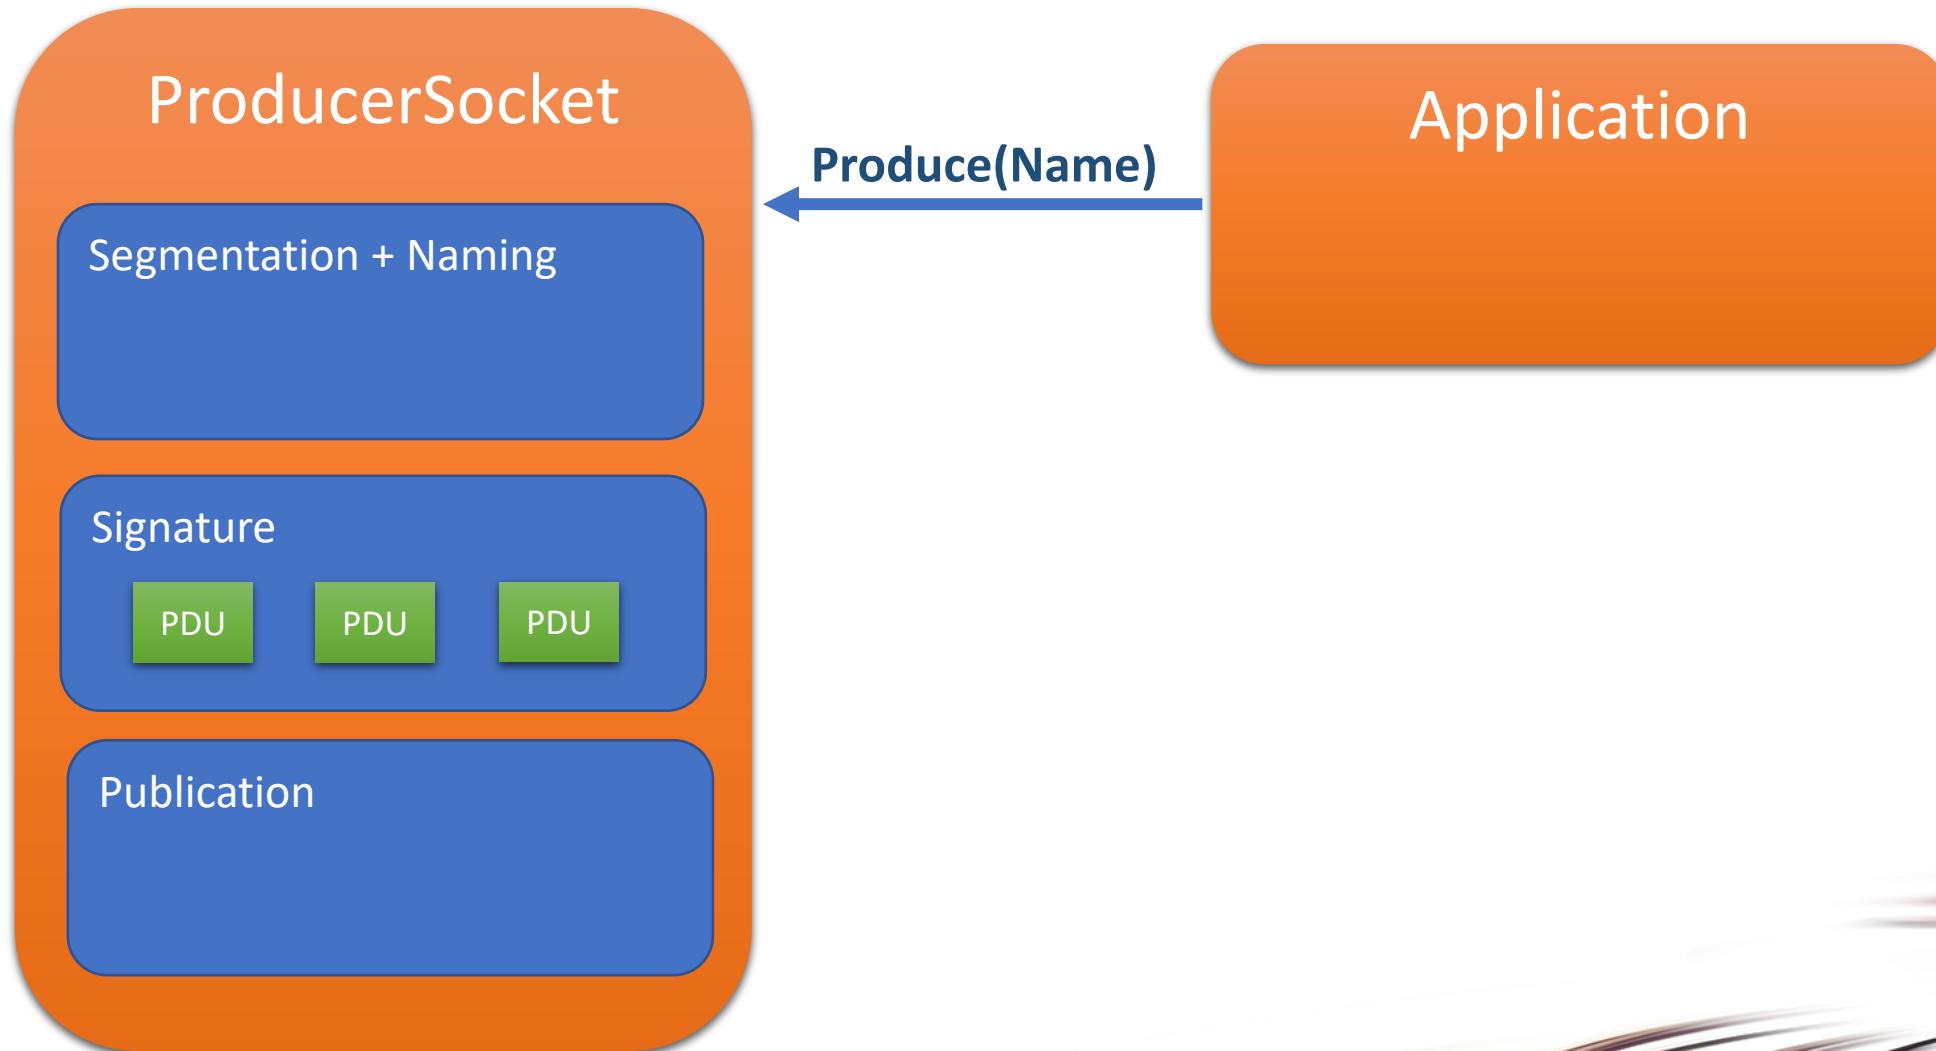


# ProducerSocket

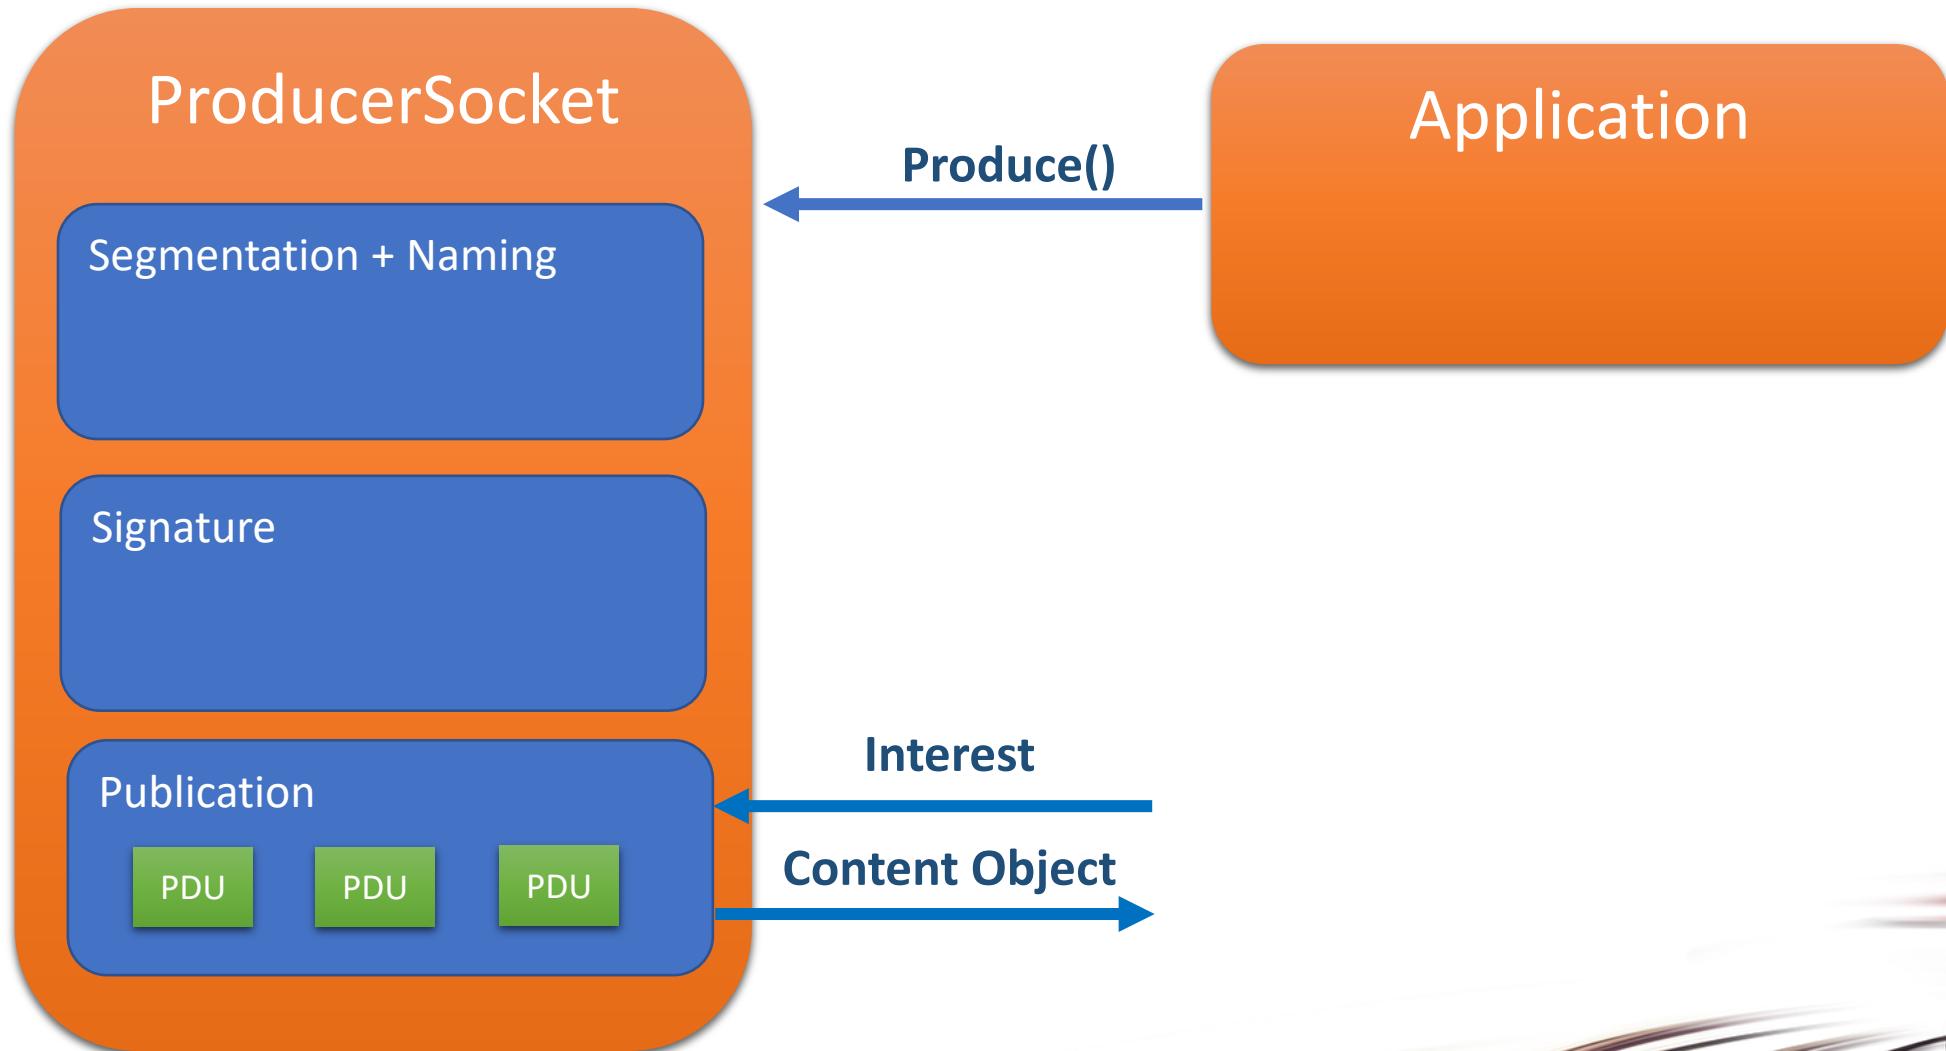




# ProducerSocket

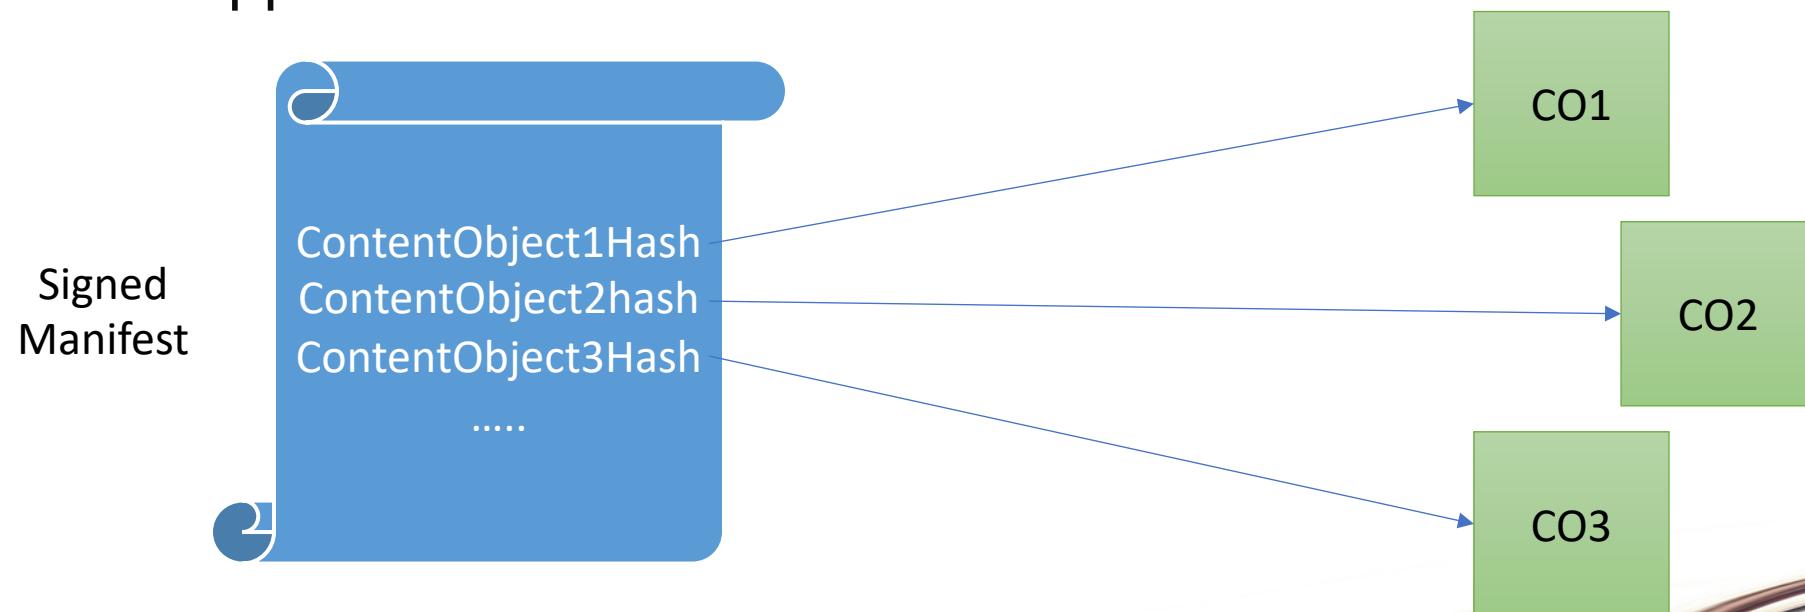


# ProducerSocket



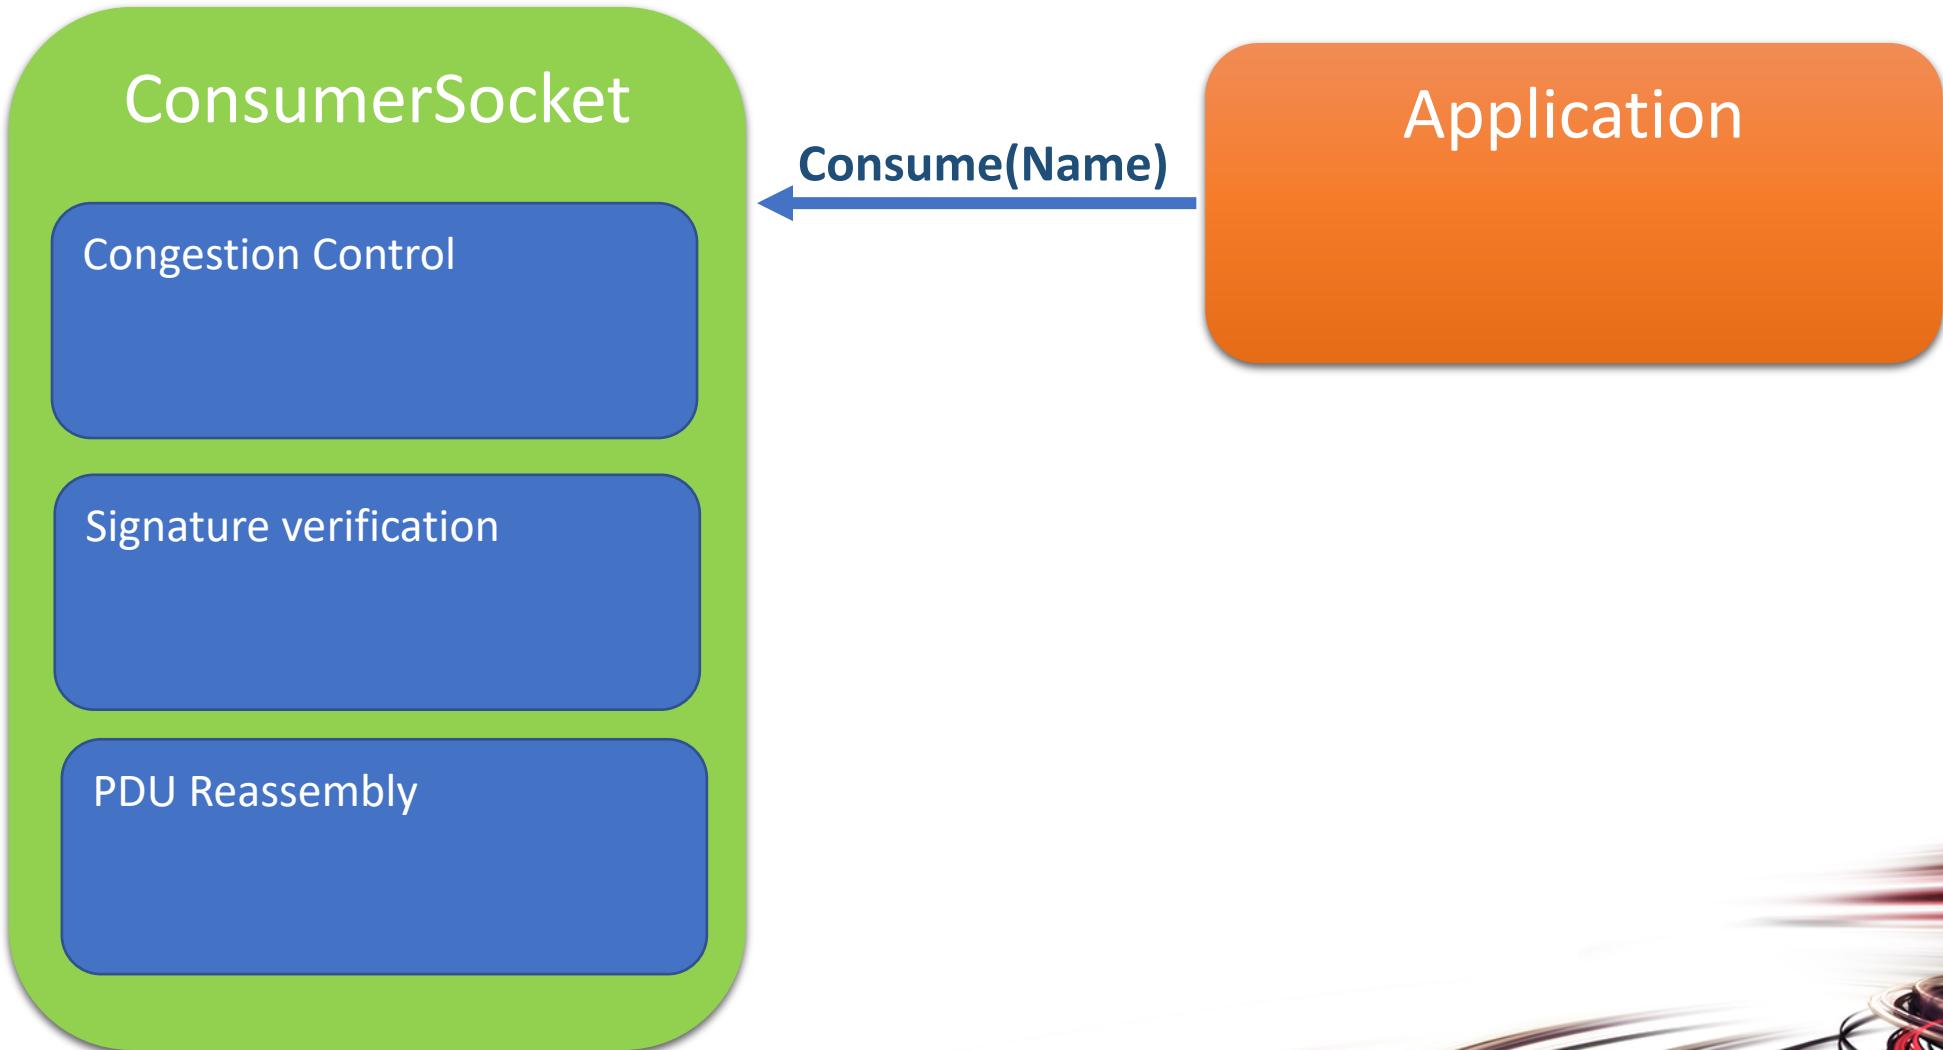
# ProducerSocket

- Signature
  - The application has to provide the library with the information for signing the content objects
  - Signing every content object is computationally expensive: we provide support for manifest

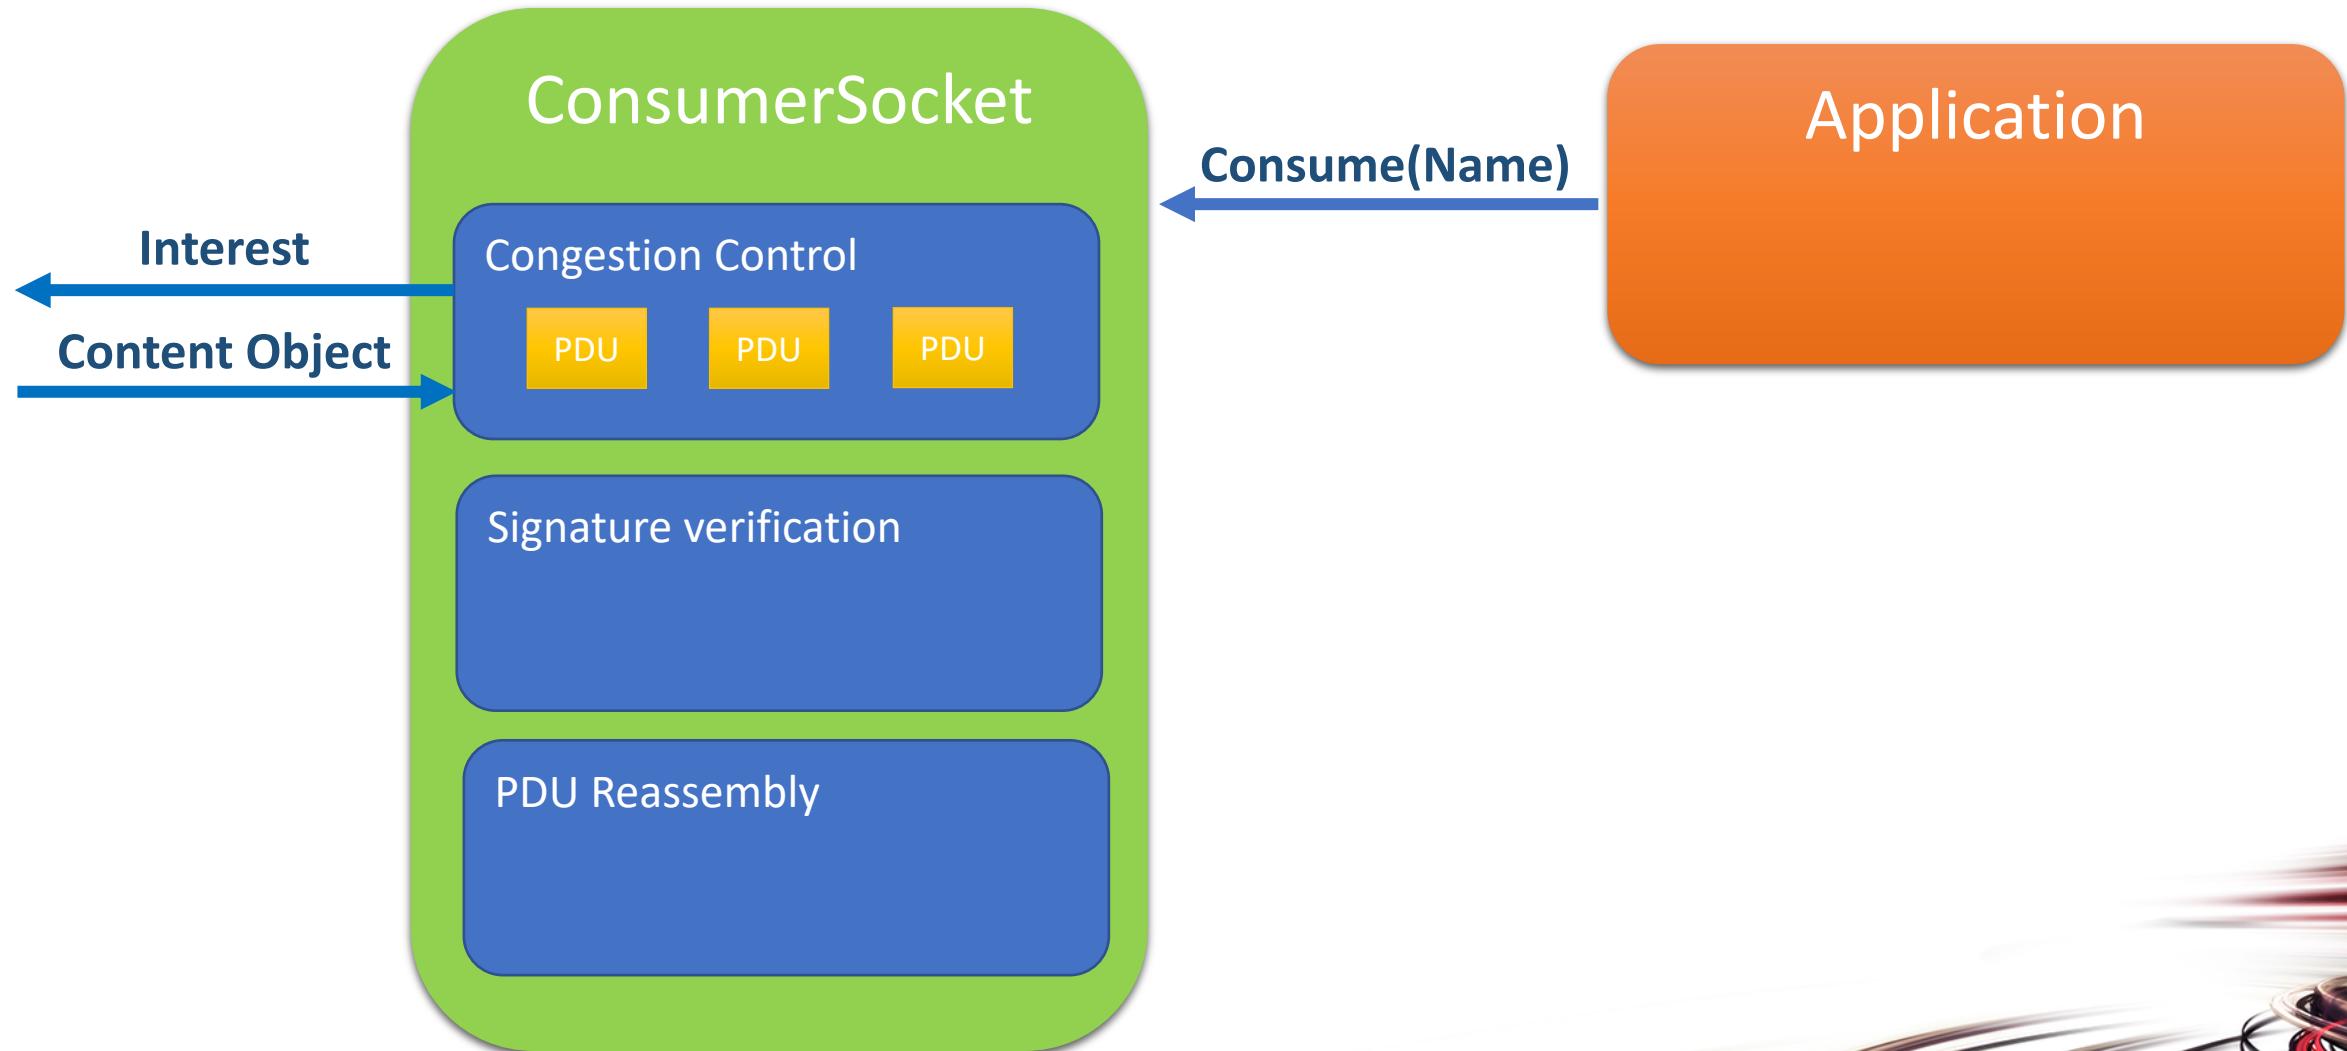




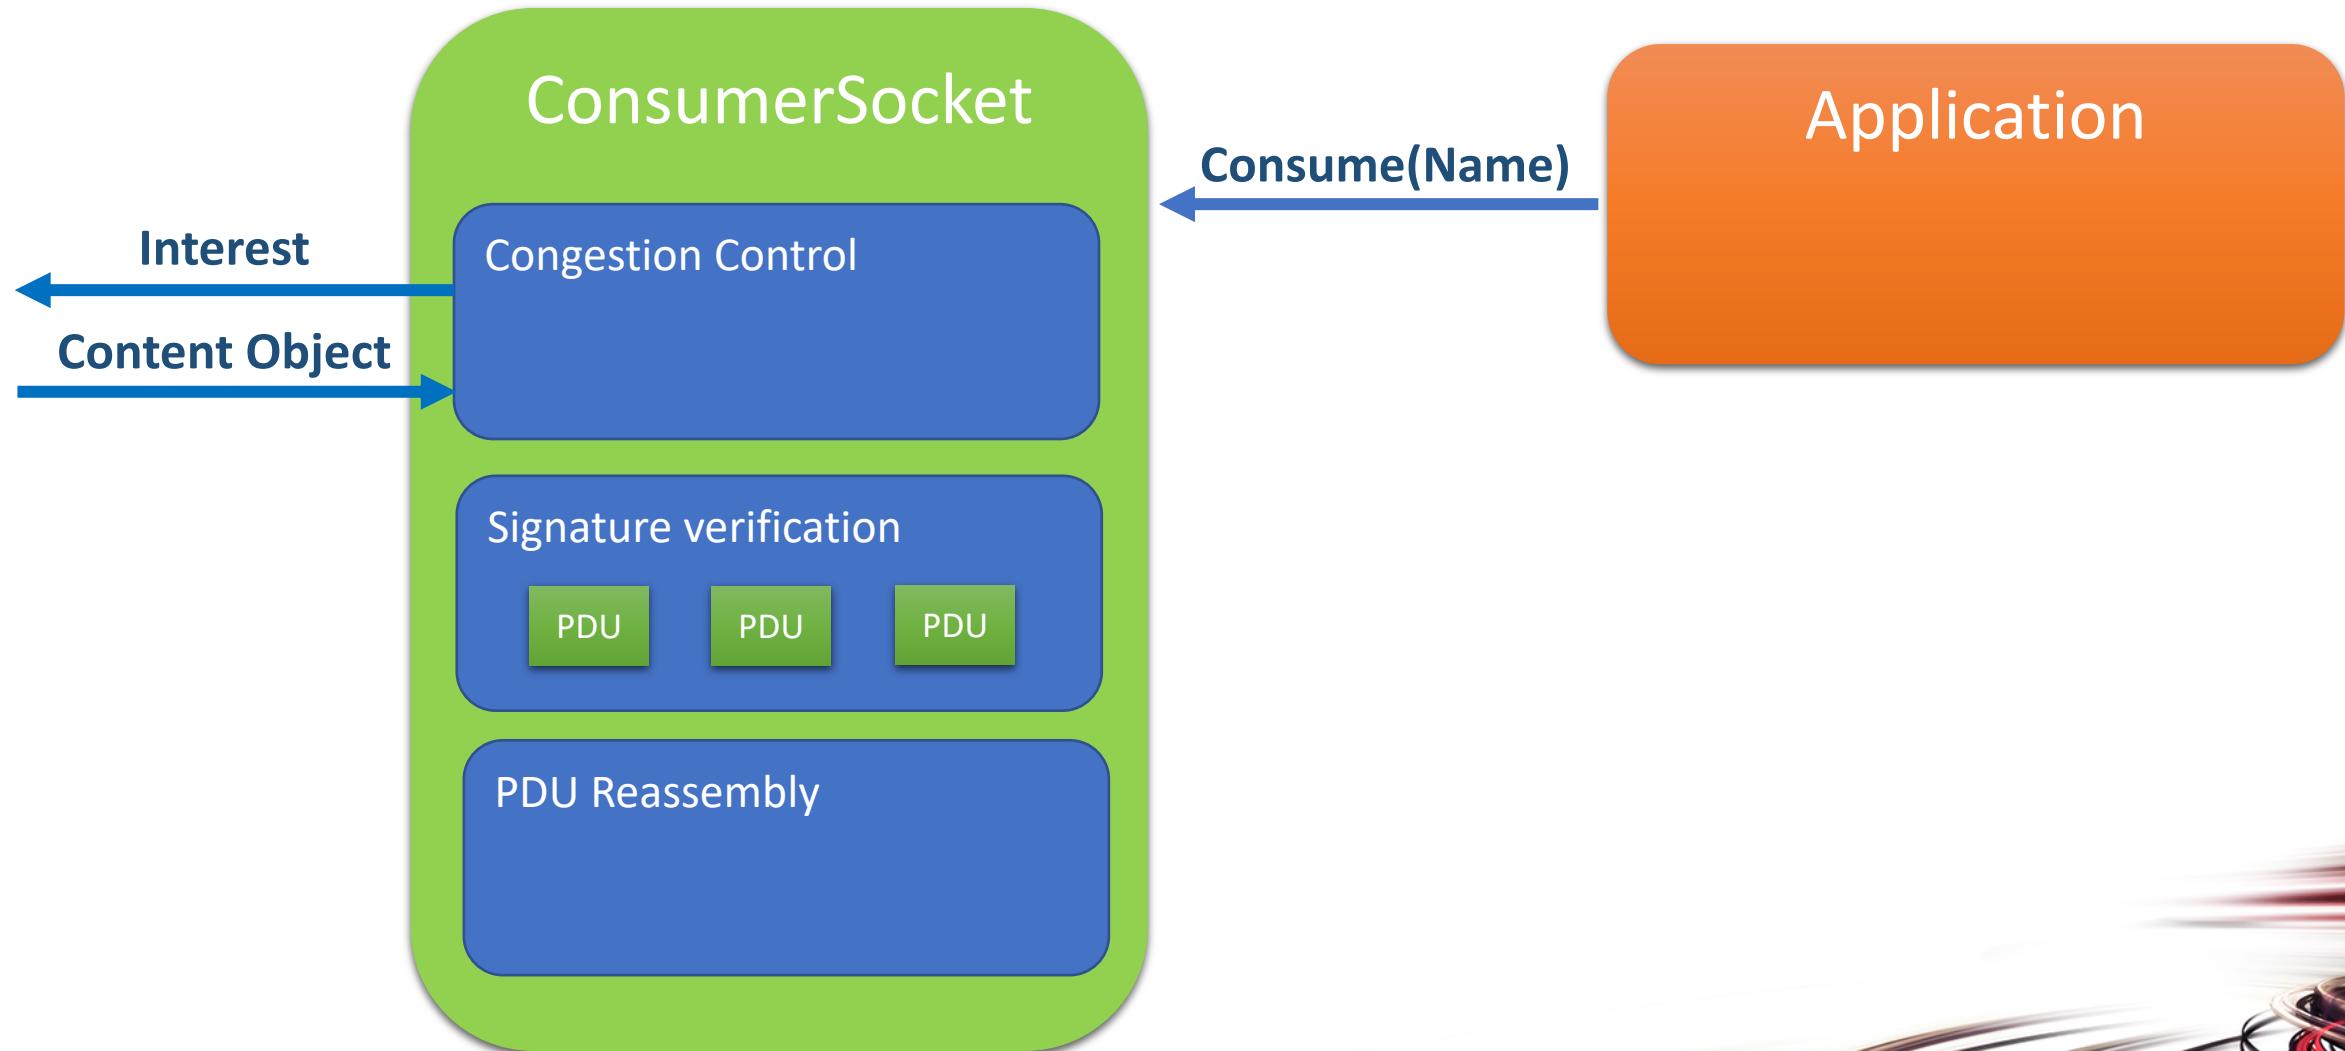
# ConsumerSocket



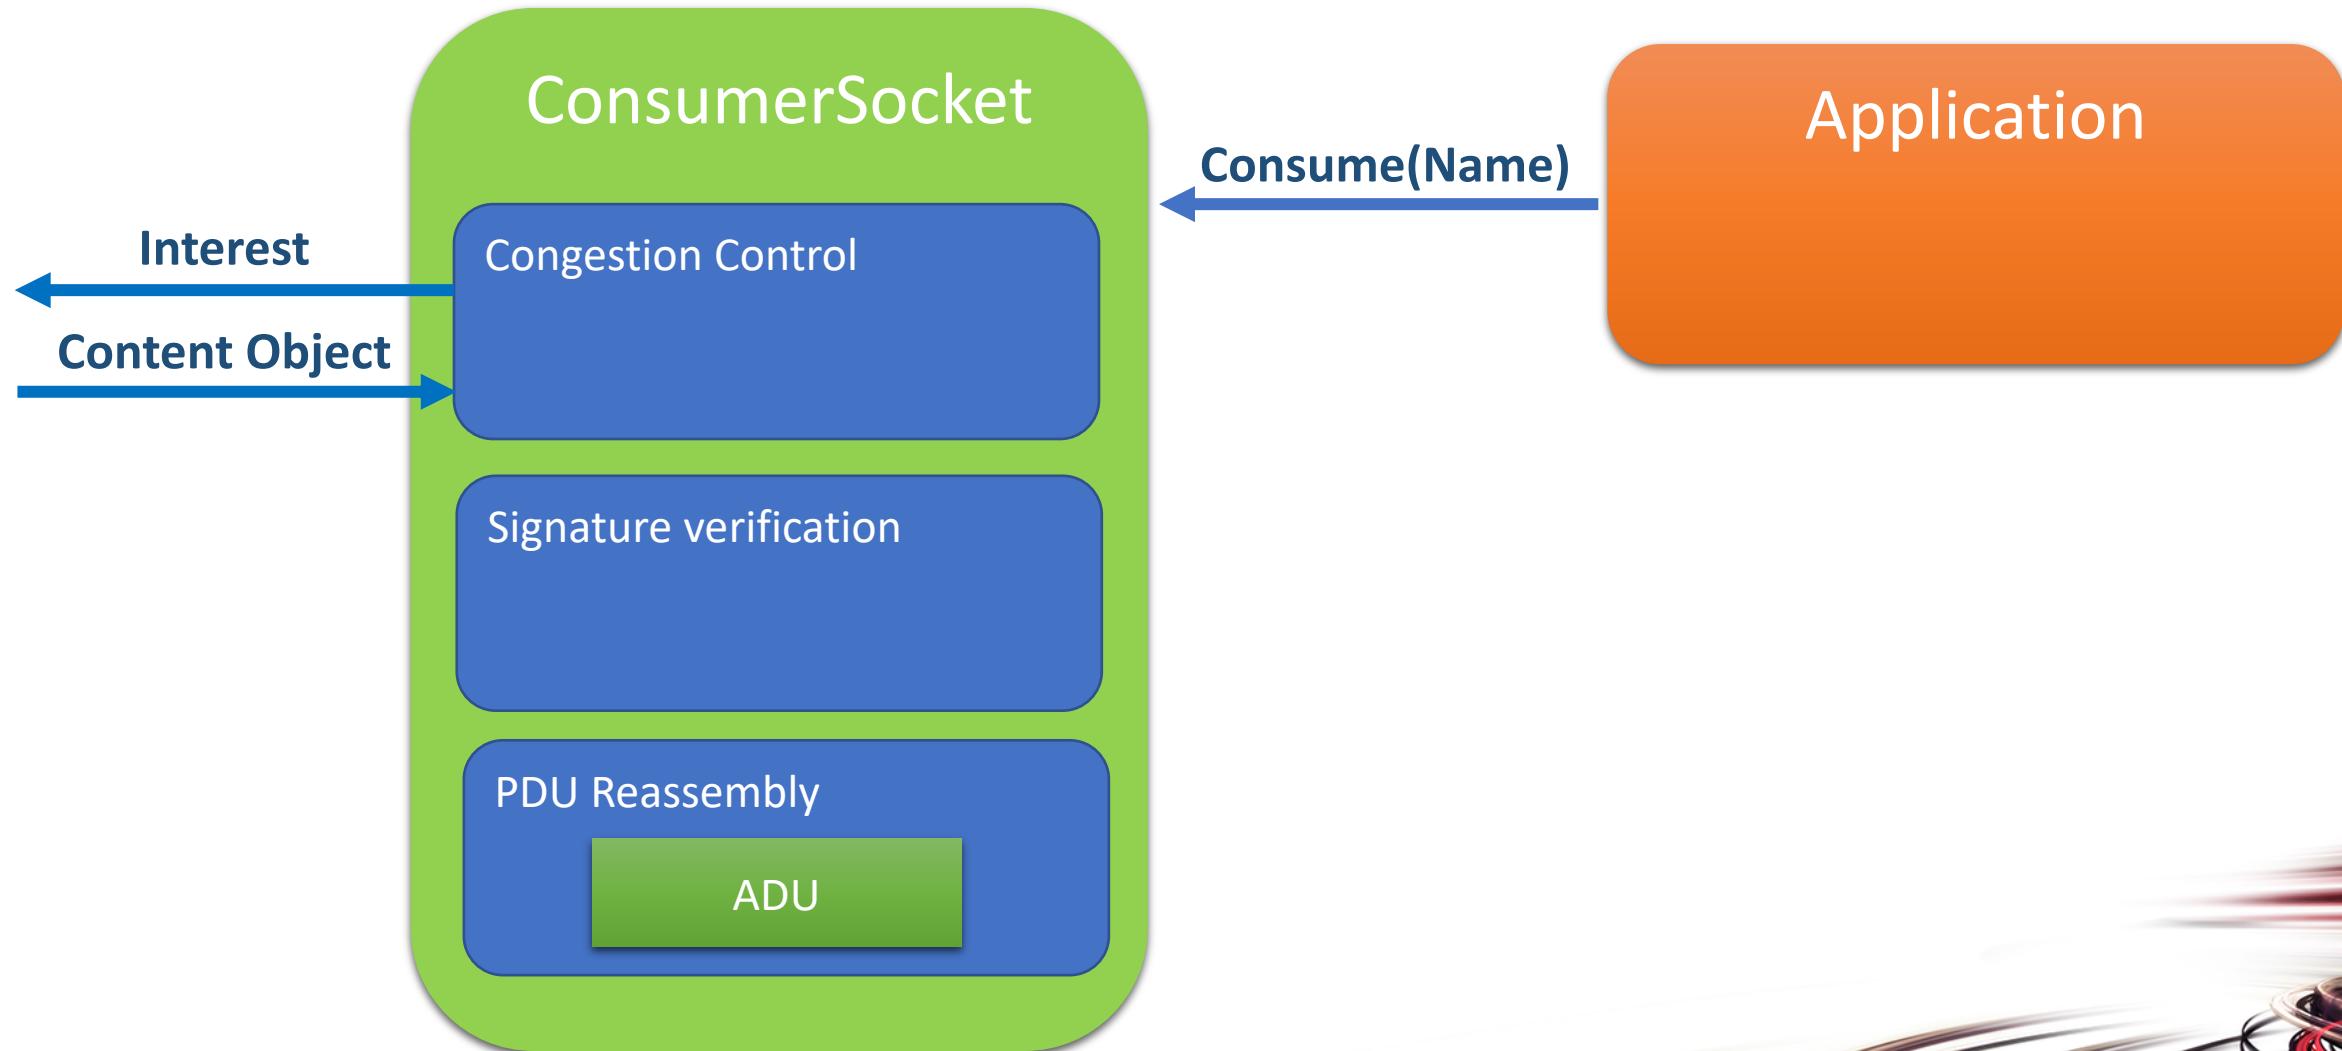
# ConsumerSocket



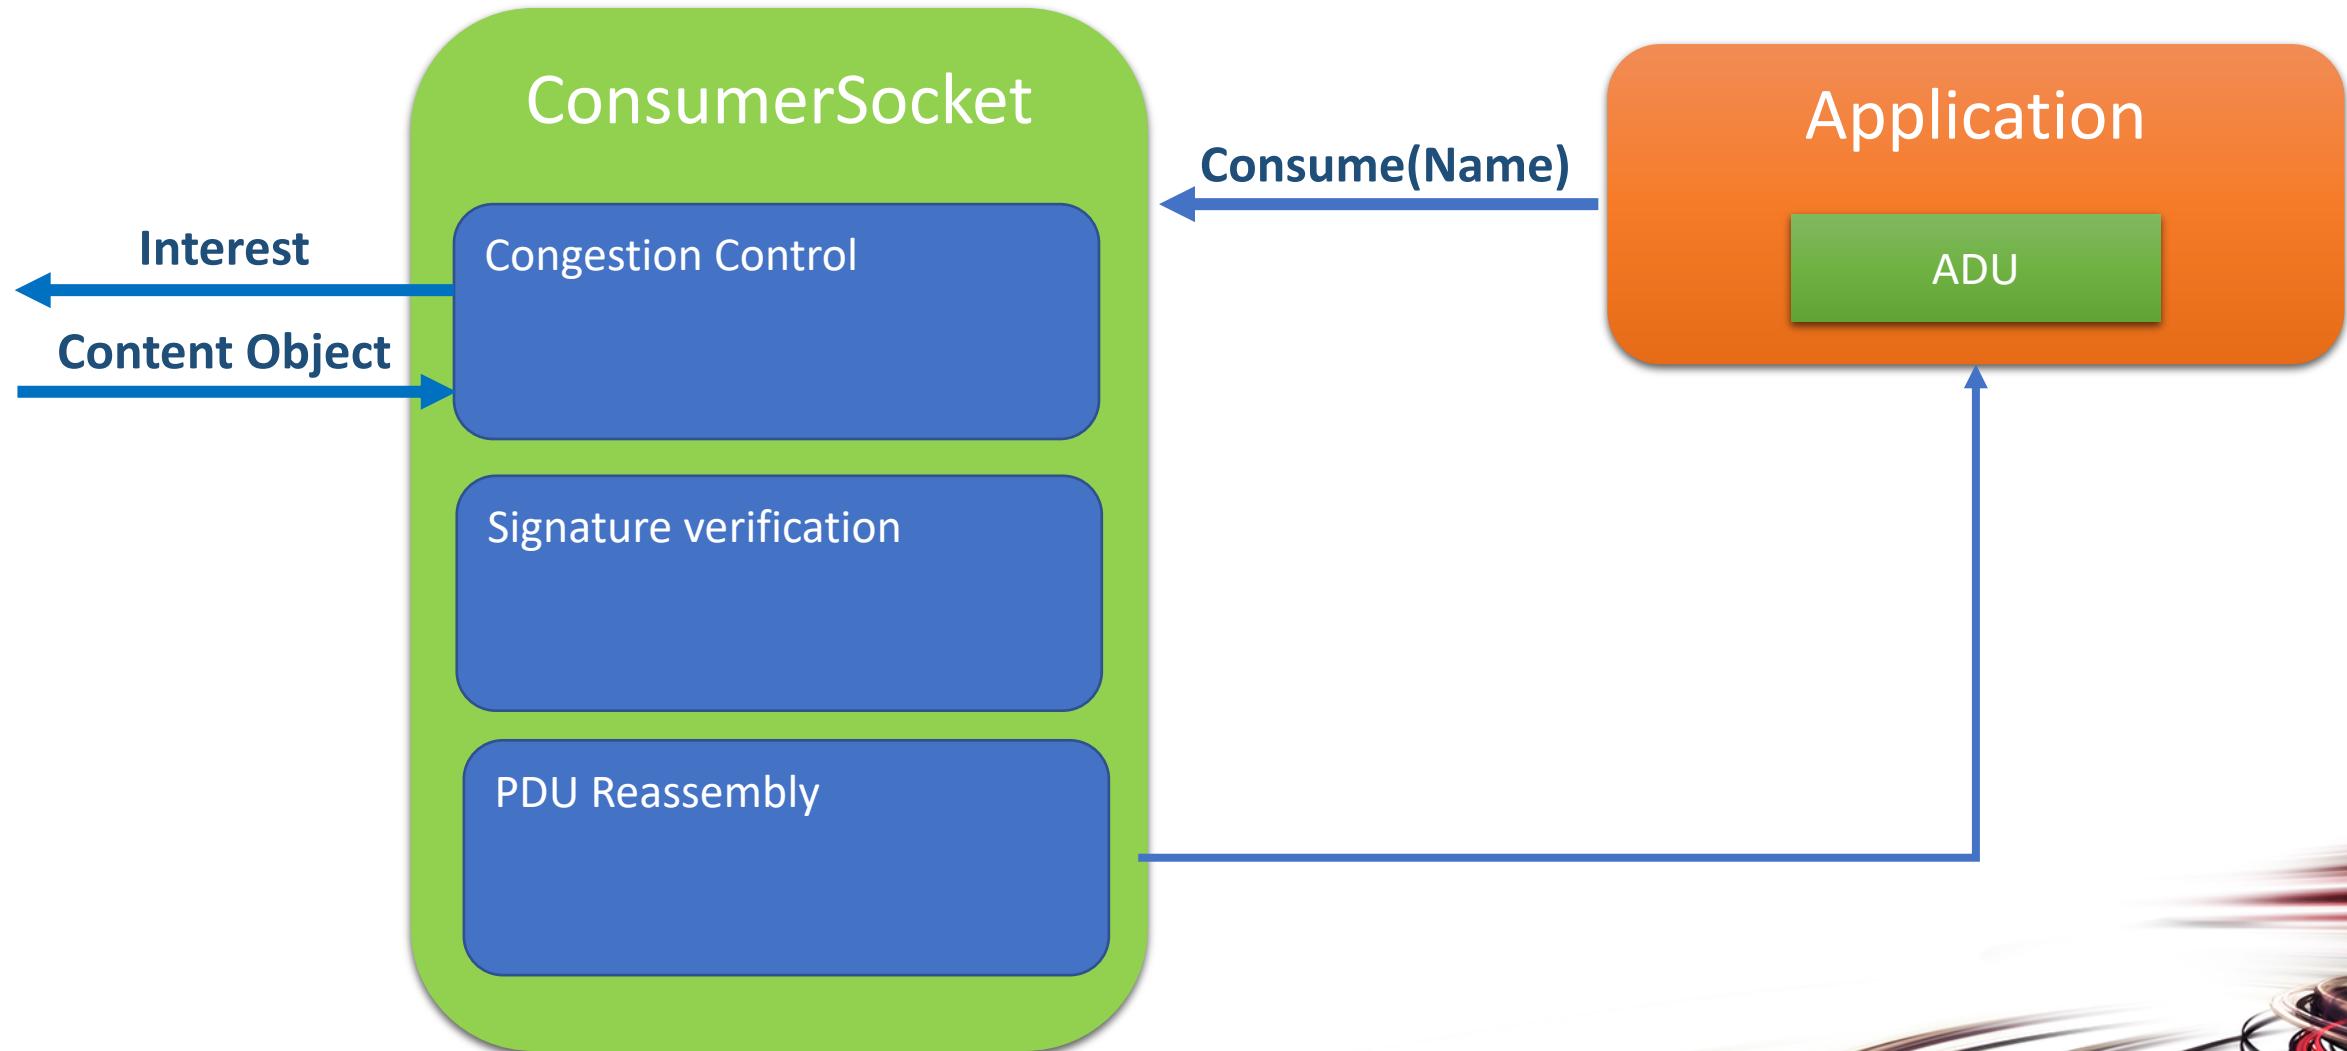
# ConsumerSocket



# ConsumerSocket



# ConsumerSocket





# ConsumerSocket

- Congestion Control
  - Application can choose among a set of algorithms: VEGAS, RAAQM<sup>1</sup>, FIXED\_WINDOW
  - Extension with new algorithms possible
- Signature
  - The application has to provide the library with the information for verifying the signature of the received content objects
  - As the producer case, verifying every content object is expensive: we verify just the manifest signature

<sup>1</sup> G. Carofiglio et al. "Multipath congestion control in content-centric networks," *2013 IEEE Conference on Computer Communications Workshops (INFOCOM WKSHPS)*

# Hands on Libicnet!





# Where to find Libicnet?

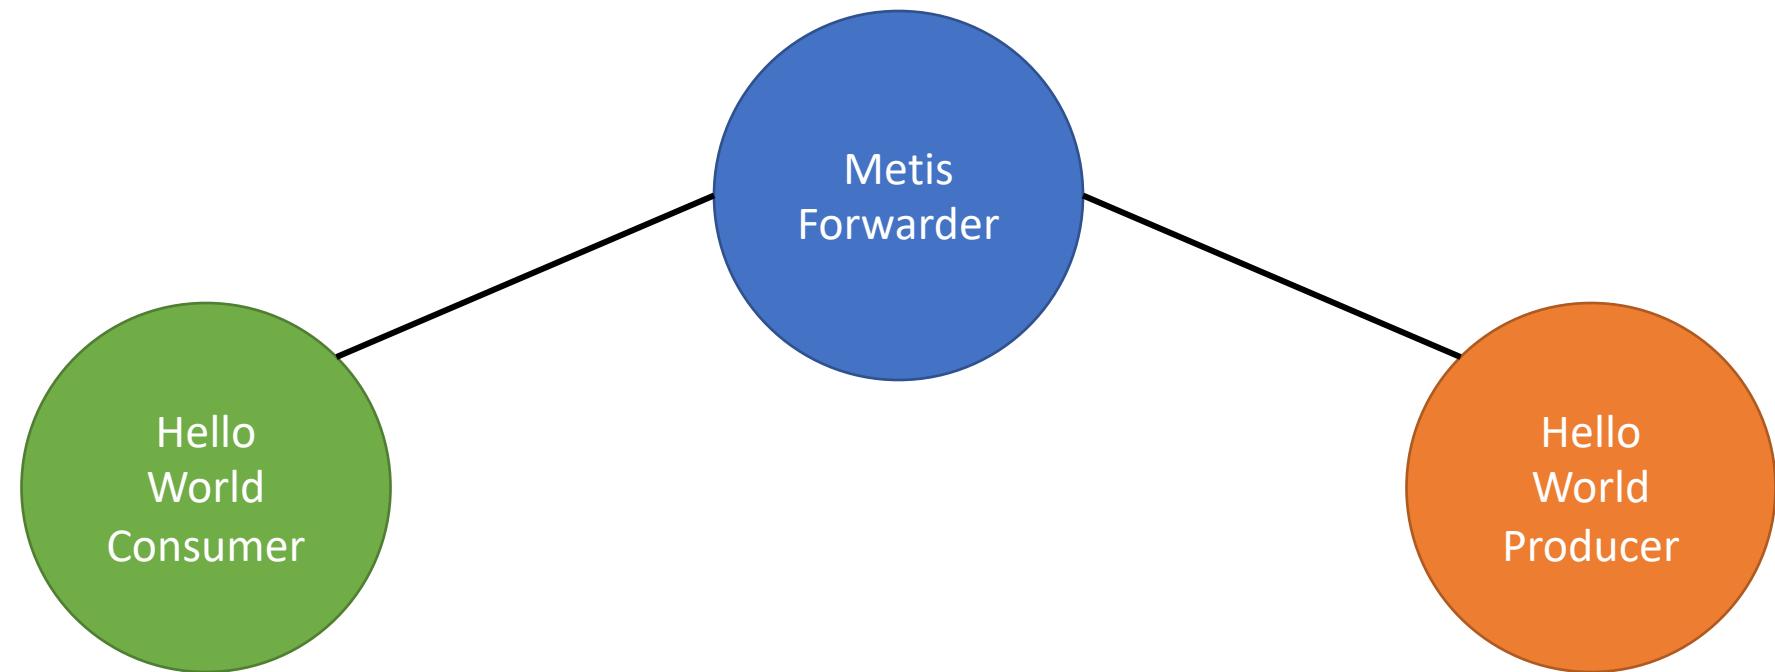
- Wiki page
  - <https://wiki.fd.io/view/Libicnet>
- Code
  - <https://git.fd.io/cicn/log/?h=libicnet/master>



# Hello World Applications

- We will see how building two trivial applications against Libicnet:
  - Hello world Producer
    - It will produce a content of a certain size
  - Hello world Consumer
    - It will pull the content published by the producer

# Topology





# Hello World Producer

```
#include <icnet/icnet_transport_socket_producer.h>
```

```
...  
Name n("ccnx://helloworld");
```

```
ProducerSocket p_(n);
```

```
std::string content(10000, 'A');
```

```
p_.produce(n, (uint8_t *)content.data(), content.size());
```

```
p_.attach();
```

```
p_.serveForever();
```

Routable prefix

Naming, Segmentation,  
Signature, Publication

Local face forwarder-producer  
establishment

# Hello World Consumer

```
#include <icnet/icnet_transport_socket_consumer.h> Congestion control algorithm
...
Consumer c_(Name(), TransportProtocolAlgorithms::RAQM);
c_.setSocketOption(GeneralTransportOptions::INTEREST_LIFETIME, 1001);
c_.setSocketOption(GeneralTransportOptions::MAX_INTEREST_RETX, 25);

c_.setSocketOption(ConsumerCallbacksOptions::CONTENT_RETRIEVED,
                   (ConsumerContentCallback) std::bind(&processContent,
                                         std::placeholders::_1,
                                         std::placeholders::_2));
Name name("ccnx://helloworld");

c_.consume(name);
```

Content Pull +  
Signature Verification +  
Reassembly

Callback called after  
whole ADU will be pulled  
and reassembled

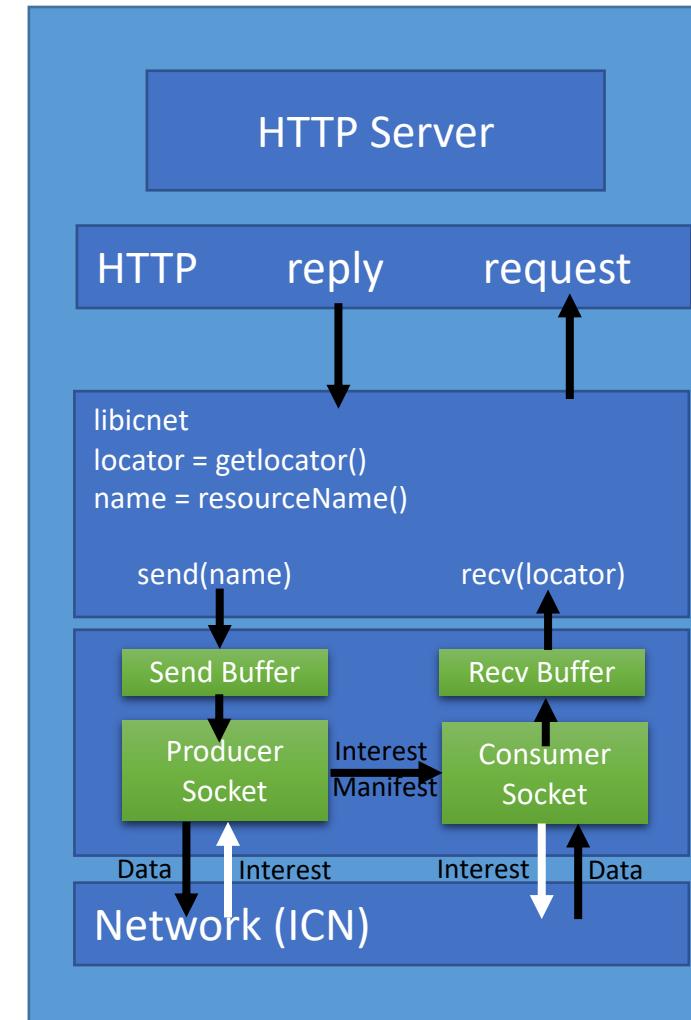
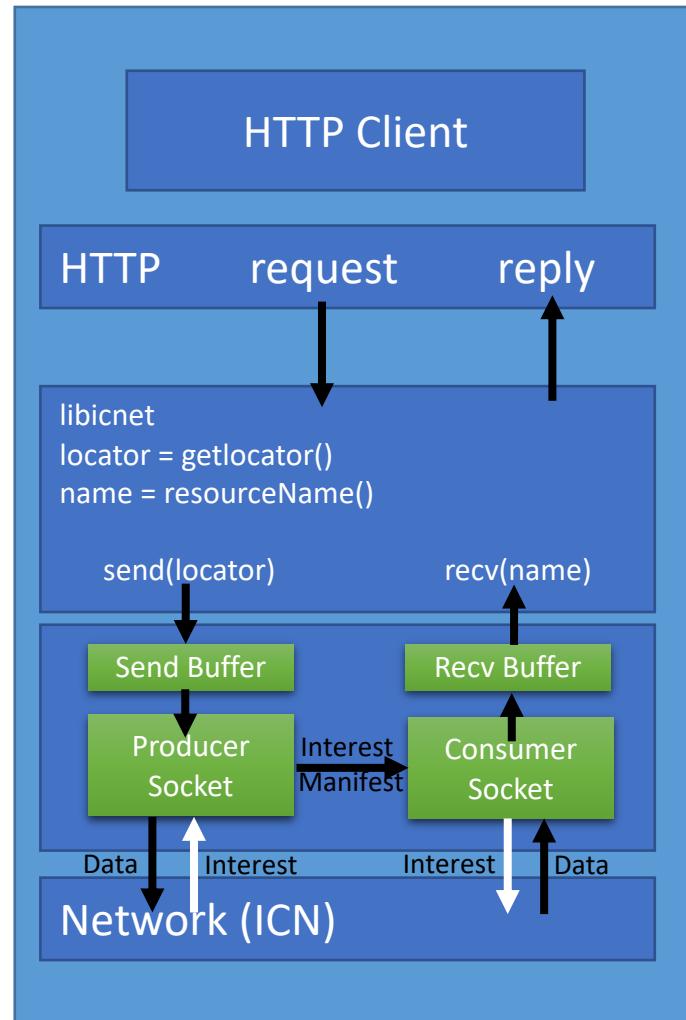
# Callbacks

- The application can register into the library a set of callback allowing to directly handle events during the download/publication.

```
typedef enum {
    INTEREST_OUTPUT = 401,
    INTEREST_RETRANSMISSION = 402,
    INTEREST_EXPIRED = 403,
    INTEREST_SATISFIED = 404,
    CONTENT_OBJECT_INPUT = 411,
    MANIFEST_INPUT = 412,
    CONTENT_OBJECT_TO_VERIFY = 413,
    CONTENT_RETRIEVED = 414,
} ConsumerCallbacksOptions;
```

```
typedef enum {
    INTEREST_INPUT = 501,
    INTEREST_DROP = 502,
    CACHE_HIT = 506,
    CACHE_MISS = 508,
    NEW_CONTENT_OBJECT = 509,
    CONTENT_OBJECT_SIGN = 513,
    CONTENT_OBJECT_READY = 510,
    CONTENT_OBJECT_OUTPUT = 511,
} ProducerCallbacksOptions;
```

# Advanced Example: HTTP support





Thank You!